

# Image artifacts, region-of-interest masking and shortcut learning in medical image classification: a review of work published from 2021 to 2026

MSc thesis *Where the Shortcut Sits* — literature review

October 2026

## 1 Introduction

A classifier takes a shortcut when it relies on something that predicts the label in the training data without being evidence of the disease (Geirhos et al., 2020). In medical imaging such signals are rarely random. Clinicians measure suspicious nodules, mark lesions they intend to excise and drain collapsed lungs, so calipers, skin markings and chest tubes appear more often in the images of sick patients. A common response is to segment the lesion or organ and let the model see only that region of interest (ROI). The thesis asks when this works, and the answer it gives is that it depends on where the artifact lies: masking removes an artifact outside the ROI but keeps one inside it, and the model then relies on it more.

This review covers peer-reviewed work from 2021 to October 2026 that bears on that question. We kept to journals and main-track conferences with an established reputation in medical imaging and machine learning, and left out preprints and workshop papers; a few older papers are cited where later work builds on them. The review is organised in seven parts: the evidence that medical classifiers take shortcuts, the artifacts documented in each modality, what is known about restricting a model to the ROI, how shortcuts are detected, how they are removed, why models prefer them, and how claims of this kind should be evaluated. A final section states where the thesis stands relative to this work.

## 2 Evidence that medical classifiers take shortcuts

The COVID-19 literature made the problem hard to ignore. DeGrave et al. (2021) showed that radiographic COVID-19 detectors relied on laterality markers, patient positioning and other differences between the source datasets, and that their accuracy fell on data from other hospitals; they used saliency maps and generative counterfactuals to show what the models attended to. A systematic review of 62 published COVID-19 imaging models found none fit for clinical use, mainly because of bias in small datasets and the variability of large ones assembled from many sources (Roberts et al., 2021). The most direct demonstration came from Maguolo and Nanni (2021), who still classified COVID-19 radiographs after blacking out the lungs. If the label can be predicted with the anatomy removed, the information must be somewhere else in the image. This is exactly the reasoning behind ROI masking, and it is worth noting that it says nothing about shortcuts that sit *inside* the anatomy.

A second line of work concerns patient attributes. Deep models predict self-reported race from medical images, even when the images are heavily degraded (Gichoya et al., 2022), and chest-radiograph classifiers underdiagnose under-served groups, most strongly at intersections such as Hispanic female patients (Seyyed-Kalantari et al., 2021). Glocker et al. (2023) found that

disease models encode protected characteristics in their features, and Yang et al. (2024) showed that the models encoding demographics most strongly have the largest fairness gaps, and that a model made fair in one hospital is not necessarily fair in another. Encoding an attribute is not the same as using it, and Brown et al. (2023) proposed shortcut testing to distinguish the two by varying how strongly the attribute is encoded. Jones et al. (2024) argue that dataset biases which look alike in the data can arise from different causal mechanisms (prevalence, presentation and annotation) and need different remedies. We find this framing useful: the artifacts in the thesis are produced by clinical actions taken on suspicion, so they are part of how the data are generated and cannot be cleaned away as noise.

Several groups study shortcuts under controlled conditions. Stanley et al. (2024) built synthetic brain images with known disease and bias effects, compared three mitigation strategies and found reweighing to work best; Sun et al. (2023) added a tag, hyperintensities or an obstruction to chest radiographs and found that explanation methods differ widely in whether they reveal the injected signal. These designs fix where the bias is placed. None of them varies its position relative to the region that matters, which is the variable the thesis manipulates.

### 3 Artifacts documented in practice

Dermoscopy has the richest record. Ruler scale bars changed the melanoma predictions of a market-approved classifier (Winkler et al., 2021), as surgical skin markings had done in an earlier study (Winkler et al., 2019). Explanation-based audits of ISIC models uncovered colour calibration patches, rulers and other artifacts (Anders et al., 2022; Pahde et al., 2023). Hair that covers the lesion degrades classifiers (Wang et al., 2024), and for dark corners, adding synthetic ones during training helped more than removing the real ones (Pewton et al., 2024). The ISIC collections also contain duplicate images across releases, which can leak between training and test sets (Cassidy et al., 2022). Read together, these studies show that the artifacts sit in different places: patches and dark corners lie mostly outside the lesion, rulers near the image edge, markings around the lesion, and hair crosses it. The studies report the effect of an artifact being present; none asks whether its position matters.

In ultrasound the evidence is thinner. Lin et al. (2024) showed that calipers and on-screen text act as shortcuts in fetal ultrasound, even for segmentation. We found no peer-reviewed study that treats calipers as a classification shortcut in thyroid or ovarian ultrasound, although calipers are placed on the nodule border by design. In chest radiography, pneumothorax classifiers respond to the chest tube that treats the condition. Rueckel et al. (2021) reduced this confounding by training with in-image annotations, and Jiménez-Sánchez et al. (2023) labelled drains on a public dataset to expose the shortcut. A drain projects over the lung field, so a lung mask cannot remove it. For capsule endoscopy we found no study of bubbles or debris as a shortcut, even though debris and lesions often occupy the same pixels.

### 4 Does restricting the model to the ROI help?

In natural images, classifiers can recognise objects from the background alone, and more accurate models tend to rely on it less (Xiao et al., 2021). The same holds for medical images. In a careful masking study on chest radiographs and fundus photographs, models trained on full images performed better on images with the ROI removed than on images containing only the ROI (Sourget et al., 2025). Training on segmented dermoscopy images improved balanced accuracy for classifiers trained on HAM10000 but not for those trained on ISIC (Maron et al., 2021), a mixed result that is easy to overlook. ISNet takes a softer route and penalises relevance outside the segmentation mask during training, which made chest-radiograph classifiers more robust to

background bias (Bassi et al., 2024).

All of these studies assume that the shortcut lies in the background. The one result that points the other way comes from natural images: when several shortcuts are present, removing one can make the model rely more on the others (Li et al., 2023). ROI masking is a mitigation of exactly this kind. If an artifact remains inside the ROI, the same reasoning predicts that masking will increase reliance on it. To our knowledge this has not been tested in medical images, and it is the starting point of the thesis.

## 5 Detecting shortcuts

Explanation methods are the most common tool. Spectral relevance analysis clusters heatmaps to reveal Clever-Hans strategies (Anders et al., 2022), and Reveal-to-Revise turns detection, correction and re-evaluation into a cycle (Pahde et al., 2023). The limits are well documented: saliency maps localise chest pathology worse than radiologists (Saporta et al., 2022), and their ability to reveal a known confounder varies from method to method (Sun et al., 2023). In our view there is a further limit for in-ROI artifacts: when calipers sit on a nodule, a heatmap that highlights the nodule cannot tell whether the model looks at the tissue or at the calipers.

Counterfactual methods avoid this by changing the image. Diffusion-based counterfactuals add or remove a suspected shortcut and measure the change in prediction (Weng et al., 2024), and RadEdit edits chest radiographs under masks to stress-test classifiers (Pérez-García et al., 2024). The thesis uses the same idea in its simplest form, by pasting one real artifact inside or outside the ROI of the same images.

## 6 Removing shortcuts

With group labels, retraining only the last layer on balanced data recovers much of the robustness lost to spurious correlations (Kirichenko et al., 2023), because standard training already learns the core features even when the classifier does not use them (Izmailov et al., 2022). Without group labels, the errors of a first model can be up-weighted (Liu et al., 2021), or the last layer reweighted automatically (Qiu et al., 2023); variance-of-risk penalties are another option (Krueger et al., 2021). The benchmarks are sobering. Across 20 algorithms and 12 datasets, methods improved subgroup robustness only for some kinds of shift (Yang et al., 2023), and in medical imaging, bias-mitigation methods did not improve fairness significantly over standard training, while the choice of model selection mattered more (Zong et al., 2023).

A more principled line starts from causal reasoning. Counterfactual invariance requires that the prediction does not change when only the nuisance changes; when the label causes the image, as in diagnosis, this implies independence of the prediction and the nuisance given the label (Veitch et al., 2021). Makar et al. (2022) enforce such conditional independences with auxiliary labels available at training time, and NuRD trains under a distribution in which the nuisance-label relation is broken (Puli et al., 2022). Kaur et al. (2023) show that the right constraint depends on how the data were generated, so no single constraint suits every shift. Linear concept erasure removes a concept from the features in closed form (Belrose et al., 2023), at the cost of also removing label signal when the two are correlated; SPLINCE was designed to keep that signal (Holstege et al., 2025). Generative data are a different route: diffusion-generated images improved fairness out of distribution in histopathology, chest radiography and dermatology (Ktena et al., 2024).

The remedy the thesis recommends, a class-conditional mean constraint on the masked features, is a first-moment, linear case of the conditional independence of Veitch et al. (2021) and Makar et al. (2022). What the thesis adds is the reason for applying it after masking, and a

registered test of whether any remedy beats masking across all cohorts. Given the benchmark results above, it is not surprising that none does.

## 7 Why models prefer shortcuts

Theory explains why a simple feature wins even when a reliable one is available. Nagarajan et al. (2021) show that geometric and statistical skews push standard training toward spurious features even when the invariant feature predicts the label perfectly. Gradient descent with cross-entropy can starve less salient features (Pezeshki et al., 2021), and its implicit preference for large margins produces shortcut reliance even when the shortcut adds no information (Puli et al., 2023). Hermann et al. (2024) separate a feature’s *predictivity* from its *availability*, that is, how easily it can be extracted, and measure shortcut bias as reliance on the more available but less predictive feature. For frozen pretrained features, as in most of the thesis, linear probing can generalise better out of distribution than full fine-tuning (Kumar et al., 2022). The thesis’s linear-Gaussian account fits this picture: masking removes competing evidence outside the ROI and so changes the relative predictivity of what is left. The difference is that the account was used to make quantitative predictions that were registered before the results existed.

## 8 How such claims should be evaluated

Varoquaux and Cheplygina (2022) review methodological failures in medical-imaging machine learning, from biased data to publication incentives, and Kapoor and Narayanan (2023) found data leakage, including non-independent splits, in at least 294 papers across 17 fields. Results also vary with data sampling, initialisation and hyperparameters, and accounting for several of these sources at once gives better estimates of that variation (Bouthillier et al., 2021). In MICCAI papers, confidence intervals are rarely reported, and their typical width is larger than the gap between the first- and second-ranked methods (Christodoulou et al., 2024). Overall metrics also hide failures in clinically important subgroups (Yang et al., 2024), a problem first described as hidden stratification (Oakden-Rayner et al., 2020). These papers shaped the evaluation in the thesis: splits grouped by patient or lesion, a bootstrap that resamples training seeds and test images together, registration before results, and evaluation on the cases that contradict the shortcut as well as on the whole test set.

## 9 Where the thesis stands

The literature establishes four things firmly. Medical classifiers take shortcuts, including ones produced by clinical practice; the artifacts involved are well documented in dermoscopy and chest radiography; image backgrounds carry signal, and restricting a model to the ROI sometimes helps; and no mitigation method is reliably better than standard training across settings. Table 1 lists what remained open and how the thesis addresses it.

Some of this is not new, and the thesis should say so. That a mask cannot remove what it keeps is self-evident; what was missing is evidence that masking makes reliance on the kept artifact grow, how large the effect is on real artifacts, and how to predict it. The recommended constraint is a special case of known conditional-invariance methods, and the group-robustness methods are used as baselines rather than proposed. The closest conceptual precedent is the whac-a-mole result of Li et al. (2023); the thesis supplies the location-specific mechanism and the medical evidence.

Table 1: Open questions in the reviewed work and the corresponding contribution of the thesis.

| Open question                                                                  | Closest prior work                                                                                                 | Thesis                                                                                                                                      |
|--------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Does the artifact’s location relative to the ROI decide whether masking works? | Sourget et al. (2025); Maron et al. (2021); Bassi et al. (2024): masking studies that assume a background shortcut | Location varied directly: controlled sweeps, traps built from real artifacts, and a transplant of the same artifact across the ROI boundary |
| Can masking increase reliance on an artifact?                                  | Li et al. (2023): removing one shortcut amplifies others (natural images)                                          | Reliance on an in-ROI artifact grows after masking, in four cohorts and three modalities                                                    |
| Can the harm be predicted?                                                     | Hermann et al. (2024); Puli et al. (2023): explanatory theories                                                    | A two-parameter account with predictions registered before the results                                                                      |
| What should be used instead of masking?                                        | Veitch et al. (2021); Makar et al. (2022); Zong et al. (2023)                                                      | A registered test against masking; a class-conditional constraint comes closest, none dominates                                             |
| Calipers and debris as classification shortcuts                                | Lin et al. (2024) (segmentation only)                                                                              | Thyroid and ovarian ultrasound and capsule endoscopy alongside dermoscopy                                                                   |

## Use of AI tools

This review was prepared with the assistance of an AI system (Claude, Anthropic), which was used to search the literature, check bibliographic details against publisher and proceedings pages, and draft the text. The choice of scope, the interpretation of the studies and any remaining errors are the author’s responsibility.

## References

- Christopher J. Anders, Leander Weber, David Neumann, Wojciech Samek, Klaus-Robert Müller, and Sebastian Lapuschkin. Finding and removing Clever Hans: using explanation methods to debug and improve deep models. *Information Fusion*, 77:261–295, 2022. doi: 10.1016/j.inffus.2021.07.015.
- Pedro R. A. S. Bassi, Sergio S. J. Dertkigil, and Andrea Cavalli. Improving deep neural network generalization and robustness to background bias via layer-wise relevance propagation optimization. *Nature Communications*, 15:291, 2024. doi: 10.1038/s41467-023-44371-z.
- Nora Belrose, David Schneider-Joseph, Shauli Ravfogel, Ryan Cotterell, Edward Raff, and Stella Biderman. LEACE: perfect linear concept erasure in closed form. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, 2023.
- Xavier Bouthillier, Pierre Delaunay, Mirko Bronzi, et al. Accounting for variance in machine learning benchmarks. In *Proceedings of Machine Learning and Systems (MLSys)*, volume 3, 2021.
- Alexander Brown, Nenad Tomasev, Jan Freyberg, Yuan Liu, Alan Karthikesalingam, and Jessica Schrouff. Detecting shortcut learning for fair medical AI using shortcut testing. *Nature Communications*, 14:4314, 2023. doi: 10.1038/s41467-023-39902-7.
- Bill Cassidy, Connah Kendrick, Andrzej Brodzicki, Joanna Jaworek-Korjakowska, and Moi Hoon Yap. Analysis of the ISIC image datasets: usage, benchmarks and recommendations. *Medical Image Analysis*, 75:102305, 2022. doi: 10.1016/j.media.2021.102305.

- Evangelia Christodoulou, Annika Reinke, Rola Houhou, et al. Confidence intervals uncovered: are we ready for real-world medical imaging AI? In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, LNCS, pages 124–132, 2024. doi: 10.1007/978-3-031-72117-5\_12.
- Alex J. DeGrave, Joseph D. Janizek, and Su-In Lee. AI for radiographic COVID-19 detection selects shortcuts over signal. *Nature Machine Intelligence*, 3:610–619, 2021. doi: 10.1038/s42256-021-00338-7.
- Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2:665–673, 2020. doi: 10.1038/s42256-020-00257-z.
- Judy Wawira Gichoya, Imon Banerjee, Ananth Reddy Bhimireddy, et al. AI recognition of patient race in medical imaging: a modelling study. *The Lancet Digital Health*, 4(6):e406–e414, 2022. doi: 10.1016/S2589-7500(22)00063-2.
- Ben Glocker, Charles Jones, Mélanie Bernhardt, and Stefan Winzeck. Algorithmic encoding of protected characteristics in chest X-ray disease detection models. *eBioMedicine*, 89:104467, 2023. doi: 10.1016/j.ebiom.2023.104467.
- Katherine L. Hermann, Hossein Mobahi, Thomas Fel, and Michael C. Mozer. On the foundations of shortcut learning. In *International Conference on Learning Representations (ICLR)*, 2024.
- Floris Holstege, Shauli Ravfogel, and Bram Wouters. Preserving task-relevant information under linear concept removal. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- Pavel Izmailov, Polina Kirichenko, Nate Gruver, and Andrew Gordon Wilson. On feature learning in the presence of spurious correlations. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 35, 2022.
- Amelia Jiménez-Sánchez, Dovile Juodelyte, Bethany Chamberlain, and Veronika Cheplygina. Detecting shortcuts in medical images – a case study in chest X-rays. In *IEEE International Symposium on Biomedical Imaging (ISBI)*, 2023.
- Charles Jones, Daniel C. Castro, Fabio De Sousa Ribeiro, Ozan Oktay, Melissa McCradden, and Ben Glocker. A causal perspective on dataset bias in machine learning for medical imaging. *Nature Machine Intelligence*, 6:138–146, 2024. doi: 10.1038/s42256-024-00797-8.
- Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in machine-learning-based science. *Patterns*, 4(9):100804, 2023. doi: 10.1016/j.patter.2023.100804.
- Jivat Neet Kaur, Emre Kiciman, and Amit Sharma. Modeling the data-generating process is necessary for out-of-distribution generalization. In *International Conference on Learning Representations (ICLR)*, 2023.
- Polina Kirichenko, Pavel Izmailov, and Andrew Gordon Wilson. Last layer re-training is sufficient for robustness to spurious correlations. In *International Conference on Learning Representations (ICLR)*, 2023.
- David Krueger, Ethan Caballero, Joern-Henrik Jacobsen, Amy Zhang, Jonathan Binas, Dinghuai Zhang, Remi Le Priol, and Aaron Courville. Out-of-distribution generalization via risk extrapolation (REx). In *International Conference on Machine Learning (ICML)*, volume 139 of *PMLR*, pages 5815–5826, 2021.

- Ira Ktena, Olivia Wiles, Isabela Albuquerque, et al. Generative models improve fairness of medical classifiers under distribution shifts. *Nature Medicine*, 30(4):1166–1173, 2024. doi: 10.1038/s41591-024-02838-6.
- Ananya Kumar, Aditi Raghunathan, Robbie Jones, Tengyu Ma, and Percy Liang. Fine-tuning can distort pretrained features and underperform out-of-distribution. In *International Conference on Learning Representations (ICLR)*, 2022.
- Zhiheng Li, Ivan Evtimov, Albert Gordo, Caner Hazirbas, Tal Hassner, Cristian Canton Ferrer, Chenliang Xu, and Mark Ibrahim. A Whac-A-Mole dilemma: shortcuts come in multiples where mitigating one amplifies others. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 20071–20082, 2023.
- Manxi Lin, Nina Weng, Kamil Mikolaj, et al. Shortcut learning in medical image segmentation. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, volume 15008 of *LNCS*, pages 623–633, 2024. doi: 10.1007/978-3-031-72111-3\_59.
- Evan Z. Liu, Behzad Haghighi, Annie S. Chen, Aditi Raghunathan, Pang Wei Koh, Shiori Sagawa, Percy Liang, and Chelsea Finn. Just train twice: improving group robustness without training group information. In *International Conference on Machine Learning (ICML)*, volume 139 of *PMLR*, pages 6781–6792, 2021.
- Gianluca Maguolo and Loris Nanni. A critic evaluation of methods for COVID-19 automatic detection from X-ray images. *Information Fusion*, 76:1–7, 2021. doi: 10.1016/j.inffus.2021.04.008.
- Maggie Makar, Ben Packer, Dan Moldovan, Davis Blalock, Yoni Halpern, and Alexander D’Amour. Causally motivated shortcut removal using auxiliary labels. In *International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 151 of *PMLR*, pages 739–766, 2022.
- Roman C. Maron, Achim Hekler, Eva Krieghoff-Henning, Max Schmitt, Justin G. Schlager, Jochen S. Utikal, and Titus J. Brinker. Reducing the impact of confounding factors on skin cancer classification via image segmentation: technical model study. *Journal of Medical Internet Research*, 23(3):e21695, 2021. doi: 10.2196/21695.
- Vaishnavh Nagarajan, Anders Andreassen, and Behnam Neyshabur. Understanding the failure modes of out-of-distribution generalization. In *International Conference on Learning Representations (ICLR)*, 2021.
- Luke Oakden-Rayner, Jared Dunnmon, Gustavo Carneiro, and Christopher Ré. Hidden stratification causes clinically meaningful failures in machine learning for medical imaging. In *Proceedings of the ACM Conference on Health, Inference, and Learning (CHIL)*, pages 151–159, 2020.
- Frederik Pahde, Maximilian Dreyer, Wojciech Samek, and Sebastian Lapuschkin. Reveal to revise: an explainable AI life cycle for iterative bias correction of deep models. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, volume 14221 of *LNCS*, 2023. doi: 10.1007/978-3-031-43895-0\_56.
- Fernando Pérez-García, Sam Bond-Taylor, Pedro P. Sanchez, et al. RadEdit: stress-testing biomedical vision models via diffusion image editing. In *European Conference on Computer Vision (ECCV)*, LNCS, pages 358–376, 2024. doi: 10.1007/978-3-031-73254-6\_21.
- Samuel William Pewton, Bill Cassidy, Connah Kendrick, and Moi Hoon Yap. Dermoscopic dark corner artifacts removal: friend or foe? *Computer Methods and Programs in Biomedicine*, 244: 107986, 2024. doi: 10.1016/j.cmpb.2023.107986.

- Mohammad Pezeshki, Sékou-Oumar Kaba, Yoshua Bengio, Aaron Courville, Doina Precup, and Guillaume Lajoie. Gradient starvation: a learning proclivity in neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, 2021.
- Aahlad Manas Puli, Lily H. Zhang, Eric Karl Oermann, and Rajesh Ranganath. Out-of-distribution generalization in the presence of nuisance-induced spurious correlations. In *International Conference on Learning Representations (ICLR)*, 2022.
- Aahlad Manas Puli, Lily Zhang, Yoav Wald, and Rajesh Ranganath. Don’t blame dataset shift! Shortcut learning due to gradients and cross entropy. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 36, 2023.
- Shikai Qiu, Andres Potapczynski, Pavel Izmailov, and Andrew Gordon Wilson. Simple and fast group robustness by automatic feature reweighting. In *International Conference on Machine Learning (ICML)*, volume 202 of *PMLR*, pages 28448–28467, 2023.
- Michael Roberts, Derek Driggs, Matthew Thorpe, et al. Common pitfalls and recommendations for using machine learning to detect and prognosticate for COVID-19 using chest radiographs and CT scans. *Nature Machine Intelligence*, 3:199–217, 2021. doi: 10.1038/s42256-021-00307-0.
- Johannes Rueckel, Christian Huemmer, Andreas Fieselmann, et al. Pneumothorax detection in chest radiographs: optimizing artificial intelligence system for accuracy and confounding bias reduction using in-image annotations in algorithm training. *European Radiology*, 31:7888–7900, 2021. doi: 10.1007/s00330-021-07833-w.
- Adriel Saporta, Xiaotong Gui, Ashwin Agrawal, et al. Benchmarking saliency methods for chest X-ray interpretation. *Nature Machine Intelligence*, 4:867–878, 2022. doi: 10.1038/s42256-022-00536-x.
- Laleh Seyyed-Kalantari, Haoran Zhang, Matthew B. A. McDermott, Irene Y. Chen, and Marzyeh Ghassemi. Underdiagnosis bias of artificial intelligence algorithms applied to chest radiographs in under-served patient populations. *Nature Medicine*, 27:2176–2182, 2021. doi: 10.1038/s41591-021-01595-0.
- Théo Sourget, Michelle Hestbek-Møller, Amelia Jiménez-Sánchez, Jack Junchi Xu, and Veronika Cheplygina. Mask of truth: model sensitivity to unexpected regions of medical images. *Journal of Imaging Informatics in Medicine*, 2025. doi: 10.1007/s10278-025-01531-5.
- Emma A. M. Stanley, Raissa Souza, Anthony Winder, Vedant Gulve, Kimberly Amador, Matthias Wilms, and Nils D. Forkert. Towards objective and systematic evaluation of bias in artificial intelligence for medical imaging. *Journal of the American Medical Informatics Association*, 31(11):2613–2621, 2024. doi: 10.1093/jamia/ocae165.
- Susu Sun, Lisa M. Koch, and Christian F. Baumgartner. Right for the wrong reason: can interpretable ML techniques detect spurious correlations? In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*, volume 14221 of *LNCS*, pages 425–434, 2023. doi: 10.1007/978-3-031-43895-0\_40.
- Gaël Varoquaux and Veronika Cheplygina. Machine learning for medical imaging: methodological failures and recommendations for the future. *npj Digital Medicine*, 5:48, 2022. doi: 10.1038/s41746-022-00592-y.
- Victor Veitch, Alexander D’Amour, Steve Yadlowsky, and Jacob Eisenstein. Counterfactual invariance to spurious correlations: why and how to pass stress tests. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 34, 2021.

- Z. Wang, Y. Xue, H. Xi, X. Tan, K. Lin, C. Wang, and J. Zhang. Influence of hair presence on dermoscopic image analysis by AI in skin lesion diagnosis. *Computers in Biology and Medicine*, 183:109335, 2024. doi: 10.1016/j.compbiomed.2024.109335.
- Nina Weng, Paraskevas Pegios, Eike Petersen, Aasa Feragen, and Siavash Bigdeli. Fast diffusion-based counterfactuals for shortcut removal and generation. In *European Conference on Computer Vision (ECCV)*, LNCS, pages 338–357, 2024. doi: 10.1007/978-3-031-73016-0\_20.
- Julia K. Winkler, Christine Fink, Ferdinand Toberer, et al. Association between surgical skin markings in dermoscopic images and diagnostic performance of a deep learning convolutional neural network for melanoma recognition. *JAMA Dermatology*, 155(10):1135–1141, 2019.
- Julia K. Winkler, Katharina Sies, Christine Fink, et al. Association between different scale bars in dermoscopic images and diagnostic performance of a market-approved deep learning convolutional neural network for melanoma recognition. *European Journal of Cancer*, 145: 146–154, 2021.
- Kai Xiao, Logan Engstrom, Andrew Ilyas, and Aleksander Mądry. Noise or signal: the role of image backgrounds in object recognition. In *International Conference on Learning Representations (ICLR)*, 2021.
- Yuzhe Yang, Haoran Zhang, Dina Katabi, and Marzyeh Ghassemi. Change is hard: a closer look at subpopulation shift. In *International Conference on Machine Learning (ICML)*, volume 202 of *PMLR*, pages 39584–39622, 2023.
- Yuzhe Yang, Haoran Zhang, Judy W. Gichoya, Dina Katabi, and Marzyeh Ghassemi. The limits of fair medical imaging AI in real-world generalization. *Nature Medicine*, 30:2838–2848, 2024. doi: 10.1038/s41591-024-03113-4.
- Yongshuo Zong, Yongxin Yang, and Timothy Hospedales. MEDFAIR: benchmarking fairness for medical imaging. In *International Conference on Learning Representations (ICLR)*, 2023.