

# Masking cannot remove what is inside the region of interest: location-dependent shortcut mitigation across dermoscopy, ultrasound and capsule endoscopy

---

## Abstract

Medical image classifiers exploit acquisition artifacts—hair, sonographer calipers, intestinal debris—that correlate with the diagnosis. The standard defence restricts the model to the region of interest (ROI). We show that this defence fails for artifacts that lie *inside* the ROI, give a model that predicts when and why, and evaluate what to do instead. In pre-registered, leakage-safe traps built from real artifacts in four diseases (dermoscopic hair, thyroid and ovarian calipers, capsule-endoscopy debris), the benefit of ROI masking is always larger for out-of-ROI than for in-ROI artifacts (crossover +0.15 to +0.37 reversed-test AUROC; all four survive Holm correction). The loss becomes harm when masking also removes disease context: in dermoscopy, in controlled overlap sweeps (ovary  $-0.20$ , capsule  $-0.10$ ), with end-to-end fine-tuning (capsule  $-0.28$ ), and on the unaltered thyroid test distribution, where masking lowers AUROC on shortcut-conflicting cases by 0.10 [0.07, 0.14]. A linear-Gaussian analysis reproduces these signs and orderings in closed form. Among remedies, a mask-then-erase procedure driven by generic synthetic overlays (U-MtE) needs no artifact example, mask or label, removes in-ROI shortcuts where the artifact is a distinct overlay (thyroid +0.13 to +0.22 over masking on three backbones; +0.14 end-to-end), and beats using the same overlays as augmentation in three of four cohorts; it fails when the artifact resembles the pathology unless disease-protected, and cannot help when the shortcut is carried by correlates of the artifact (hair, chest drains), where label-based reweighting is required. A validation-driven selector over these options, using image-level artifact labels on validation data only, beats masking in all five settings and matches the best candidate in each. Code, splits, pre-registrations and a backbone-agnostic tool are public.

*Keywords:* shortcut learning, spurious correlation, artifacts, dermoscopy, chest radiography, concept erasure

---

## 1. Introduction

Deep classifiers for medical images frequently rely on features that are correlated with the label in the training distribution but are not evidence of disease [9, 7]. In dermoscopy, surgical skin markings [24] and ruler scale bars [25] changed the melanoma predictions of a market-approved system; in chest radiography, pneumothorax classifiers detect the chest drain that treats the condition rather than the condition itself

[16, 10]. Such correlations are causal in origin—suspicious lesions are measured and marked, collapsed lungs are drained—so they are present in every retrospectively collected dataset.

The most widely recommended defence is to restrict the classifier to the region of interest (ROI): segment the lesion or the lungs and discard the rest of the image. The rationale assumes that the shortcut lives in the background. Rulers, hair, ink and drains do not respect that assumption. When the artifact overlaps the ROI, masking cannot remove it; it removes everything *else*. We ask whether the effect of a mitigation depends on where the artifact sits, which quantity governs that dependence, what supervision each family of mitigations actually needs, and whether the in-ROI regime can be mitigated without localising the artifact.

### *Contributions.*

1. **A location law for ROI masking**, established with pre-registered real-artifact traps in four diseases and three imaging modalities, replicated on four frozen backbones and with end-to-end fine-tuning, confirmed on an unaltered test distribution, and robust to the artifact-label definition (Secs. 7–8).
2. **A closed-form account** (Sec. 6) that predicts when in-ROI masking harms, why erasure fails for pathology-like artifacts, and why group reweighting is location-invariant.
3. **An annotation-free remedy (U-MtE)** and a disease-protected variant, with the evidence that erasure, not augmentation, is what works, and a delineation of the regimes in which it does not.
4. **A practical decision procedure**: a validation-driven selector and a decision rule based on where the artifact sits and what carries the shortcut, benchmarked against DFR, GroupDRO, JTT, LEACE and SPLICE with gaming-aware reporting.
5. **Reproducibility**: pre-registration before every experiment, Holm-corrected primary claims, replication of the original pilot (73/88 numbers matched), public code and a two-command tool.

## **2. Related work**

*Artifacts in dermoscopy.* Bissoto et al. [4] annotated seven artifact types in ISIC images and showed that debiasing is harder than it looks; Nauta et al. [15] neutralised coloured calibration patches by inpainting and retraining—an artifact that a colour rule localises and that lies largely outside the lesion. MaskMedPaint [11] inpaints regions *outside* the classification-relevant area. None of these treats artifact location relative to the ROI as a variable.

*Chest radiography.* Drain-driven pneumothorax shortcuts are documented on NIH-CXR14 and CheXpert [16, 10]; recent work localises them to late layers of probed encoders [17] and mitigates them by prevalence-equalised recalibration [12], which requires the artifact label at test time.

*Group robustness and erasure.* Reweighting and GroupDRO [19], last-layer retraining [13] and linear concept erasure [18, 2] require artifact labels or paired data; erasure from unpaired labels removes label-correlated signal when the concept is correlated with the label, which is why we build pairs by insertion. Task-preserving linear erasure (SPLICE; 1) and projection of biased directions in vision–language models [5] are the closest erasure methods; both require concept labels or prompts, whereas our concept subspace is estimated from synthetic insertions. Random erasing and occlusion augmentation [28, 8] use similar overlays as training noise; we show that this does not remove the shortcut. Test-time selection [3] steers dermoscopy models with user keypoints.

### 3. Data

*Controlled dermoscopy cohort.* ISIC 2018 Task 1/2 images with manual lesion masks and melanoma labels from the union of the ISIC 2016 and 2017 tables (three images with conflicting labels dropped); perceptual-hash (Hamming  $\leq 8$ ) and lesion grouping; 2,437 images for which a  $65 \times 19$  px ruler (0.46% of a  $518 \times 518$  image) can be placed at 0/25/50/75/100% overlap with the lesion within 10 percentage points (train/val/test 1,849/371/371 before common support).

*Real hair cohort (ISIC 2019).* All 25,331 ISIC 2019 training images; hair/ruler masks from the archive of Wegley et al. [23]; lesion masks manual for HAM10000 [21] and predicted elsewhere by a U-Net trained on ISIC 2018 Task 1. Two traps share one hair-free group ( $\leq 30$  hair pixels at native resolution): Trap A contains hair with overlap  $r = |\text{hair} \cap \text{lesion}|/|\text{hair}| \geq 0.5$ , Trap B  $r < 0.1$ . Cells are subsampled to the hair-free source mix (HAM/BCN/MSK) within each label, giving Trap A 704/157/2,738/989 and Trap B 704/157/4,327/930 images (hair-free benign/melanoma, hair benign/melanoma); five seeds of five-fold group-safe cross-validation (leakage groups: lesion identifiers and pHash); 172 reversed-test melanomas per trap and seed.

*Thyroid ultrasound.* TN3K images with TNCD benign/malignant labels (3,493 images; 1,210 malignant) and manual nodule masks. Sonographer calipers were segmented by a frozen rule-based detector (dotted measurement lines and ‘+’ ends; audit in Sec. 11). Trap A: calipers with  $r \geq 0.5$  (1,150 images); Trap B:  $r < 0.1$  (326); caliper-free: 1,801. No patient identifiers are published; leakage groups are pHash chains.

*Ovarian tumour ultrasound.* MMOTU OTU\_2d [27] (CC BY 4.0): 1,202 images with tumour masks, labelled benign cystic (chocolate cyst, serous cystadenoma, simple cyst) versus lesions with solid components (teratoma, theca-cell tumour, mucinous cystadenoma, high-grade serous carcinoma). The thyroid detector is applied unchanged: Trap A 341/276, Trap B 91/103, caliper-free 170/178 (benign-cystic/solid).

*Capsule endoscopy.* SEE-AI frames [26] (CC BY 4.0) containing only erosions (3,550) or only polyp-like lesions (1,931); the ROI is the union of the expert lesion boxes. Contamination (debris, bubbles, bile) is segmented by a linear probe on DINOv2 patch tokens trained on 2,160 expert-masked frames (held-out pixel accuracy 0.92, IoU 0.62; expert masks are used where available). Debris-free:  $< 3\%$  of the frame; Trap A: debris  $\geq 10\%$  with  $r \geq 0.5$ ; Trap B:  $r < 0.1$ . Leakage groups: blocks of 100 consecutive frame numbers plus pHash duplicates.

*Chest radiography.* NIH ChestX-ray14 [22]; lung ROIs from the TorchXRayVision PSPNet [6]. *Controlled cohort:* 1,906 pneumothorax images without a chest drain (NEATX annotations) and 5,718 *No Finding* images from other patients; a  $100 \times 12$  px radio-opaque tube is placed at controlled overlap with the lungs (7,610 images in common support). *Real drains* (in-ROI only): 29,687 images; drains are labelled by NEATX for pneumothorax-positive images and by a drain detector (precision 0.95) for negatives; environments are built within view position (AP/PA); five patient-grouped folds.

#### 4. Problem setup and protocol

Let  $Y \in \{0, 1\}$  be the diagnosis and  $A \in \{0, 1\}$  the presence of an artifact. A *trap* trains with  $P(A=1 | Y=1) = 0.9$  and  $P(A=1 | Y=0) = 0.1$  and tests on a correlated environment with the same rates, a *reversed* environment with the rates swapped (0.1/0.9), and a clean environment without artifacts. A model that uses  $A$  gains on the correlated test and collapses on the reversed one; the primary metric is reversed-test AUROC, and clean AUROC measures the cost of a mitigation. The overlap  $r = |\text{artifact} \cap \text{ROI}| / |\text{artifact}|$  is the controlled variable. All backbones are frozen and only linear heads are fitted (an end-to-end fine-tuning check is reported in Sec. 10.2); regularisation strength and the decision threshold (balanced accuracy) are selected on clean validation data only. Uncertainty is a hierarchical paired bootstrap that resamples training seeds (or cross-validation folds), then paired test images within each, computes the AUROC difference inside each cluster and averages; predictions of different models are never pooled into one ROC curve. Splits are grouped by lesion/patient identifier and perceptual hash.

#### 5. Mask-then-erase with a generic insertion library

*Insertion pairs.* Removing an in-ROI artifact requires its pixels; inserting one does not. Given a frozen encoder  $f$  and an overlay generator  $\mathcal{G}$ , every training image  $x_i$  yields a counterfactual  $x_i^+ = \mathcal{G}(x_i; \xi_i)$  with random parameters  $\xi_i$  and random position, half of the elements centred inside the ROI. Because insertion ignores the label, the differences  $\Delta_i = f(x_i^+) - f(x_i)$  are independent of  $Y$ .

*Generic library.*  $\mathcal{G}$  is one procedural family used unchanged for every modality and artifact: thin curvilinear strokes, thick bands, solid or translucent patches, blobs, dashed lines and small glyphs, with random colour (half greyscale), opacity and width. It contains no ruler, hair, caliper or debris template; it is a prior over “foreign overlays”, not a model of any particular artifact.

*Erasure.* I2E erases the span of the differences:

$$\Sigma_{\Delta} = \frac{1}{n} \sum_i \Delta_i \Delta_i^{\top}, \quad U_k = \text{top-}k \text{ eigvecs}(\Sigma_{\Delta}), \quad P(z) = z - U_k U_k^{\top} (z - \mu), \quad (1)$$

with  $k$  chosen *without labels* as the smallest rank whose held-out paired-difference energy  $\mathbb{E}\|P\Delta\|^2/\mathbb{E}\|\Delta\|^2$  falls below 0.10 (fixed before any experiment).

*Mask-then-erase (U-MtE).* Masking and erasure have complementary reach: masking removes any artifact outside the ROI but none inside; erasure acts on the representation regardless of location. U-MtE therefore computes the pairs *on masked images* ( $\Delta_i = f(m(x_i^+)) - f(m(x_i))$ ) so that the erased subspace targets exactly what masking leaves behind, and fits the head on  $P(f(m(x)))$ . When image-level artifact labels exist, the head is group-balanced (U-MtE<sub>bal</sub>).

*Disease protection.* If the artifact resembles the pathology (intestinal debris versus erosion fibrin),  $U_k$  can contain disease evidence. Let  $W$  be an orthonormal basis of logistic label directions fitted on artifact-free training images; we erase  $U'_k = \text{orth}((I - WW^{\top})U_k)$  instead, which guarantees  $w^{\top}P(z) = w^{\top}z$  for every  $w \in \text{span}(W)$  (cf. task-preserving erasure, 1).

*Tool.* The released package exposes insertion, erasure and a patch-token artifact probe for any frozen backbone (ViTs and CNNs) through a two-call API.

## 6. Why location matters: a linear-Gaussian account

Let a frozen encoder yield a disease block with class-separation signal-to-noise ratio  $S$  and an artifact block with SNR  $A$  per unit of artifact presence; training uses  $P(A=1 \mid Y=1) = p_1$ ,  $P(A=1 \mid Y=0) = p_0$  and the reversed test swaps them. For the LDA-optimal linear head the artifact’s effective training SNR is  $\tilde{A} = (p_1 - p_0)^2 A / (1 + p_1(1 - p_1)A)$  and

$$\text{AUROC}_{\text{rev}}(S, \tilde{A}) = \Phi\left(\frac{S - \tilde{A}}{\sqrt{2(S + \tilde{A})}}\right), \quad (2)$$

increasing in  $S$  and decreasing in  $\tilde{A}$  (verified by simulation to within 0.009 AUROC; Fig. 1). Masking maps  $(S, \tilde{A}) \mapsto (S_m, \tilde{A}_m)$ . Out of the ROI,  $\tilde{A}_m = 0$ ; inside,  $\tilde{A}_m \geq \tilde{A}$ . Hence (i) in-ROI masking harms exactly when the mask removes disease-informative context ( $S_m < S$ ) and helps only by denoising ( $S_m > S$ ); (ii)




Figure 1: Reversed-test AUROC predicted by the linear-Gaussian model (lines) and simulated with logistic heads (points) as masking removes a growing share of disease signal ( $S = 1$ ,  $A = 2$ , 90/10 training). Generated by `scripts/analysis/theory_sim.py`.

for any  $S_m$  the benefit of masking is larger out of the ROI than inside it—the crossover is always positive; (iii) erasing a subspace that contains the artifact direction sets  $\tilde{A} \rightarrow 0$  but scales  $S$  by  $1 - \rho^2$ , where  $\rho$  is its overlap with the disease direction, so erasure beats in-ROI masking unless the artifact resembles the pathology, which disease protection prevents; (iv) group balancing removes the artifact–label association and everything correlated with it, giving  $\Phi(\sqrt{S/2})$  irrespective of location; (v) a shortcut carried by correlates that neither the mask nor pixel-level erasure touch can only be removed by label-based reweighting. Each prediction is borne out in Sections 7–9.

## 7. Results: dermoscopy

### 7.1. Masking reverses with location (controlled)

With a synthetic ruler whose overlap with the lesion is swept from 0 to 100% (2,437 ISIC 2018 images, DINOv2 ViT-B/14 at 518 px, three seeds), ROI masking improves reversed-environment AUROC over ERM by  $+0.184$  [ $+0.146$ ,  $+0.224$ ] when the ruler lies outside the lesion and *lowers* it by  $-0.097$  [ $-0.138$ ,  $-0.056$ ],  $-0.112$  [ $-0.155$ ,  $-0.068$ ] and  $-0.129$  [ $-0.170$ ,  $-0.087$ ] at 50, 75 and 100% overlap (location interaction  $+0.313$  [ $+0.272$ ,  $+0.353$ ]; Table 1). The sign flip is reproduced at 224 px ( $+0.361$  at 0%;  $-0.095$  at 100%) and under randomised ruler colour, opacity, ticks and scale. At 0% overlap the masked model’s correlated and reversed AUROC coincide (0.849/0.849): the shortcut is removed. At 100% the masked model’s correlated AUROC exceeds ERM’s (0.961 vs 0.937): it uses the ruler *more*.

Table 1: Synthetic ruler, ISIC 2018, DINOv2 ViT-B/14 @518: AUROC and reversed-test  $\Delta$  vs ERM [95% CI]

| Method                         | Needs                        | clean@0 | rev@0 | clean@50 | rev@50 | clean@100 | rev@100 | $\Delta$ rev@0          | $\Delta$ rev@50         | $\Delta$ rev@100        |
|--------------------------------|------------------------------|---------|-------|----------|--------|-----------|---------|-------------------------|-------------------------|-------------------------|
| ERM                            | —                            | 0.797   | 0.665 | 0.793    | 0.573  | 0.792     | 0.515   |                         |                         |                         |
| ROI masking                    | ROI mask (train+test)        | 0.849   | 0.849 | 0.829    | 0.476  | 0.825     | 0.386   | +0.184 [+0.146, +0.224] | -0.097 [-0.138, -0.056] | -0.129 [-0.170, -0.087] |
| Inpainting (oracle)            | artifact pixels (train+test) | 0.810   | 0.806 | 0.809    | 0.783  | 0.807     | 0.759   | +0.141 [+0.120, +0.162] | +0.211 [+0.187, +0.234] | +0.243 [+0.218, +0.269] |
| Group-balanced                 | image-level A (train)        | 0.814   | 0.802 | 0.815    | 0.792  | 0.806     | 0.796   | +0.137 [+0.109, +0.168] | +0.219 [+0.188, +0.252] | +0.281 [+0.248, +0.315] |
| DFR                            | image-level A (val)          | 0.786   | 0.789 | 0.787    | 0.794  | 0.789     | 0.794   | +0.124 [+0.091, +0.160] | +0.221 [+0.184, +0.261] | +0.279 [+0.240, +0.320] |
| LEACE (paired, oracle renders) | paired renders (train)       | 0.829   | 0.832 | 0.810    | 0.810  | 0.809     | 0.809   | +0.167 [+0.139, +0.197] | +0.237 [+0.212, +0.263] | +0.294 [+0.265, +0.325] |
| Repair consistency             | artifact pixels (train)      | 0.828   | 0.707 | 0.822    | 0.616  | 0.814     | 0.558   | +0.042 [+0.015, +0.069] | +0.043 [+0.014, +0.072] | +0.042 [+0.012, +0.073] |
| I2E (ours)                     | template library             | 0.848   | 0.844 | 0.850    | 0.833  | 0.847     | 0.818   | +0.179 [+0.152, +0.206] | +0.260 [+0.229, +0.291] | +0.303 [+0.271, +0.336] |
| I2E + balanced (ours)          | templates + A (train)        | 0.841   | 0.842 | 0.843    | 0.840  | 0.844     | 0.842   | +0.177 [+0.151, +0.205] | +0.267 [+0.238, +0.298] | +0.327 [+0.293, +0.362] |
| I2E rank-1 (ablation)          | template library             | 0.811   | 0.799 | 0.809    | 0.770  | 0.806     | 0.730   | +0.134 [+0.115, +0.154] | +0.197 [+0.174, +0.220] | +0.215 [+0.190, +0.240] |
| Insertion augmentation         | template library             | 0.814   | 0.726 | 0.812    | 0.654  | 0.810     | 0.603   | +0.061 [+0.049, +0.074] | +0.081 [+0.068, +0.094] | +0.088 [+0.077, +0.099] |

### 7.2. Mechanism: retention, not occlusion or context

Three explanations were tested separately. *Occlusion*: with the ruler present but uncorrelated with the label (50/50), masking helps at every overlap (+0.036, +0.040, +0.039 at 0/50/100%); covering lesion pixels costs nothing measurable, so the harm requires the correlation. *Lost context*: dilating the mask at 50% overlap *increases* harm ( $-0.097 \rightarrow -0.138 \rightarrow -0.154$  for margins 0/10/25 px) as the margin restores ruler pixels (retention  $0.50 \rightarrow 0.73 \rightarrow 0.95$ ); at 100% overlap, where retention is already 1, harm is flat ( $-0.129$  to  $-0.144$ ) although the ruler’s share of visible pixels falls threefold. *Amplified reliance*: masking doubles the same-head counterfactual sensitivity to the ruler ( $|\Delta p|$   $0.235 \rightarrow 0.491$  at 100%). Harm is largest for small lesions ( $-0.364$  vs  $+0.064$  for large lesions at 100%), where the ruler is a larger share of what survives the mask.

### 7.3. Real hair (ISIC 2019), pre-registered

Two traps share one hair-free group (at most 30 hair pixels): Trap A contains hair lying mostly inside the lesion ( $r \geq 0.5$ ), Trap B hair mostly outside ( $r < 0.1$ ); cells are matched to the hair-free source mix, and five seeds of five-fold group-safe cross-validation give 172 reversed-test melanomas per trap. Masking helps in Trap B (+0.107 [+0.078, +0.135]), harms in Trap A ( $-0.037$  [ $-0.057, -0.019$ ], negative in all five seeds) and the crossover is +0.145 [+0.105, +0.181]. The harm holds in BCN images ( $-0.051$  [ $-0.075, -0.027$ ]); in HAM it is not distinguishable from zero ( $-0.009$  [ $-0.036, +0.017$ ]). Oracle hair inpainting helps in both traps (+0.103, +0.096). Group-balanced training and DFR, which need only an image-level hair label, give the largest gains (+0.225/+0.251 in Trap A, +0.209/+0.215 in Trap B). Paired linear erasure recovers less than half of that (+0.086/+0.081). The magnitude of the in-ROI harm depends on the definition of the comparison group: with a looser hair-free group ( $< 0.1\%$  of pixels) and only heavily haired images, masking gives no benefit but no measurable harm inside the lesion (+0.012 [ $-0.042, +0.062$ ]) while still helping outside (+0.178); the benefit shrinks to zero as more of the hair lies inside the lesion.

### 7.4. What the real-hair shortcut is made of

Removing hair pixels at test time only closes 44–47% of the ERM clean–reversed gap: most of what a frozen encoder exploits in a real hair trap is not the hair pixels but what accompanies hair (hair-bearing

images differ in sex, age and body site). This bounds every pixel-level method, including oracle inpainting and I2E, and explains why image-level-label methods, which remove everything correlated with the label, recover more.

### 7.5. *Reversed-environment gains can be gamed*

Erasing the artifact concept from unpaired image-level labels yields the largest reversed gain in the study (+0.326 in Trap A) while *inverting* the correlated/reversed ordering (0.618 vs 0.817). Post-hoc prevalence-equalised recalibration [12] behaves the same way under the trap’s 90/10 shift (+0.427; correlated 0.433 vs reversed 0.918). Both remove diagnosis-correlated signal or over-correct the group prior rather than removing reliance on the artifact; we therefore read every reversed gain together with clean AUROC and the correlated/reversed ordering.

### 7.6. *Leakage*

Holding cohort, features and protocol fixed, replacing the grouped split with random image-level splits (93 near-duplicate pairs across train/test) inflates the measured shortcut gap by 18% (0% overlap) and 10% (100%) and distorts both masking effects (+0.184  $\rightarrow$  +0.223;  $-0.129 \rightarrow -0.088$ ).

## 8. Results: ultrasound, capsule endoscopy and chest radiography

### 8.1. *Thyroid ultrasound: sonographer calipers*

TN3K images with TNCD labels (3,493 images, 1,210 malignant) carry measurement calipers that a sonographer places on or around the nodule. Calipers were segmented by an audited rule-based detector (image-level precision and recall 0.86–0.89). With the E13 protocol unchanged, masking the nodule helps greatly when the calipers lie outside it (Trap B: +0.341 [+0.320, +0.363]) and much less when they lie inside it (Trap A: +0.111 [+0.097, +0.124]); the crossover is +0.230 [+0.204, +0.255] with DINOv2 and +0.163 [+0.132, +0.194] with MedSigLIP. After masking, the Trap A model still relies on the calipers (correlated 0.902 versus reversed 0.400). With synthetic calipers swept across the nodule boundary (1,663 caliper-free images), masking gains +0.402 [+0.347, +0.458] at no overlap and  $-0.030 [-0.079, +0.016]$  at full overlap (interaction +0.432 [+0.377, +0.492]): the entire benefit of masking disappears, without significant harm.

### 8.2. *Capsule endoscopy: intestinal debris*

SEE-AI frames (erosion 3,550 versus polyp-like 1,931; expert lesion boxes as ROI) with contamination masks from a patch-token probe trained on 2,160 expert-masked frames (held-out pixel accuracy 0.92, IoU 0.62). Leakage groups are contiguous blocks of 100 frame numbers plus perceptual-hash duplicates (no video identifiers are published). Masking helps against out-of-ROI debris (+0.465 [+0.440, +0.490]) and much less against debris on the lesion (+0.097), a crossover of +0.368 [+0.340, +0.397]. With synthetic debris on the

Table 2: In-ROI artifacts (Trap A): reversed-test AUROC, with min(reversed, correlated) in parentheses (penalises shortcut flipping). Best robust value per column in bold. Generated by `scripts/make_main_table.py`.

| Method                         | Needs              | Hair/DINOv2        | Hair/DermLIP       | Thyroid/DINOv2     | Thyroid/MedSigLIP  | Thyroid/ConvNeXt   | Capsule/DINOv2     |            |
|--------------------------------|--------------------|--------------------|--------------------|--------------------|--------------------|--------------------|--------------------|------------|
| ERM                            | —                  | 0.49 (0.49)        | 0.67 (0.67)        | 0.29 (0.29)        | 0.28 (0.28)        | 0.35 (0.35)        | 0.59 (0.59)        | 0.5        |
| ROI masking                    | ROI                | 0.45 (0.45)        | 0.58 (0.58)        | 0.40 (0.40)        | 0.35 (0.35)        | 0.36 (0.36)        | 0.68 (0.68)        | 0.7        |
| Mask + overlay augmentation    | ROI                | 0.46 (0.46)        | —                  | 0.42 (0.42)        | 0.38 (0.38)        | 0.37 (0.37)        | 0.70 (0.70)        | 0.7        |
| <b>U-MtE (ours)</b>            | ROI                | 0.51 (0.51)        | 0.51 (0.51)        | 0.62 (0.62)        | 0.55 (0.55)        | 0.49 (0.49)        | 0.57 (0.57)        | 0.4        |
| <b>U-MtE, protected (ours)</b> | ROI + $Y$ of $A=0$ | 0.49 (0.49)        | 0.59 (0.59)        | 0.62 (0.62)        | 0.61 (0.61)        | 0.55 (0.55)        | 0.67 (0.67)        | 0.6        |
| Group-balanced                 | $A$                | 0.72 (0.72)        | <b>0.81 (0.81)</b> | 0.68 (0.68)        | 0.70 (0.70)        | 0.66 (0.66)        | 0.78 (0.78)        | 0.8        |
| DFR                            | $A$                | <b>0.74 (0.74)</b> | 0.78 (0.78)        | 0.74 (0.66)        | 0.76 (0.72)        | 0.71 (0.67)        | 0.80 (0.80)        | 0.8        |
| Mask + DFR                     | ROI + $A$          | 0.69 (0.69)        | —                  | 0.83 (0.74)        | 0.81 (0.76)        | 0.77 (0.72)        | <b>0.88 (0.88)</b> | <b>0.9</b> |
| <b>U-MtE + balanced (ours)</b> | ROI + $A$          | 0.70 (0.70)        | 0.65 (0.65)        | <b>0.79 (0.79)</b> | <b>0.79 (0.79)</b> | <b>0.74 (0.74)</b> | 0.86 (0.86)        | 0.8        |

796 cleanest frames the reversal appears:  $+0.103$  [ $+0.060, +0.149$ ] at no overlap,  $-0.102$  [ $-0.162, -0.040$ ] at full overlap (interaction  $+0.205$  [ $+0.150, +0.259$ ]).

### 8.3. Chest radiography: a boundary case

With a synthetic tube swept into the lung fields of NIH pneumothorax radiographs (RAD-DINO), the encoder does not use an extra-pulmonary tube at all (reversed 0.913 versus clean 0.915 at no overlap), so masking has nothing to remove; inside the lungs the tube is used (reversed 0.614) and lung masking only partly removes it (0.725). On real devices (RANZCR-CLiP linked to NIH; central lines and endotracheal tubes) masking helped for both locations. The location-dependent *harm* of masking is therefore not universal; what is universal in our data is that masking leaves an in-ROI shortcut largely intact.

## 9. Results: artifact-agnostic mask-then-erase

### 9.1. U-MtE repairs what masking leaves behind

Without any artifact example, mask or label, U-MtE improves reversed AUROC over masking in the in-ROI regime of every modality: thyroid calipers  $+0.222$  [ $+0.197, +0.251$ ] (DINOv2) and  $+0.204$  [ $+0.174, +0.237$ ] (MedSigLIP); dermoscopic hair  $+0.052$  [ $+0.044, +0.062$ ]; synthetic capsule debris  $+0.081$  [ $+0.050, +0.112$ ]; synthetic calipers  $+0.278$  [ $+0.232, +0.325$ ]; synthetic chest tubes  $+0.134$  [ $+0.120, +0.151$ ] (DINOv2 / RAD-DINO). With MedSigLIP the unbalanced method *loses* to masking in both controlled sweeps (debris  $-0.114$  [ $-0.147, -0.082$ ], calipers  $-0.076$  [ $-0.110, -0.043$ ]) although it wins on real thyroid calipers ( $+0.204$ ); the balanced variant beats masking in all four controlled sweeps ( $+0.193$  to  $+0.471$ ) and

has no location interaction in any of them. Disease protection turns the MedSigLIP debris failure into a gain (+0.061 [+0.032, +0.092] over masking) but not the caliper one (−0.041 [−0.089, −0.005]).

### 9.2. Erasure, not augmentation

Using the same generic overlays as training augmentation (overlaid copies of a random half of the training images, in the masked view) gains at most +0.03 over masking; erasing their subspace instead gains +0.207 [+0.182, +0.235] (thyroid, DINOv2), +0.177 [+0.148, +0.211] (thyroid, MedSigLIP) and +0.045 [+0.037, +0.054] (hair) more than augmentation.

### 9.3. When the artifact resembles the disease

On real capsule debris U-MtE is worse than masking (−0.115): erosions carry yellow-white fibrin that resembles debris, and a subspace estimated from real-debris templates is worse still (−0.395). Disease protection removes the failure (−0.010 [−0.020, −0.001] generic, −0.012 [−0.028, +0.003] template) at a cost of at most 0.012 where erasure already works (thyroid −0.005 [−0.036, +0.024], hair −0.012 [−0.020, −0.003]).

### 9.4. With image-level artifact labels

Adding group balancing, U-MtE<sub>bal</sub> is the most robust model in thyroid ultrasound by the minimum of correlated and reversed AUROC (0.788 DINOv2, 0.792 MedSigLIP; DFR 0.659 and 0.719); in real capsule endoscopy masking plus last-layer retraining is higher (0.885 versus 0.857); on dermoscopic hair DFR alone is best (0.741) because most of the hair shortcut is not in the hair pixels (removing every hair patch with the oracle mask gains only +0.015). In both controlled sweeps the benefit of U-MtE<sub>bal</sub> no longer depends on location (interaction −0.000 [−0.023, +0.023] capsule; −0.054 [−0.098, −0.004] thyroid), whereas that of masking does (+0.205, +0.432). The capsule findings replicate with MedSigLIP (crossover +0.332 [+0.296, +0.365]; masking plus DFR 0.901). All thyroid and capsule findings replicate with a CNN (ConvNeXt-Base): crossovers +0.284 and +0.292; U-MtE − mask +0.128 [+0.112, +0.145] and erase − augment +0.120 [+0.102, +0.139] on thyroid.

### 9.5. Label-free and erasure baselines

Among methods that need no artifact labels, U-MtE beats JTT [14] by +0.198 [+0.167, +0.231] (thyroid), +0.137 [+0.098, +0.179] (ovary) and +0.094 [+0.077, +0.109] (chest drains, RAD-DINO; not with DINOv2, −0.010); JTT on the unmasked view is best on hair. Task-preserving concept removal (SPLICE; 1) fitted with artifact labels *flips* the shortcut in thyroid, capsule and hair (reversed 0.83–0.95 but correlated 0.60–0.73): when artifact and label are strongly associated in training their covariance directions are nearly collinear. By min(reversed, correlated) U-MtE<sub>bal</sub> exceeds SPLICE in all three (0.788 vs 0.646; 0.857 vs 0.734; 0.704 vs 0.601).

### 9.6. A single procedure: validation-driven selection

Selecting, per fold, the candidate with the best group-aware validation score (image-level artifact labels needed on validation images only) beats masking in all five cohorts (+0.17 to +0.40), is within 0.02 of the best candidate in each, and beats DFR in three (thyroid +0.057, capsule +0.047, ovary +0.010); it chooses U-MtE<sub>bal</sub> where the artifact is an overlay and plain balancing on hair and drains, where DFR remains better. A split-validation variant that admits DFR as a candidate is reported in the supplement.

### 9.7. Modern inpainting

Inpainting the detected caliper pixels with LaMa [20] and then masking is a strong baseline when artifact masks are available (thyroid +0.226 [+0.213, +0.238], ovary +0.123 [+0.089, +0.155] over masking). U-MtE<sub>protect</sub>, which needs no artifact mask, matches it on thyroid (−0.009 [−0.045, +0.023]) and is below it on ovary (−0.080); with image-level artifact labels, U-MtE<sub>bal</sub> exceeds it on both (+0.162, +0.067).

## 10. Results: natural distribution, fine-tuning and label robustness

### 10.1. Unaltered test distribution

In TN3K, in-ROI calipers occur naturally more often in benign nodules (37.6% vs 24.1%). Training on the official training split and testing on the untouched official test split (614 images; five validation-split seeds), masking leaves overall AUROC unchanged (−0.005 [−0.027, +0.018]) but lowers AUROC on shortcut-conflicting cases (malignant nodules with in-ROI calipers vs benign nodules without) by −0.102 [−0.139, −0.066] and raises it on shortcut-consistent cases (+0.048 [+0.019, +0.079]): masking amplifies the natural in-ROI shortcut. U-MtE recovers half of the loss (+0.051 [+0.033, +0.072] over masking) with the best overall AUROC (0.742); group balancing recovers all of it (+0.142 [+0.107, +0.178]). The disease-protected variant did not help here (pre-registered claims N1–N2 not supported).

### 10.2. End-to-end fine-tuning

Fine-tuning ResNet-50 end-to-end (five folds) reproduces the frozen-backbone findings. Capsule: masking harms in-ROI (−0.284 [−0.347, −0.226]) and helps out-of-ROI (+0.301), crossover +0.584 [+0.469, +0.691]. Thyroid: the crossover points the same way (+0.057 [−0.021, +0.130], five folds), and an end-to-end analogue of U-MtE—an invariance penalty between masked images with and without a generic overlay—improves on masking by +0.143 [+0.099, +0.188] without shortcut flipping (reversed 0.679, correlated 0.703), matching the best label-using arm. It fails on capsule (−0.021), where the artifact resembles the pathology. Ovary (1,202 images) is underpowered: U-MtE end-to-end has the best reversed, correlated and clean AUROC, with CIs overlapping masking.

### 10.3. Robustness to artifact-label definitions

Artifact masks for calipers come from a detector (visual audit: 95% of thyroid detections lie on real calipers; 92–97% of caliper-free images contain no visible caliper). Re-running the thyroid and ovary analyses with large-marker-only ( $\geq 50$  px) and strict-location ( $r \geq 0.7$  /  $r < 0.05$ ) definitions leaves the crossover (thyroid +0.215 to +0.262; ovary +0.170 to +0.195) and the U-MtE<sub>protect</sub> gain over masking (thyroid +0.217 to +0.282; ovary +0.043 to +0.067) significant in every variant; both grow under stricter definitions, as expected if detector errors dilute rather than create the effects.

## 11. Discussion and limitations

*What carries the shortcut decides the method.* Across six real and controlled settings three regimes emerge. (1) *Distinct overlays that carry the shortcut in their pixels* (sonographer calipers, synthetic rulers, debris and tubes): masking removes an out-of-ROI copy but leaves an in-ROI copy intact, and U-MtE removes most of what masking leaves (thyroid +0.22, +0.20, +0.13 on three backbones; controlled sweeps +0.08 to +0.28). (2) *Artifacts that resemble the pathology* (capsule debris versus erosion fibrin): any subspace estimated from artifact appearance removes disease evidence; disease protection makes erasure safe, and with image-level artifact labels masking plus last-layer retraining is the most robust option. (3) *Artifacts whose shortcut is carried by correlates* (dermoscopic hair co-varies with body site, age and sex; a chest drain comes with a treated, re-expanded lung): removing even every artifact patch with an oracle mask gains little (+0.015 for hair), no pixel- or insertion-based method can close the gap, and label-based group methods (DFR, balancing, JTT on hair) are preferable. The oracle-removal ceiling is a cheap diagnostic of the regime.

*Why erasure works where masking and patch removal do not.* In thyroid ultrasound, dropping every patch token that contains a caliper (oracle masks) before mean-pooling improves the reversed AUROC by only +0.047 over pooling all ROI tokens, whereas erasing the generic-overlay subspace from the class-token representation improves on masking by +0.22 (the two representations differ, so the comparison is indicative). Self-attention distributes the information about a salient overlay across the whole representation; an in-ROI artifact is therefore not contained in “its” pixels or patches, and a representation-level operation is needed to remove it.

*Erasure, not augmentation.* The same overlays used as training augmentation gain at most +0.03 over masking in every cohort; used to estimate and erase a subspace they gain up to +0.21 more. Augmentation teaches invariance to the *generated* overlays; erasure removes the directions along which *any* overlay moves the representation, which transfers to real artifacts the generator never produced.

*Reading reversed-environment gains.* A reversed-test gain can be manufactured by removing diagnosis-correlated signal together with the artifact: erasure fitted on unpaired image-level labels and post-hoc prevalence recalibration invert the correlated/reversed ordering. We report reversed, correlated and clean AUROC for every method and summarise robustness by the minimum of correlated and reversed AUROC.

*Decision rule for practitioners.* (1) Measure where artifacts sit relative to the ROI; out-of-ROI artifacts are removed by masking. (2) For in-ROI artifacts without any artifact annotation, use protected U-MtE; it matches oracle-mask LaMa inpainting where the artifact is a distinct overlay. (3) With image-level artifact labels, combine masking or U-MtE with group balancing or last-layer retraining, choosing among them on a group-aware validation split; prefer label-based reweighting alone when the shortcut is carried by correlates of the artifact, which the oracle-removal ceiling (a handful of annotated masks) reveals.

*Limitations.* (i) Backbones are frozen and heads linear; an end-to-end fine-tuning check is reported but not a full sweep. (ii) Artifact masks for thyroid calipers and capsule debris come from an audited rule-based detector (precision/recall 0.86–0.89) and a patch-token probe (IoU 0.62); errors act as label noise on  $A$  and on the overlap strata. (iii) Hair and ruler share one annotation channel in the ISIC 2019 archive. (iv) No patient or video identifiers are published for TN3K or SEE-AI; leakage groups use perceptual hashes and frame blocks. (v) Chest-drain labels for pneumothorax-negative radiographs come from a detector. (vi) The controlled chest-X-ray sweep did not show the sign reversal: the encoder ignored extra-pulmonary tubes. (vii) Reversed environments are constructed by subsampling; real deployment shifts are less extreme, and the clean results bound the cost without shift. (viii) The disease-protection step needs diagnosis labels of artifact-free images, and the balanced variants need image-level artifact labels.

## References

- [1] , 2025. Preserving task-relevant information under linear concept removal, in: OpenReview preprint. SPLICE; <https://openreview.net/pdf/cacf1140306e3ba01f21e975186d4998dffe7a61.pdf>.
- [2] Belrose, N., Schneider-Joseph, D., Ravfogel, S., Cotterell, R., Raff, E., Biderman, S., 2023. LEACE: Perfect linear concept erasure in closed form, in: NeurIPS.
- [3] Bissoto, A., Barata, C., Valle, E., Avila, S., 2023. Test-time selection for robust skin lesion analysis, in: MICCAI Workshops (ISIC).
- [4] Bissoto, A., Valle, E., Avila, S., 2020. Debiasing skin lesion datasets and models? not so fast, in: CVPR Workshops, pp. 740–741.
- [5] Chuang, C.Y., Jampani, V., Li, Y., Torralba, A., Jegelka, S., 2023. Debiasing vision-language models via biased prompts. arXiv preprint arXiv:2302.00070 .
- [6] Cohen, J.P., Viviano, J.D., Bertin, P., et al., 2022. TorchXRayVision: A library of chest x-ray datasets and models. MIDL .
- [7] DeGrave, A.J., Janizek, J.D., Lee, S.I., 2021. AI for radiographic COVID-19 detection selects shortcuts over signal. Nature Machine Intelligence 3, 610–619.

- [8] Fong, R., Vedaldi, A., 2019. Occlusions for effective data augmentation in image classification, in: ICCV Workshops.
- [9] Geirhos, R., Jacobsen, J.H., Michaelis, C., Zemel, R., Brendel, W., Bethge, M., Wichmann, F.A., 2020. Shortcut learning in deep neural networks. *Nature Machine Intelligence* 2, 665–673.
- [10] Jiménez-Sánchez, A., Juodelyte, D., Chamberlain, B., Cheplygina, V., 2023. Detecting shortcuts in medical images – a case study in chest x-rays, in: IEEE ISBI.
- [11] Jin, Q., Gerych, W., Ghassemi, M., 2024. MaskMedPaint: Masked medical image inpainting with diffusion models for mitigation of spurious correlations. [arXiv:2411.10686](#).
- [12] Kina, M.A., Petersen, E., 2026. Prevalence calibration as shortcut mitigation. [arXiv:2609.07922](#).
- [13] Kirichenko, P., Izmailov, P., Wilson, A.G., 2023. Last layer re-training is sufficient for robustness to spurious correlations, in: ICLR.
- [14] Liu, E.Z., Haghighi, B., Chen, A.S., Raghunathan, A., Koh, P.W., Sagawa, S., Liang, P., Finn, C., 2021. Just train twice: Improving group robustness without training group information, in: ICML.
- [15] Nauta, M., Walsh, R., Dubowski, A., Seifert, C., 2022. Uncovering and correcting shortcut learning in machine learning models for skin cancer diagnosis. *Diagnostics* 12, 40.
- [16] Oakden-Rayner, L., Dunnmon, J., Carneiro, G., Ré, C., 2020. Hidden stratification causes clinically meaningful failures in machine learning for medical imaging, in: ACM CHIL, pp. 151–159.
- [17] Pedersen, N., Sydendal, R., Cheplygina, V., Sourget, T., 2026. Look what the probes dragged in! real-world chest x-ray shortcuts in MedCLIP. [arXiv:2608.12086](#).
- [18] Ravfogel, S., Elazar, Y., Gonen, H., Twiton, M., Goldberg, Y., 2020. Null it out: Guarding protected attributes by iterative nullspace projection, in: ACL.
- [19] Sagawa, S., Koh, P.W., Hashimoto, T.B., Liang, P., 2020. Distributionally robust neural networks for group shifts, in: ICLR.
- [20] Suvorov, R., Logacheva, E., Mashikhin, A., Remizova, A., Ashukha, A., Silvestrov, A., Kong, N., Goka, H., Park, K., Lempitsky, V., 2022. Resolution-robust large mask inpainting with fourier convolutions, in: WACV.
- [21] Tschandl, P., Rosendahl, C., Kittler, H., 2018. The HAM10000 dataset, a large collection of multi-source dermatoscopic images of common pigmented skin lesions. *Scientific Data* 5, 180161.
- [22] Wang, X., Peng, Y., Lu, L., Lu, Z., Bagheri, M., Summers, R.M., 2017. ChestX-ray8: Hospital-scale chest x-ray database and benchmarks, in: CVPR.
- [23] Wegley, V., et al., 2026. Dermatology skin lesion image analysis: artifact annotation masks for ISIC 2019. *Scholars' Mine, Missouri S&T*, doi:10.71674/man1-qa33.
- [24] Winkler, J.K., Fink, C., Toberer, F., et al., 2019. Association between surgical skin markings in dermoscopic images and diagnostic performance of a deep learning convolutional neural network for melanoma recognition. *JAMA Dermatology* 155, 1135–1141.
- [25] Winkler, J.K., Sies, K., Fink, C., et al., 2021. Association between different scale bars in dermoscopic images and diagnostic performance of a market-approved deep learning convolutional neural network for melanoma recognition. *European Journal of Cancer* 145, 146–154.
- [26] Yokote, A., et al., 2023. The SEE-AI project dataset: capsule endoscopy images with lesion bounding boxes. *Kaggle dataset capsuleyolo/kyucapsule*, CC BY 4.0.
- [27] Zhao, Q., Lyu, S., Bai, W., Cai, L., Liu, B., Wu, M., Sang, X., Yang, M., Chen, L., 2022. A multi-modality ovarian tumor ultrasound image dataset for unsupervised cross-domain semantic segmentation. [arXiv preprint arXiv:2207.06799](#).
- [28] Zhong, Z., Zheng, L., Kang, G., Li, S., Yang, Y., 2020. Random erasing data augmentation, in: AAAI.