
Masking Cannot Remove What Is Inside the Region of Interest: Location-Dependent Shortcut Mitigation in Medical Imaging

Anonymous Author(s)

Affiliation

Address

email

Abstract

Medical image classifiers exploit acquisition artifacts—hair, sonographer calipers, intestinal debris—that correlate with the diagnosis. The standard defence restricts the model to the region of interest (ROI). We show that this defence fails for artifacts that lie *inside* the ROI, explain why, and evaluate what to do instead. In pre-registered, leakage-safe traps built from real artifacts in four diseases and three imaging modalities, the benefit of ROI masking is always larger for out-of-ROI than for in-ROI artifacts (crossover +0.15 to +0.37 reversed-test AUROC, all Holm-significant), across frozen and fine-tuned backbones and even in the zero-shot scores of vision–language foundation models. Masking becomes harmful when it also removes disease context: in controlled overlap sweeps, with end-to-end fine-tuning (−0.28), and on unaltered test distributions (thyroid: −0.10 on shortcut-conflicting cases; two of three held-out dermoscopy hospitals). A linear-Gaussian analysis predicts these signs in closed form. As a remedy we propose mask-then-erase with a generic insertion library (U-MtE): synthetic overlays that resemble no particular artifact are inserted into masked images and the feature subspace they move is erased, requiring no artifact example, mask or label. U-MtE removes in-ROI shortcuts carried by distinct overlays (+0.13 to +0.22 over masking on three backbones), erasure beats augmentation with the same overlays, a disease-protected variant prevents failure when the artifact resembles the pathology, and text-prompted erasure—the obvious vision–language alternative—is a near no-op. Where the shortcut is carried by correlates of the artifact, label-based reweighting is required. Code, pre-registrations and a backbone-agnostic tool are released.

1 Introduction

Deep classifiers for medical images rely on features that correlate with the label but are not evidence of disease [Geirhos et al., 2020, DeGrave et al., 2021]: surgical markings and rulers change melanoma predictions [Winkler et al., 2019, 2021], and pneumothorax classifiers detect the chest drain that treats the condition [Oakden-Rayner et al., 2020, Jiménez-Sánchez et al., 2023]. The most common defence is to restrict the classifier to the region of interest (ROI) by masking the background. Its rationale assumes that shortcuts live outside the ROI. Hair, calipers and debris do not respect that assumption: when the artifact overlaps the ROI, masking cannot remove it—it removes everything *else*.

We ask whether the value of a mitigation depends on where the artifact sits, which quantity governs that dependence, and how in-ROI shortcuts can be removed without localising the artifact. Our contributions:




Figure 1: Real artifacts inside and outside the region of interest (green: ROI; red: artifact mask). Masking the background removes an out-of-ROI artifact (2nd, 4th) but keeps an in-ROI one (1st, 3rd) while deleting the context around it. Left: ovarian tumours (MMOTU); right: capsule endoscopy (SEE-AI).

1. **A location law for ROI masking**, established with pre-registered real-artifact traps in four diseases (dermoscopic hair; thyroid and ovarian sonographer calipers; capsule-endoscopy debris), replicated with three frozen backbones, two fine-tuned architectures, the zero-shot scores of vision–language models and on unaltered test data.
2. **A closed-form account** (Sec. 5) predicting when in-ROI masking harms, when erasure fails, and why group reweighting is location-invariant.
3. **U-MtE**, an annotation-free remedy, with evidence that erasure (not augmentation) is what works, a disease-protected variant, and a delineation of the regimes in which it does not apply.
4. **Rigour**: pre-registration before every experiment, Holm-corrected primary claims (12/16 supported), replication of an earlier pilot (73/88 numbers), and gaming-aware reporting.

2 Related work

Bissoto et al. [2020] showed that artifact debiasing in dermoscopy is harder than it appears and built trap sets; test-time selection [Bissoto et al., 2023] and generative repair [Jin et al., 2024] address specific artifacts. None treats the artifact’s location relative to the ROI as a variable. Group-robust training [Sagawa et al., 2020, Kirichenko et al., 2023, Liu et al., 2021] needs group labels or error sets; linear concept erasure [Ravfogel et al., 2020, Belrose et al., 2023] and its task-preserving form SPLINCE [Holstege et al., 2025] need concept labels; biased-prompt projection [Chuang et al., 2023] derives directions from text. Random erasing and occlusion augmentation [Zhong et al., 2020, Fong and Vedaldi, 2019] use overlays as training noise. We estimate an artifact subspace from synthetic *insertion* pairs and remove it only after masking.

3 Setup

Traps Let Y be the diagnosis and A the presence of an artifact. A trap trains with $P(A=1 \mid Y=1) = 0.9$, $P(A=1 \mid Y=0) = 0.1$ and tests on a correlated (same rates), a reversed (swapped) and a clean (artifact-free) environment. The overlap $r = |\text{artifact} \cap \text{ROI}|/|\text{artifact}|$ defines Trap A (in-ROI, $r \geq 0.5$) and Trap B (out-of-ROI, $r < 0.1$); both share one artifact-free group, cells are matched on source and label, and splits are grouped by lesion/patient identifiers and perceptual hashes (five seeds $\times$ five folds). The primary metric is reversed-test AUROC; we always report correlated and clean AUROC, and summarise robustness by $\min(\text{reversed}, \text{correlated})$, which penalises methods that merely flip the shortcut. Uncertainty is a hierarchical paired bootstrap (10,000 replicates, clusters = seeds or folds).

Cohorts (Table 1; details in Appendix A.) Heads are linear on frozen DINOv2 ViT-B/14, Med-SigLIP, DermLIP, RAD-DINO or ConvNeXt features; regularisation and threshold are chosen on clean validation data only. Controlled sweeps paste a synthetic artifact (ruler, caliper, debris, tube) at overlap $r \in \{0, 0.25, 0.5, 0.75, 1\}$.

Table 1: Real-artifact cohorts. Artifact masks: published (Wegley et al.), audited detector (95% precision on 40 images), or expert-trained probe (pixel accuracy 0.92).

Cohort	Modality	Images	Artifact (mask source)	ROI
ISIC 2019	dermoscopy	25,331	hair/ruler (published masks)	lesion
TN3K/TNCD	thyroid ultrasound	3,493	calipers (rule-based detector)	nodule
MMOTU	ovarian ultrasound	1,202	calipers/overlays (same detector)	tumour
SEE-AI	capsule endoscopy	5,481	debris/bubbles (patch-token probe)	lesion boxes
NIH ChestX-ray14	radiography	29,687	chest drains (image-level)	lungs

4 Mask-then-erase with a generic insertion library

Removing an in-ROI artifact requires its pixels; inserting one does not. A procedural generator $\mathcal{G}$ (strokes, bands, patches, blobs, dashed lines, glyphs; random colour, opacity and position, half centred in the ROI) turns every masked training image $m(x_i)$ into a counterfactual $m(\mathcal{G}(x_i))$; the feature differences $\Delta_i = f(m(\mathcal{G}(x_i))) - f(m(x_i))$ are independent of Y . U-MtE erases their span,

$$U_k = \text{top-}k \text{ eigvecs} \left(\frac{1}{n} \sum_i \Delta_i \Delta_i^\top \right), \quad P(z) = z - U_k U_k^\top (z - \mu), \quad (1)$$

with k chosen without labels (90% held-out difference energy), and fits the head on $P(f(m(x)))$. The generator contains no ruler, hair, caliper or debris template. *Disease protection* erases $U'_k = \text{orth}((I - WW^\top)U_k)$, where W spans logistic label directions fitted on artifact-free training images, guaranteeing $w^\top P(z) = w^\top z$ for $w \in \text{span}(W)$. With image-level artifact labels the head is group-balanced (U-MtE_{bal}).

5 Why location matters

Let the disease block have class-separation SNR S and the artifact block SNR A per unit of presence. For the LDA-optimal linear head trained with (p_1, p_0) , the artifact's effective SNR is $\tilde{A} = (p_1 - p_0)^2 A / (1 + p_1(1 - p_1)A)$ and

$$\text{AUROC}_{\text{rev}} = \Phi \left(\frac{S - \tilde{A}}{\sqrt{2(S + \tilde{A})}} \right), \quad (2)$$

increasing in S and decreasing in $\tilde{A}$ (simulation agreement ≤ 0.009 AUROC; Appendix B). Masking maps $(S, \tilde{A}) \mapsto (S_m, \tilde{A}_m)$ with $\tilde{A}_m = 0$ out of the ROI and $\tilde{A}_m \geq \tilde{A}$ inside. Hence in-ROI masking harms exactly when it removes disease context ($S_m < S$); the crossover is positive for every S_m ; erasure sets $\tilde{A} \rightarrow 0$ but scales S by $1 - \rho^2$ (ρ : overlap of the erased subspace with the disease direction), so it fails for pathology-like artifacts unless protected; and group reweighting yields $\Phi(\sqrt{S/2})$ regardless of location (Fig. 2).

6 Results

6.1 A location law for ROI masking

In all four real-artifact cohorts the crossover is positive and survives Holm correction (Table 2); it replicates with MedSigLIP and ConvNeXt, with end-to-end fine-tuning (ResNet-50 capsule; ViT-S thyroid), and in the zero-shot scores of MedSigLIP and DermLIP, i.e. before any task training. Masking becomes *harmful* for in-ROI artifacts when it also removes disease context: real dermoscopic hair (-0.037 [$-0.057, -0.019$]), controlled sweeps at full overlap (ovary -0.197 [$-0.292, -0.100$]; capsule -0.102 [$-0.162, -0.040$]; Fig. 3), fine-tuned capsule (-0.284 [$-0.347, -0.226$]) and on unaltered data, where masking lowers AUROC on shortcut-conflicting cases for thyroid and two of three held-out ISIC hospitals (Table 3). The claims are robust to the artifact-label definition (large-marker and strict-location variants) and to a visual audit of the masks (Appendix E). Chest radiography is a boundary case: the encoder ignores extra-pulmonary tubes, so masking has nothing to remove outside the lungs.




Figure 2: Theory (lines) versus simulation (points): reversed-test AUROC as masking removes a growing share of disease signal ($S = 1$, $A = 2$, 90/10 training). In-ROI masking falls below ERM; erasure and reweighting recover.




Figure 3: Controlled overlap sweeps: Δ reversed-test AUROC versus ERM as a synthetic artifact moves into the ROI (95% CIs). Masking (blue) declines with overlap except in chest radiography, where the encoder ignores extra-pulmonary tubes; U-MtE with group balancing (green) is flat.

104 **Mechanism.** Three explanations were tested separately in the controlled dermoscopy sweep (Ap-
 105 pendix C). *Occlusion* is not the cause: with the artifact present but uncorrelated with the label,
 106 masking helps at every overlap. *Lost context* alone is not the cause: dilating the mask at 50% over-
 107 lap increases harm as it restores artifact pixels, and at full overlap harm is flat although the artifact's
 108 share of the visible pixels falls threefold. The governing quantity is artifact *retention*: masking keeps
 109 an in-ROI artifact, and the masked model's counterfactual sensitivity to it roughly doubles.

110 6.2 Removing in-ROI shortcuts

111 Without any artifact example, mask or label, U-MtE improves on masking where the artifact is
 112 a distinct overlay: thyroid $+0.222$ [$+0.197, +0.251$] (DINOv2), $+0.204$ [$+0.174, +0.237$] (Med-
 113 SigLIP), $+0.128$ [$+0.112, +0.145$] (ConvNeXt); ovary $+0.215$ [$+0.182, +0.246$] (MedSigLIP);
 114 controlled sweeps $+0.08$ to $+0.28$; end-to-end (ResNet-50, thyroid) $+0.143$ [$+0.099, +0.188$].
 115 The same overlays used as training augmentation gain at most $+0.03$; erasing their subspace gains
 116 $+0.207$ [$+0.182, +0.235$] more on thyroid. On capsule debris, which resembles erosion fibrin, un-
 117 protected erasure removes disease signal; disease protection removes the failure. With image-level
 118 artifact labels, U-MtE_{bal} is the most robust model on thyroid with all three backbones (Table 4; full

Table 2: Masking gain (mask – ERM, reversed AUROC) is always larger for out-of-ROI than in-ROI artifacts. Crossover = gain(out) – gain(in), 95% CI. Zero-shot: no task training. Natural: shortcut-conflicting pairs of untouched test data.

Setting	Cohort (backbone)	Crossover [95% CI]
Real traps, frozen	ISIC hair (DINOv2)	+0.145 [+0.105, +0.181]
	Thyroid (DINOv2 / MedSigLIP / ConvNeXt)	+0.230 / +0.163 / +0.284
	Ovary (DINOv2)	+0.170 [+0.116, +0.224]
	Capsule (DINOv2 / MedSigLIP / ConvNeXt)	+0.368 / +0.332 / +0.292
Fine-tuned	Capsule (ResNet-50)	+0.584 [+0.469, +0.691]
	Thyroid (ViT-S)	+0.168 [+0.074, +0.265]
Zero-shot VLM	Thyroid / Capsule (MedSigLIP)	+0.103 / +0.049
	Ovary (MedSigLIP) / Hair (DermLIP)	+0.181 / +0.051
Natural, mask – ERM	Thyroid (official split)	−0.102 [−0.139, −0.066]
	ISIC held-out BCN / MSK	−0.016 / −0.056




Figure 4: The 16 pre-registered primary tests with Holm correction (blue: supported, 12/16; red: not supported).

119 values in Table 6). Where the shortcut is carried by correlates of the artifact (hair co-varies with
120 body site, age and sex; a drain comes with a treated lung) no pixel- or insertion-based method helps
121 much and DFR [Kirichenko et al., 2023] is best.

122 **Baselines** SPLINCE [Holstege et al., 2025], fitted with artifact labels, removes much disease sig-
123 nal when artifact and label are strongly associated (clean AUROC 0.56–0.78) and flips the short-
124 cut on capsule; U-MtE_{bal} is more robust on thyroid, capsule and hair (0.788 vs 0.596; 0.857 vs
125 0.683; 0.704 vs 0.564). Text-prompted erasure (biased prompts, Chuang et al., 2023) is a near no-op
126 ($\leq +0.008$ over masking): the direction along which a *described* artifact moves the text embed-
127 ding is not the one along which the real artifact moves the image embedding. Oracle-mask LaMa
128 inpainting [Suvorov et al., 2022] is strong when artifact masks exist; protected U-MtE matches it
129 on thyroid without masks. Label-free, U-MtE beats JTT [Liu et al., 2021] on thyroid (+0.198)
130 and ovary (+0.137). A validation-driven selector beats masking in all five cohorts but not DFR
131 everywhere (Appendix F).

Table 3: Unaltered data (no resampling): change in AUROC on shortcut-conflicting pairs. Thyroid: official TN3K split; ISIC: each hospital held out in turn; capsule: held-out frame blocks. 95% CIs.

Test set	mask – ERM	U-MtE – mask	U-MtE _{bal} – mask
Thyroid (official split)	−0.102 [−0.139, −0.066]	+0.051 [+0.033, +0.072]	–
ISIC, BCN held out	−0.016 [−0.027, −0.005]	+0.003 [+0.000, +0.006]	+0.034 [+0.029, +0.039]
ISIC, MSK held out	−0.056 [−0.084, −0.029]	−0.008 [−0.015, −0.002]	+0.041 [+0.029, +0.053]
ISIC, HAM held out	+0.053 [+0.037, +0.068]	+0.023 [+0.017, +0.032]	+0.065 [+0.056, +0.075]
Capsule (held-out frames)	+0.021 [+0.000, +0.044]	+0.000 [−0.010, +0.009]	+0.014 [+0.003, +0.025]

Table 4: In-ROI artifacts (Trap A): robustness = min(reversed, correlated) AUROC, which penalises both shortcut use and shortcut flipping. Bold: best per column. Full reversed/correlated/clean values for all cohorts and backbones: Table 6.

Method	Needs	Hair DINOv2	Thyroid DINOv2	Thyroid MedSigLIP	Ovary DINOv2	Capsule DINOv2	Drain RAD-DINO
ROI masking	ROI	0.45	0.40	0.35	0.55	0.68	0.24
Mask + overlay augmentation	ROI	0.46	0.42	0.38	0.53	0.70	0.24
U-MtE (ours)	ROI	0.51	0.62	0.55	0.56	0.57	0.27
U-MtE, protected (ours)	ROI	0.49	0.62	0.61	0.59	0.67	0.27
Group-balanced	<i>A</i>	0.72	0.68	0.70	0.71	0.78	0.54
DFR	<i>A</i>	0.74	0.66	0.72	0.69	0.80	0.68
Mask + DFR	ROI + <i>A</i>	0.69	0.74	0.76	0.75	0.88	0.65
U-MtE + balanced (ours)	ROI + <i>A</i>	0.70	0.79	0.79	0.74	0.86	0.46

7 Discussion and limitations

What carries the shortcut decides the method Across six real and controlled settings three regimes emerge. (1) *Distinct overlays that carry the shortcut in their pixels* (sonographer calipers, synthetic rulers, debris and tubes): masking removes an out-of-ROI copy but leaves an in-ROI copy intact, and U-MtE removes most of what masking leaves (thyroid +0.22, +0.20, +0.13 on three backbones; controlled sweeps +0.08 to +0.28). (2) *Artifacts that resemble the pathology* (capsule debris versus erosion fibrin): any subspace estimated from artifact appearance removes disease evidence; disease protection makes erasure safe, and with image-level artifact labels masking plus last-layer retraining is the most robust option. (3) *Artifacts whose shortcut is carried by correlates* (dermoscopic hair co-varies with body site, age and sex; a chest drain comes with a treated, re-expanded lung): removing even every artifact patch with an oracle mask gains little (+0.015 for hair), no pixel- or insertion-based method can close the gap, and label-based group methods (DFR, balancing, JTT on hair) are preferable. The oracle-removal ceiling is a cheap diagnostic of the regime.

Why erasure works where masking and patch removal do not In thyroid ultrasound, dropping every patch token that contains a caliper (oracle masks) before mean-pooling improves the reversed AUROC by only +0.047 over pooling all ROI tokens, whereas erasing the generic-overlay subspace from the class-token representation improves on masking by +0.22 (the two representations differ, so the comparison is indicative). Self-attention distributes the information about a salient overlay across the whole representation; an in-ROI artifact is therefore not contained in “its” pixels or patches, and a representation-level operation is needed to remove it.

Erasure, not augmentation The same overlays used as training augmentation gain at most +0.03 over masking in every cohort; used to estimate and erase a subspace they gain up to +0.21 more. Augmentation teaches invariance to the *generated* overlays; erasure removes the directions along which *any* overlay moves the representation, which transfers to real artifacts the generator never produced.

Reading reversed-environment gains A reversed-test gain can be manufactured by removing diagnosis-correlated signal together with the artifact: erasure fitted on unpaired image-level labels and post-hoc prevalence recalibration invert the correlated/reversed ordering. We report reversed,

161 correlated and clean AUROC for every method and summarise robustness by the minimum of corre-
162 lated and reversed AUROC.

163 **Decision rule for practitioners** (1) Measure where artifacts sit relative to the ROI; out-of-ROI
164 artifacts are removed by masking. (2) For in-ROI artifacts without any artifact annotation, use
165 protected U-MtE; it matches oracle-mask LaMa inpainting where the artifact is a distinct overlay.
166 (3) With image-level artifact labels, combine masking or U-MtE with group balancing or last-layer
167 retraining, choosing among them on a group-aware validation split; prefer label-based reweighting
168 alone when the shortcut is carried by correlates of the artifact, which the oracle-removal ceiling (a
169 handful of annotated masks) reveals.

170 **Limitations** (i) Backbones are frozen and heads linear; an end-to-end fine-tuning check is reported
171 but not a full sweep. (ii) Artifact masks for thyroid calipers and capsule debris come from an audited
172 rule-based detector (precision/recall 0.86–0.89) and a patch-token probe (IoU 0.62); errors act as
173 label noise on A and on the overlap strata. (iii) Hair and ruler share one annotation channel in
174 the ISIC 2019 archive. (iv) No patient or video identifiers are published for TN3K or SEE-AI;
175 leakage groups use perceptual hashes and frame blocks. (v) Chest-drain labels for pneumothorax-
176 negative radiographs come from a detector. (vi) The controlled chest-X-ray sweep did not show
177 the sign reversal: the encoder ignored extra-pulmonary tubes. (vii) Reversed environments are
178 constructed by subsampling; real deployment shifts are less extreme, and the clean results bound the
179 cost without shift. (viii) The disease-protection step needs diagnosis labels of artifact-free images,
180 and the balanced variants need image-level artifact labels.

181 **Broader impact** Shortcut audits that report only overall accuracy can certify masked models that
182 still exploit in-ROI artifacts, and a mitigation that flips a shortcut can look like a success on a
183 reversed test. We recommend reporting artifact location relative to the ROI and performance on
184 shortcut-conflicting cases together with clean and correlated performance.

185 8 Conclusion

186 ROI masking is not neutral: it removes shortcuts outside the ROI and leaves—or amplifies—those in-
187 side it. This holds across diseases, modalities, training regimes and foundation models, and follows
188 from a two-parameter model. Erasing a subspace learned from generic synthetic overlays removes
189 in-ROI shortcuts carried by distinct artifacts without any annotation of the artifact.

190 References

- 191 N. Belrose, D. Schneider-Joseph, S. Ravfogel, R. Cotterell, E. Raff, and S. Biderman. LEACE:
192 Perfect linear concept erasure in closed form. In *NeurIPS*, 2023.
- 193 A. Bissoto, E. Valle, and S. Avila. Debiasing skin lesion datasets and models? not so fast. In *CVPR*
194 *Workshops*, pages 740–741, 2020.
- 195 A. Bissoto, C. Barata, E. Valle, and S. Avila. Test-time selection for robust skin lesion analysis. In
196 *MICCAI Workshops (ISIC)*, 2023.
- 197 C.-Y. Chuang, V. Jampani, Y. Li, A. Torralba, and S. Jegelka. Debiasing vision-language models via
198 biased prompts. *arXiv preprint arXiv:2302.00070*, 2023.
- 199 J. P. Cohen, J. D. Viviano, P. Bertin, et al. TorchXRayVision: A library of chest x-ray datasets and
200 models. *MIDL*, 2022.
- 201 A. J. DeGrave, J. D. Janizek, and S.-I. Lee. AI for radiographic COVID-19 detection selects short-
202 cuts over signal. *Nature Machine Intelligence*, 3:610–619, 2021.
- 203 R. Fong and A. Vedaldi. Occlusions for effective data augmentation in image classification. In *ICCV*
204 *Workshops*, 2019.
- 205 R. Geirhos, J.-H. Jacobsen, C. Michaelis, R. Zemel, W. Brendel, M. Bethge, and F. A. Wichmann.
206 Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2:665–673, 2020.

207 F. Holstege, S. Ravfogel, and B. Wouters. Preserving task-relevant information under linear
208 concept removal. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
209 arXiv:2506.10703.

210 A. Jiménez-Sánchez, D. Juodelyte, B. Chamberlain, and V. Cheplygina. Detecting shortcuts in
211 medical images – a case study in chest x-rays. In *IEEE ISBI*, 2023.

212 Q. Jin, W. Gerych, and M. Ghassemi. MaskMedPaint: Masked medical image inpainting with
213 diffusion models for mitigation of spurious correlations, 2024.

214 M. A. Kina and E. Petersen. Prevalence calibration as shortcut mitigation, 2026.

215 P. Kirichenko, P. Izmailov, and A. G. Wilson. Last layer re-training is sufficient for robustness to
216 spurious correlations. In *ICLR*, 2023.

217 E. Z. Liu, B. Haghighi, A. S. Chen, A. Raghunathan, P. W. Koh, S. Sagawa, P. Liang, and C. Finn.
218 Just train twice: Improving group robustness without training group information. In *ICML*, 2021.

219 L. Oakden-Rayner, J. Dunnmon, G. Carneiro, and C. Ré. Hidden stratification causes clinically
220 meaningful failures in machine learning for medical imaging. In *ACM CHIL*, pages 151–159,
221 2020.

222 S. Ravfogel, Y. Elazar, H. Gonen, M. Twiton, and Y. Goldberg. Null it out: Guarding protected
223 attributes by iterative nullspace projection. In *ACL*, 2020.

224 S. Sagawa, P. W. Koh, T. B. Hashimoto, and P. Liang. Distributionally robust neural networks for
225 group shifts. In *ICLR*, 2020.

226 R. Suvorov, E. Logacheva, A. Mashikhin, A. Remizova, A. Ashukha, A. Silvestrov, N. Kong,
227 H. Goka, K. Park, and V. Lempitsky. Resolution-robust large mask inpainting with fourier convo-
228 lutions. In *WACV*, 2022.

229 P. Tschandl, C. Rosendahl, and H. Kittler. The HAM10000 dataset, a large collection of multi-source
230 dermatoscopic images of common pigmented skin lesions. *Scientific Data*, 5:180161, 2018.

231 X. Wang, Y. Peng, L. Lu, Z. Lu, M. Bagheri, and R. M. Summers. ChestX-ray8: Hospital-scale
232 chest x-ray database and benchmarks. In *CVPR*, 2017.

233 V. Wegley et al. Dermatology skin lesion image analysis: artifact annotation masks for ISIC 2019.
234 Scholars’ Mine, Missouri S&T, doi:10.71674/man1-qa33, 2026.

235 J. K. Winkler, C. Fink, F. Toberer, et al. Association between surgical skin markings in dermoscopic
236 images and diagnostic performance of a deep learning convolutional neural network for melanoma
237 recognition. *JAMA Dermatology*, 155(10):1135–1141, 2019.

238 J. K. Winkler, K. Sies, C. Fink, et al. Association between different scale bars in dermoscopic images
239 and diagnostic performance of a market-approved deep learning convolutional neural network for
240 melanoma recognition. *European Journal of Cancer*, 145:146–154, 2021.

241 A. Yokote et al. The SEE-AI project dataset: capsule endoscopy images with lesion bounding boxes.
242 Kaggle dataset capsuleyolo/kyucapsule, CC BY 4.0, 2023.

243 Q. Zhao, S. Lyu, W. Bai, L. Cai, B. Liu, M. Wu, X. Sang, M. Yang, and L. Chen. A multi-modality
244 ovarian tumor ultrasound image dataset for unsupervised cross-domain semantic segmentation.
245 *arXiv preprint arXiv:2207.06799*, 2022.

246 Z. Zhong, L. Zheng, G. Kang, S. Li, and Y. Yang. Random erasing data augmentation. In *AAAI*,
247 2020.

248 A Data

249 **Controlled dermoscopy cohort** ISIC 2018 Task 1/2 images with manual lesion masks and
 250 melanoma labels from the union of the ISIC 2016 and 2017 tables (three images with conflicting
 251 labels dropped); perceptual-hash (Hamming ≤ 8) and lesion grouping; 2,437 images for which a
 252 65×19 px ruler (0.46% of a 518×518 image) can be placed at 0/25/50/75/100% overlap with the
 253 lesion within 10 percentage points (train/val/test 1,849/371/371 before common support).

254 **Real hair cohort (ISIC 2019)** All 25,331 ISIC 2019 training images; hair/ruler masks from the
 255 archive of Wegley et al. [2026]; lesion masks manual for HAM10000 [Tschandl et al., 2018] and pre-
 256 dicted elsewhere by a U-Net trained on ISIC 2018 Task 1. Two traps share one hair-free group (≤ 30
 257 hair pixels at native resolution): Trap A contains hair with overlap $r = |\text{hair} \cap \text{lesion}|/|\text{hair}| \geq 0.5$,
 258 Trap B $r < 0.1$. Cells are subsampled to the hair-free source mix (HAM/BCN/MSK) within
 259 each label, giving Trap A 704/157/2,738/989 and Trap B 704/157/4,327/930 images (hair-free be-
 260 nign/melanoma, hair benign/melanoma); five seeds of five-fold group-safe cross-validation (leakage
 261 groups: lesion identifiers and pHash); 172 reversed-test melanomas per trap and seed.

262 **Thyroid ultrasound** TN3K images with TNCD benign/malignant labels (3,493 images; 1,210
 263 malignant) and manual nodule masks. Sonographer calipers were segmented by a frozen rule-based
 264 detector (dotted measurement lines and ‘+’ ends; audit in Sec. 7). Trap A: calipers with $r \geq 0.5$
 265 (1,150 images); Trap B: $r < 0.1$ (326); caliper-free: 1,801. No patient identifiers are published;
 266 leakage groups are pHash chains.

267 **Ovarian tumour ultrasound** MMOTU OTU_2d [Zhao et al., 2022] (CC BY 4.0): 1,202 images
 268 with tumour masks, labelled benign cystic (chocolate cyst, serous cystadenoma, simple cyst) versus
 269 lesions with solid components (teratoma, theca-cell tumour, mucinous cystadenoma, high-grade
 270 serous carcinoma). The thyroid detector is applied unchanged: Trap A 341/276, Trap B 91/103,
 271 caliper-free 170/178 (benign-cystic/solid).

272 **Capsule endoscopy** SEE-AI frames [Yokote et al., 2023] (CC BY 4.0) containing only erosions
 273 (3,550) or only polyp-like lesions (1,931); the ROI is the union of the expert lesion boxes. Contam-
 274 ination (debris, bubbles, bile) is segmented by a linear probe on DINOv2 patch tokens trained on
 275 2,160 expert-masked frames (held-out pixel accuracy 0.92, IoU 0.62; expert masks are used where
 276 available). Debris-free: $< 3\%$ of the frame; Trap A: debris $\geq 10\%$ with $r \geq 0.5$; Trap B: $r < 0.1$.
 277 Leakage groups: blocks of 100 consecutive frame numbers plus pHash duplicates.

278 **Chest radiography** NIH ChestX-ray14 [Wang et al., 2017]; lung ROIs from the TorchXRyVision
 279 PSPNet [Cohen et al., 2022]. *Controlled cohort*: 1,906 pneumothorax images without a chest drain
 280 (NEATX annotations) and 5,718 *No Finding* images from other patients; a 100×12 px radio-opaque
 281 tube is placed at controlled overlap with the lungs (7,610 images in common support). *Real drains*
 282 (in-ROI only): 29,687 images; drains are labelled by NEATX for pneumothorax-positive images
 283 and by a drain detector (precision 0.95) for negatives; environments are built within view position
 284 (AP/PA); five patient-grouped folds.

285 B Why location matters: a linear-Gaussian account

286 Let a frozen encoder yield a disease block with class-separation signal-to-noise ratio S and an artifact
 287 block with SNR A per unit of artifact presence; training uses $P(A=1 | Y=1) = p_1$, $P(A=1 |$
 288 $Y=0) = p_0$ and the reversed test swaps them. For the LDA-optimal linear head the artifact’s
 289 effective training SNR is $\tilde{A} = (p_1 - p_0)^2 A / (1 + p_1(1 - p_1)A)$ and

$$\text{AUROC}_{\text{rev}}(S, \tilde{A}) = \Phi \left(\frac{S - \tilde{A}}{\sqrt{2(S + \tilde{A})}} \right), \quad (3)$$

290 increasing in S and decreasing in $\tilde{A}$ (verified by simulation to within 0.009 AUROC; Fig. 2). Mask-
 291 ing maps $(S, \tilde{A}) \mapsto (S_m, \tilde{A}_m)$. Out of the ROI, $\tilde{A}_m = 0$; inside, $\tilde{A}_m \geq \tilde{A}$. Hence (i) in-ROI
 292 masking harms exactly when the mask removes disease-informative context ($S_m < S$) and helps

Table 5: Synthetic ruler, ISIC 2018, DINOv2 ViT-B/14 @518: AUROC and reversed-test vs ERM [95% CI]

Method	Needs	clean@0	rev@0	clean@50	rev@50	clean@100	rev@100	rev@0	rev@50	rev@100
ERM		0.797	0.665	0.793	0.573	0.792	0.515			
ROI masking	ROI mask (train+test)	0.849	0.849	0.829	0.476	0.825	0.386	+0.184 [+0.146, +0.224]	-0.097 [-0.138, -0.056]	-0.129 [-0.170, -0.087]
Inpainting (oracle)	artifact pixels (train+test)	0.810	0.806	0.809	0.783	0.807	0.759	+0.141 [+0.120, +0.162]	+0.211 [+0.187, +0.234]	+0.243 [+0.218, +0.269]
Group-balanced	image-level A (train)	0.814	0.802	0.815	0.792	0.806	0.796	+0.137 [+0.109, +0.168]	+0.219 [+0.188, +0.252]	+0.281 [+0.248, +0.315]
DFR	image-level A (val)	0.786	0.789	0.787	0.794	0.789	0.794	+0.124 [+0.091, +0.160]	+0.221 [+0.184, +0.261]	+0.279 [+0.240, +0.320]
LEACE (paired, oracle renders)	paired renders (train)	0.829	0.832	0.810	0.810	0.809	0.809	+0.167 [+0.139, +0.197]	+0.237 [+0.212, +0.263]	+0.294 [+0.265, +0.325]
Repair consistency	artifact pixels (train)	0.828	0.707	0.822	0.616	0.814	0.558	+0.042 [+0.015, +0.069]	+0.043 [+0.014, +0.072]	+0.042 [+0.012, +0.073]
I2E (ours)	template library	0.848	0.844	0.850	0.833	0.847	0.818	+0.179 [+0.152, +0.206]	+0.260 [+0.229, +0.291]	+0.303 [+0.271, +0.336]
I2E + balanced (ours)	templates + A (train)	0.841	0.842	0.843	0.840	0.844	0.842	+0.177 [+0.151, +0.205]	+0.267 [+0.238, +0.298]	+0.327 [+0.293, +0.362]
I2E rank-1 (ablation)	template library	0.811	0.799	0.809	0.770	0.806	0.730	+0.134 [+0.115, +0.154]	+0.197 [+0.174, +0.220]	+0.215 [+0.190, +0.240]
Insertion augmentation	template library	0.814	0.726	0.812	0.654	0.810	0.603	+0.061 [+0.049, +0.074]	+0.081 [+0.068, +0.094]	+0.088 [+0.077, +0.099]

only by denoising ($S_m > S$); (ii) for any S_m the benefit of masking is larger out of the ROI than inside it—the crossover is always positive; (iii) erasing a subspace that contains the artifact direction sets $\tilde{A} \rightarrow 0$ but scales S by $1 - \rho^2$, where ρ is its overlap with the disease direction, so erasure beats in-ROI masking unless the artifact resembles the pathology, which disease protection prevents; (iv) group balancing removes the artifact–label association and everything correlated with it, giving $\Phi(\sqrt{S/2})$ irrespective of location; (v) a shortcut carried by correlates that neither the mask nor pixel-level erasure touch can only be removed by label-based reweighting. Each prediction is borne out in Sections C–F.

C Results: dermoscopy

C.1 Masking reverses with location (controlled)

With a synthetic ruler whose overlap with the lesion is swept from 0 to 100% (2,437 ISIC 2018 images, DINOv2 ViT-B/14 at 518px, three seeds), ROI masking improves reversed-environment AUROC over ERM by +0.184 [+0.146, +0.224] when the ruler lies outside the lesion and *lowers* it by −0.097 [−0.138, −0.056], −0.112 [−0.155, −0.068] and −0.129 [−0.170, −0.087] at 50, 75 and 100% overlap (location interaction +0.313 [+0.272, +0.353]; Table 5). The sign flip is reproduced at 224px (+0.361 at 0%; −0.095 at 100%) and under randomised ruler colour, opacity, ticks and scale. At 0% overlap the masked model’s correlated and reversed AUROC coincide (0.849/0.849): the shortcut is removed. At 100% the masked model’s correlated AUROC exceeds ERM’s (0.961 vs 0.937): it uses the ruler *more*.

C.2 Mechanism: retention, not occlusion or context

Three explanations were tested separately. *Occlusion*: with the ruler present but uncorrelated with the label (50/50), masking helps at every overlap (+0.036, +0.040, +0.039 at 0/50/100%); covering lesion pixels costs nothing measurable, so the harm requires the correlation. *Lost context*: dilating the mask at 50% overlap *increases* harm (−0.097 → −0.138 → −0.154 for margins 0/10/25 px) as the margin restores ruler pixels (retention 0.50 → 0.73 → 0.95); at 100% overlap, where retention is already 1, harm is flat (−0.129 to −0.144) although the ruler’s share of visible pixels falls threefold. *Amplified reliance*: masking doubles the same-head counterfactual sensitivity to the ruler ($|\Delta p|$ 0.235 → 0.491 at 100%). Harm is largest for small lesions (−0.364 vs +0.064 for large lesions at 100%), where the ruler is a larger share of what survives the mask.

C.3 Real hair (ISIC 2019), pre-registered

Two traps share one hair-free group (at most 30 hair pixels): Trap A contains hair lying mostly inside the lesion ($r \geq 0.5$), Trap B hair mostly outside ($r < 0.1$); cells are matched to the hair-free source mix, and five seeds of five-fold group-safe cross-validation give 172 reversed-test melanomas per trap. Masking helps in Trap B (+0.107 [+0.078, +0.135]), harms in Trap A (−0.037 [−0.057, −0.019], negative in all five seeds) and the crossover is +0.145 [+0.105, +0.181]. The harm holds in BCN images (−0.051 [−0.075, −0.027]); in HAM it is not distinguishable from zero (−0.009 [−0.036, +0.017]). Oracle hair inpainting helps in both traps (+0.103, +0.096). Group-balanced training and DFR, which need only an image-level hair label, give the largest gains (+0.225/+0.251 in Trap A, +0.209/+0.215 in Trap B). Paired linear erasure recovers less than half of that (+0.086/+0.081). The magnitude of the in-ROI harm depends on the definition of the comparison group: with a looser hair-free group (< 0.1% of pixels) and only heavily haired images,

masking gives no benefit but no measurable harm inside the lesion (+0.012 [−0.042, +0.062]) while still helping outside (+0.178); the benefit shrinks to zero as more of the hair lies inside the lesion.

With the dermatology foundation model DermLIP, the controlled ruler sweep reverses as with DINOv2 (+0.161 at no overlap, −0.263 to −0.309 at 25–100%), but on real hair masking harms at both locations (Trap A −0.085 [−0.107, −0.063], Trap B −0.088): with this encoder masking removes disease context regardless of where the hair lies, so no location gap appears.

C.4 What the real-hair shortcut is made of

Removing hair pixels at test time only closes 44–47% of the ERM clean–reversed gap: most of what a frozen encoder exploits in a real hair trap is not the hair pixels but what accompanies hair (hair-bearing images differ in sex, age and body site). This bounds every pixel-level method, including oracle inpainting and U-MtE, and explains why image-level-label methods, which remove everything correlated with the label, recover more.

C.5 Reversed-environment gains can be gamed

Erasing the artifact concept from unpaired image-level labels yields the largest reversed gain in the study (+0.326 in Trap A) while *inverting* the correlated/reversed ordering (0.618 vs 0.817). Post-hoc prevalence-equalisation [Kina and Petersen, 2026] behaves the same way under the trap’s 90/10 shift (+0.427; correlated 0.433 vs reversed 0.918). Both remove diagnosis-correlated signal or over-correct the group prior rather than removing reliance on the artifact; we therefore read every reversed gain together with clean AUROC and the correlated/reversed ordering.

D Results: ultrasound, capsule endoscopy and chest radiography

D.1 Thyroid ultrasound: sonographer calipers

TN3K images with TNCD labels (3,493 images, 1,210 malignant) carry measurement calipers that a sonographer places on or around the nodule. Calipers were segmented by an audited rule-based detector (image-level precision and recall 0.86–0.89). With the E13 protocol unchanged, masking the nodule helps greatly when the calipers lie outside it (Trap B: +0.341 [+0.320, +0.363]) and much less when they lie inside it (Trap A: +0.111 [+0.097, +0.124]); the crossover is +0.230 [+0.204, +0.255] with DINOv2 and +0.163 [+0.132, +0.194] with MedSigLIP. After masking, the Trap A model still relies on the calipers (correlated 0.902 versus reversed 0.400). With synthetic calipers swept across the nodule boundary (1,663 caliper-free images), masking gains +0.402 [+0.347, +0.458] at no overlap and −0.030 [−0.079, +0.016] at full overlap (interaction +0.432 [+0.377, +0.492]): the entire benefit of masking disappears, without significant harm.

D.2 Capsule endoscopy: intestinal debris

SEE-AI frames (erosion 3,550 versus polyp-like 1,931; expert lesion boxes as ROI) with contamination masks from a patch-token probe trained on 2,160 expert-masked frames (held-out pixel accuracy 0.92, IoU 0.62). Leakage groups are contiguous blocks of 100 frame numbers plus perceptual-hash duplicates (no video identifiers are published). Masking helps against out-of-ROI debris (+0.465 [+0.440, +0.490]) and much less against debris on the lesion (+0.097), a crossover of +0.368 [+0.340, +0.397]. With synthetic debris on the 796 cleanest frames the reversal appears: +0.103 [+0.060, +0.149] at no overlap, −0.102 [−0.162, −0.040] at full overlap (interaction +0.205 [+0.150, +0.259]).

D.3 Chest radiography: a boundary case

With a synthetic tube swept into the lung fields of NIH pneumothorax radiographs (RAD-DINO), the encoder does not use an extra-pulmonary tube at all (reversed 0.913 versus clean 0.915 at no overlap), so masking has nothing to remove; inside the lungs the tube is used (reversed 0.614) and lung masking only partly removes it (0.725). On real devices (RANZCR-CLiP linked to NIH; central lines and endotracheal tubes) masking helped for both locations. The location-dependent *harm* of

masking is therefore not universal; what is universal in our data is that masking leaves an in-ROI shortcut largely intact.

E Results: natural distribution, fine-tuning and label robustness

E.1 Unaltered test distributions

Without any resampling, in-ROI artifacts are naturally associated with the label in every cohort (thyroid calipers: 37.6% of benign vs 24.1% of malignant nodules; in-lesion hair: 31% of melanomas vs 20% of benign lesions; capsule debris on the lesion: 23% of polyp-like lesions vs 10% of erosions). We trained on natural training data and tested on untouched data: the official TN3K test split, each ISIC source held out in turn (a real change of hospital and device), and held-out capsule frame blocks. On the shortcut-conflicting pairs, masking *lowers* AUROC for thyroid (-0.102 [$-0.139, -0.066$]) and for two of the three held-out hospitals (BCN -0.016 [$-0.027, -0.005$], MSK -0.056 [$-0.084, -0.029$]); it helps for HAM ($+0.053$) and does not differ significantly from ERM on capsule. U-MtE with group balancing improves on masking for these pairs in every cohort in which it was evaluated ($+0.014$ to $+0.065$) and gives the best or tied-best cross-group AUROC in three of five settings (HAM 0.786, BCN 0.734 vs 0.735, capsule 0.947); plain group balancing is best on thyroid and MSK. Without artifact labels, U-MtE improves on masking for thyroid ($+0.051$) and HAM ($+0.023$) only.

E.2 End-to-end fine-tuning

Fine-tuning ResNet-50 end-to-end (five folds) reproduces the frozen-backbone findings. Capsule: masking harms in-ROI (-0.284 [$-0.347, -0.226$]) and helps out-of-ROI ($+0.301$), crossover $+0.584$ [$+0.469, +0.691$]. Thyroid: the crossover points the same way ($+0.057$ [$-0.021, +0.130$], five folds), and an end-to-end analogue of U-MtE—an invariance penalty between masked images with and without a generic overlay—improves on masking by $+0.143$ [$+0.099, +0.188$] without shortcut flipping (reversed 0.679, correlated 0.703), matching the best label-using arm. It fails on capsule (-0.021), where the artifact resembles the pathology. Ovary (1,202 images) is underpowered: U-MtE end-to-end has the best reversed, correlated and clean AUROC, with CIs overlapping masking. A fine-tuned ViT-S reproduces the crossover on thyroid ($+0.168$ [$+0.074, +0.265$]), but the end-to-end U-MtE analogue does not improve on masking with this architecture ($+0.009$ [$-0.021, +0.040$]).

E.3 Zero-shot vision-language models

Without any task training, the foundation models’ own zero-shot scores (similarity to “benign” versus “malignant” text prompts) are already swayed by artifacts: on capsule frames the correlated and reversed AUROC differ by 0.40. Masking changes these zero-shot scores exactly as the location law predicts: for MedSigLIP it leaves reversed AUROC unchanged for in-ROI artifacts but improves it for out-of-ROI artifacts (thyroid $+0.013$ vs $+0.116$, crossover $+0.103$ [$+0.079, +0.127$]; capsule crossover $+0.049$ [$+0.024, +0.076$]), and it harms more for in-ROI artifacts on ovarian ultrasound (crossover $+0.181$ [$+0.135, +0.228$]) and for DermLIP on hair ($+0.051$ [$+0.031, +0.071$]). The location effect is therefore a property of the representation, not of the classifier trained on top of it.

E.4 Robustness to artifact-label definitions

Artifact masks for calipers come from a detector (visual audit: 95% of thyroid detections lie on real calipers; 92–97% of caliper-free images contain no visible caliper). Re-running the thyroid and ovary analyses with large-marker-only (≥ 50 px) and strict-location ($r \geq 0.7 / r < 0.05$) definitions leaves the crossover (thyroid $+0.215$ to $+0.262$; ovary $+0.170$ to $+0.195$) and the U-MtE_{protect} gain over masking (thyroid $+0.217$ to $+0.282$; ovary $+0.043$ to $+0.067$) significant in every variant; both grow under stricter definitions, as expected if detector errors dilute rather than create the effects.

Table 6: In-ROI artifacts (Trap A): reversed-test AUROC, with min(reversed, correlated) in parentheses (penalises shortcut flipping). Best robust value per column in bold. Generated by scripts/make_main_table.py.

Method	Needs	Hair/DINOv2	Hair/DeinLIP	Thyroid/DINOv2	Thyroid/MedSigLIP	Thyroid/ConvNeXt	Ovary/DINOv2	Ovary/MedSigLIP	Capsule/DINOv2	Capsule/MedSigLIP	Capsule/ConvNeXt	Debris/RAD-DINO	Debris/DINOv2
ERM		0.49 (0.49)	0.67 (0.67)	0.29 (0.29)	0.28 (0.28)	0.35 (0.35)	0.41 (0.41)	0.42 (0.42)	0.59 (0.59)	0.59 (0.59)	0.55 (0.55)	0.19 (0.19)	0.28 (0.28)
ROI masking	ROI	0.45 (0.45)	0.58 (0.58)	0.40 (0.40)	0.35 (0.35)	0.36 (0.36)	0.55 (0.55)	0.43 (0.43)	0.68 (0.68)	0.70 (0.70)	0.64 (0.64)	0.24 (0.24)	0.29 (0.29)
Mask + overlay augmentation	ROI	0.46 (0.46)		0.42 (0.42)	0.38 (0.38)	0.37 (0.37)	0.53 (0.53)		0.70 (0.70)	0.72 (0.72)	0.65 (0.65)	0.24 (0.24)	0.29 (0.29)
U-MtE (ours)	ROI	0.51 (0.51)	0.51 (0.51)	0.62 (0.62)	0.55 (0.55)	0.49 (0.49)	0.56 (0.56)	0.64 (0.64)	0.57 (0.57)	0.47 (0.47)	0.56 (0.56)	0.27 (0.27)	0.28 (0.28)
U-MtE, protected (ours)	ROI + Y of A=0	0.49 (0.49)	0.59 (0.59)	0.62 (0.62)	0.61 (0.61)	0.55 (0.55)	0.59 (0.59)	0.57 (0.57)	0.67 (0.67)	0.67 (0.67)	0.64 (0.64)	0.27 (0.27)	0.28 (0.28)
Group-balanced	A	0.72 (0.72)	0.81 (0.81)	0.68 (0.68)	0.70 (0.70)	0.66 (0.66)	0.71 (0.71)	0.73 (0.73)	0.78 (0.78)	0.81 (0.81)	0.72 (0.72)	0.54 (0.54)	0.40 (0.40)
DFR	A	0.74 (0.74)	0.78 (0.78)	0.74 (0.74)	0.76 (0.72)	0.71 (0.67)	0.72 (0.69)		0.80 (0.80)	0.80 (0.80)	0.77 (0.77)	0.68 (0.68)	0.55 (0.55)
Mask + DFR	ROI + A	0.69 (0.69)		0.83 (0.74)	0.81 (0.76)	0.77 (0.72)	0.75 (0.75)	0.88 (0.88)	0.90 (0.90)	0.86 (0.86)	0.65 (0.65)	0.57 (0.57)	
U-MtE + balanced (ours)	ROI + A	0.70 (0.70)	0.65 (0.65)	0.79 (0.79)	0.79 (0.79)	0.74 (0.74)	0.74 (0.74)	0.79 (0.79)	0.86 (0.86)	0.85 (0.85)	0.80 (0.80)	0.46 (0.46)	0.44 (0.44)

F Results: artifact-agnostic mask-then-erase

F.1 U-MtE repairs what masking leaves behind

Without any artifact example, mask or label, U-MtE improves reversed AUROC over masking in the in-ROI regime of every modality: thyroid calipers +0.222 [+0.197, +0.251] (DINOv2) and +0.204 [+0.174, +0.237] (MedSigLIP); dermoscopic hair +0.052 [+0.044, +0.062]; synthetic capsule debris +0.081 [+0.050, +0.112]; synthetic calipers +0.278 [+0.232, +0.325]; synthetic chest tubes +0.134 [+0.120, +0.151] (DINOv2 / RAD-DINO). With MedSigLIP the unbalanced method *loses* to masking in both controlled sweeps (debris $-0.114 [-0.147, -0.082]$, calipers $-0.076 [-0.110, -0.043]$) although it wins on real thyroid calipers (+0.204); the balanced variant beats masking in all four controlled sweeps (+0.193 to +0.471) and has no location interaction in any of them. Disease protection turns the MedSigLIP debris failure into a gain (+0.061 [+0.032, +0.092] over masking) but not the caliper one ($-0.041 [-0.089, -0.005]$).

F.2 Erasure, not augmentation

Using the same generic overlays as training augmentation (overlaid copies of a random half of the training images, in the masked view) gains at most +0.03 over masking; erasing their subspace instead gains +0.207 [+0.182, +0.235] (thyroid, DINOv2), +0.177 [+0.148, +0.211] (thyroid, MedSigLIP) and +0.045 [+0.037, +0.054] (hair) more than augmentation.

F.3 When the artifact resembles the disease

On real capsule debris U-MtE is worse than masking (-0.115): erosions carry yellow-white fibrin that resembles debris, and a subspace estimated from real-debris templates is worse still (-0.395). Disease protection removes the failure ($-0.010 [-0.020, -0.001]$ generic, $-0.012 [-0.028, +0.003]$ template) at a cost of at most 0.012 where erasure already works (thyroid $-0.005 [-0.036, +0.024]$, hair $-0.012 [-0.020, -0.003]$).

F.4 With image-level artifact labels

Adding group balancing, U-MtE_{bal} is the most robust model in thyroid ultrasound by the minimum of correlated and reversed AUROC (0.788 DINOv2, 0.792 MedSigLIP; DFR 0.659 and 0.719); in real capsule endoscopy masking plus last-layer retraining is higher (0.885 versus 0.857); on dermoscopic hair DFR alone is best (0.741) because most of the hair shortcut is not in the hair pixels (removing every hair patch with the oracle mask gains only +0.015). In both controlled sweeps the benefit of U-MtE_{bal} no longer depends on location (interaction $-0.000 [-0.023, +0.023]$ capsule; $-0.054 [-0.098, -0.004]$ thyroid), whereas that of masking does (+0.205, +0.432). The capsule findings replicate with MedSigLIP (crossover +0.332 [+0.296, +0.365]; masking plus DFR 0.901). All thyroid and capsule findings replicate with a CNN (ConvNeXt-Base): crossovers +0.284 and +0.292; U-MtE – mask +0.128 [+0.112, +0.145] and erase – augment +0.120 [+0.102, +0.139] on thyroid.

F.5 Label-free and erasure baselines

Among methods that need no artifact labels, U-MtE beats JTT [Liu et al., 2021] by +0.198 [+0.167, +0.231] (thyroid), +0.137 [+0.098, +0.179] (ovary) and +0.094 [+0.077, +0.109] (chest drains, RAD-DINO; not with DINOv2, -0.010); JTT on the unmasked view is best on hair. Task-preserving concept removal (SPLINCE; Holstege et al., 2025), fitted with artifact labels in the whitened, minimal-distortion form of the original method, removes the artifact covariance at a large cost in disease signal (clean AUROC 0.56–0.78 versus 0.69–0.87 for masking) and still flips the shortcut on capsule endoscopy (reversed 0.868, correlated 0.683): when artifact and label are strongly associated in training, removing one removes much of the other. By min(reversed, correlated) U-MtE_{bal} exceeds SPLINCE on thyroid, capsule and hair (0.788 vs 0.596; 0.857 vs 0.683; 0.704 vs 0.564); on chest drains SPLINCE is higher (0.609 vs 0.460) but below DFR (0.683).

F.6 Text-prompted erasure

With vision–language backbones an artifact direction can be obtained from text alone (“a thyroid ultrasound image with measurement calipers” minus “a thyroid ultrasound image”; biased-prompt construction, Chuang et al., 2023). Erasing that subspace after masking is a near no-op in every cohort (+0.001 to +0.008 over masking; MedSigLIP thyroid, ovary, capsule and DermLIP hair), whereas U-MtE with the same MedSigLIP features gains +0.204 [+0.174, +0.237] (thyroid) and +0.215 [+0.182, +0.246] (ovary). The direction along which a described artifact moves the *text* embedding is not the direction along which the real artifact moves the *image* embedding; synthetic insertions measure the latter directly.

F.7 A single procedure: validation-driven selection

Selecting, per fold, the candidate with the best group-aware validation score (image-level artifact labels needed on validation images only) beats masking in all five cohorts (+0.17 to +0.40), is within 0.02 of the best candidate in each, and beats DFR in three (thyroid +0.057, capsule +0.047, ovary +0.010); it chooses U-MtE_{bal} where the artifact is an overlay and plain balancing on hair and drains, where DFR remains better. A split-validation variant that admits DFR as a candidate is reported in the supplement.

F.8 Modern inpainting

Inpainting the detected caliper pixels with LaMa [Suvorov et al., 2022] and then masking is a strong baseline when artifact masks are available (thyroid +0.226 [+0.213, +0.238], ovary +0.123 [+0.089, +0.155] over masking). U-MtE_{protect}, which needs no artifact mask, matches it on thyroid (-0.009 [-0.045, +0.023]) and is below it on ovary (-0.080) and on dermoscopic hair, where published hair masks are available (-0.086); with image-level artifact labels, U-MtE_{bal} exceeds it in all three (+0.162, +0.067, +0.123).

G Statistical plan and pre-registration

Every experiment was specified in a dated, publicly versioned document before any model was fitted on its data. The 16 primary tests (location crossover, protected U-MtE vs masking, erase vs augment, U-MtE_{bal} vs balanced; four cohorts each) are Holm-corrected; 12/16 are supported (Fig. 4). Everything else is secondary. Datasets considered and not used (CANDID-PTX: no access; breast ultrasound: no artifact-free group; colonoscopy: no reflection-free images; ISIC ink, cardiomegaly lines: too few images; pathology pen marks: unreliable masks) and the leakage analysis are documented in the repository.

H Compute

All experiments ran on one RTX 3090 (24 GB), 8 CPU cores and 31 GB RAM; the full pipeline takes about two days, dominated by feature extraction and fine-tuning.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and Sec. 1 state the location law, the theory, U-MtE and its limits; each claim is backed in Sec. 6 and Appendices C–F, including the negative results.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Sec. 7 (Limitations) and the regime analysis in Sec. 6 state where masking does not harm, where U-MtE fails and which datasets were infeasible.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Sec. 5 states the assumptions (Gaussian class-conditional features, LDA-optimal linear head, additive artifact block); the derivation and a simulation check are in Appendix B.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Protocol, splits, seeds, thresholds and arms are specified in Sec. 3 and the appendices; every experiment has a dated pre-registration and a make target in the released code.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in

some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All datasets are public; code, frozen splits, pre-registrations and a two-command tool are released (anonymised link at submission).

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so No is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Sec. 3 and Appendix A: cohorts, backbones, regularisation grid, threshold selection on clean validation data, fine-tuning hyper-parameters.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All effects are reported with 95% hierarchical paired bootstrap CIs (10,000 replicates; clusters = seeds/folds); the primary family is Holm-corrected (Appendix G).

Guidelines:

- The answer NA means that the paper does not include experiments.

- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix H: one RTX 3090, 8 CPU cores, 31 GB RAM; about two days end to end.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Only public, de-identified datasets under their licences; no human subjects; results reported whichever way they went.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Sec. 7: masked models can be wrongly certified as shortcut-free; we recommend reporting artifact location and shortcut-conflicting performance. No foreseeable misuse beyond general model deployment risks.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: No high-risk models or data are released; the released artifacts are code and feature-level tools.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Datasets are cited with licences in Appendix A and the repository (ISIC CC BY-NC, MMOTU and SEE-AI CC BY 4.0, NIH ChestX-ray14, TN3K/TNCD, NEATX, RANZCR-CLiP); backbones are used under their model licences.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The code base and the U-MtE/SLAS tools are documented in the repository README with usage examples and tests.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[NA\]](#)

Justification: No crowdsourcing or new human-subject data.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[NA\]](#)

817 Justification: Only publicly released, de-identified retrospective datasets are used; no new
818 data were collected.

819 Guidelines:

- 820 • The answer NA means that the paper does not involve crowdsourcing nor research
821 with human subjects.
- 822 • Depending on the country in which research is conducted, IRB approval (or equiva-
823 lent) may be required for any human subjects research. If you obtained IRB approval,
824 you should clearly state this in the paper.
- 825 • We recognize that the procedures for this may vary significantly between institutions
826 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
827 guidelines for their institution.
- 828 • For initial submissions, do not include any information that would break anonymity
829 (if applicable), such as the institution conducting the review.

830 16. Declaration of LLM usage

831 Question: Does the paper describe the usage of LLMs if it is an important, original, or
832 non-standard component of the core methods in this research? Note that if the LLM is used
833 only for writing, editing, or formatting purposes and does not impact the core methodology,
834 scientific rigorousness, or originality of the research, declaration is not required.

835 Answer: [NA]

836 Justification: LLMs are not a component of the proposed methods. (An AI coding assistant
837 supported implementation and drafting under author supervision; authors should state this
838 per venue policy.)

839 Guidelines:

- 840 • The answer NA means that the core method development in this research does not
841 involve LLMs as any important, original, or non-standard components.
- 842 • Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>)
843 for what should or should not be described.