
Masking Cannot Remove What Is Inside the Region of Interest: Location-Dependent Shortcut Mitigation in Medical Imaging

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Medical image classifiers exploit acquisition artifacts—hair, sonographer calipers,
2 intestinal debris—that correlate with the diagnosis; the standard defence restricts
3 the model to the region of interest (ROI). We show that this defence fails for arti-
4 facts inside the ROI, give a closed-form account of when and why, and evaluate
5 alternatives. In pre-registered, leakage-controlled traps built from real artifacts
6 in four diseases, masking always helps more for out-of-ROI than for in-ROI arti-
7 facts (crossover $+0.15$ to $+0.37$ reversed-test AUROC, all Holm-significant), with
8 frozen and fine-tuned backbones and in the zero-shot scores of vision–language
9 foundation models. Masking becomes harmful when it also removes disease con-
10 text: in dermoscopy, controlled overlap sweeps, fine-tuning (capsule -0.28) and
11 unaltered test data (official thyroid split -0.10 on shortcut-conflicting cases; two
12 of three held-out dermoscopy hospitals). A linear-Gaussian model fitted only to
13 each classifier’s clean and correlated performance matches its held-out reversed-
14 test AUROC (mean absolute error 0.039) and the location crossover across 14
15 cohort–backbone pairs ($r = 0.997$). For frozen foundation-model features we
16 propose mask-then-erase with generic synthetic overlays (U-MtE), which needs
17 no artifact examples, artifact masks or artifact labels—only the ROI mask. It re-
18 moves in-ROI shortcuts carried by distinct overlays (thyroid $+0.13$ to $+0.22$ over
19 masking on three backbones); erasure beats augmentation with the same overlays
20 (thyroid, dermoscopy); a disease-protected variant prevents failure on pathology-
21 like artifacts; text-prompted erasure is a near no-op. Shortcuts carried by cor-
22 relates of the artifact (hair, chest drains; chest radiography is a boundary case)
23 require label-based reweighting. Code, pre-registrations and a backbone-agnostic
24 tool are public.

25 1 Introduction

26 Deep classifiers for medical images rely on features that correlate with the label but are not evi-
27 dence of disease [Geirhos et al., 2020, DeGrave et al., 2021]: surgical markings and rulers change
28 melanoma predictions [Winkler et al., 2019, 2021], and pneumothorax classifiers detect the chest
29 drain that treats the condition [Oakden-Rayner et al., 2020, Jiménez-Sánchez et al., 2023]. The most
30 common defence is to restrict the classifier to the region of interest (ROI) by masking the back-
31 ground. Its rationale assumes that shortcuts live outside the ROI. Hair, calipers and debris do not
32 respect that assumption: when the artifact overlaps the ROI, masking cannot remove it—it removes
33 everything *else*.




Figure 1: Real artifacts inside and outside the region of interest (green: ROI; red: artifact mask). Masking the background removes an out-of-ROI artifact (2nd, 4th) but keeps an in-ROI one (1st, 3rd) while deleting the context around it. Left: ovarian tumours (MMOTU); right: capsule endoscopy (SEE-AI).

We ask whether the value of a mitigation depends on where the artifact sits, which quantity governs that dependence, and how in-ROI shortcuts can be removed without localising the artifact. Our contributions:

1. **A location law for ROI masking**, established with pre-registered real-artifact traps in four diseases (dermoscopic hair; thyroid and ovarian sonographer calipers; capsule-endoscopy debris), replicated with three frozen backbones, two fine-tuned architectures, the zero-shot scores of vision–language models and on unaltered test data.
2. **A closed-form account** (Sec. 5) of when in-ROI masking harms, when erasure fails, and why group reweighting is location-invariant, which matches held-out results of real encoders in a pre-registered test.
3. **U-MtE**, a representation-level remedy for frozen foundation-model features that needs no artifact annotation (only the ROI mask), with evidence that erasure (not augmentation) is what works, a disease-protected variant, and a delineation of the regimes in which it does not apply.
4. **Rigour**: pre-registration before every experiment, Holm-corrected primary claims (12/16 supported), replication of an earlier pilot (73/88 numbers), and gaming-aware reporting.

2 Related work

Bissoto et al. [2020] showed that artifact debiasing in dermoscopy is harder than it appears and built trap sets; test-time selection [Bissoto et al., 2023] and generative repair [Jin et al., 2024] address specific artifacts. Hair occlusion degrades HAM10000 classifiers [Wang et al., 2024]; adding synthetic dark-corner artifacts to training improves classification [Pewton et al., 2024]; calipers are shortcuts even for fetal-ultrasound segmentation [Lin et al., 2024]; and classifiers on frozen foundation-model features show shortcut-like behaviour that pooled evaluation hides [Germani et al., 2026]. None treats the artifact’s location relative to the ROI as a variable. Group-robust training [Sagawa et al., 2020, Kirichenko et al., 2023, Liu et al., 2021] needs group labels or error sets; linear concept erasure [Ravfogel et al., 2020, Belrose et al., 2023] and its task-preserving form SPLINCE [Holstege et al., 2025] need concept labels; biased-prompt projection [Chuang et al., 2023] derives directions from text. Random erasing and occlusion augmentation [Zhong et al., 2020, Fong and Vedaldi, 2019] use overlays as training noise. We estimate an artifact subspace from synthetic *insertion* pairs and remove it only after masking.

3 Setup

Traps Let Y be the diagnosis and A the presence of an artifact. A trap trains with $P(A=1 \mid Y=1) = 0.9$, $P(A=1 \mid Y=0) = 0.1$ and tests on a correlated (same rates), a reversed (swapped) and a clean environment (artifact independent of the label: present in half of each class in real traps, absent in sweeps). The overlap $r = |\text{artifact} \cap \text{ROI}|/|\text{artifact}|$ defines Trap A (in-ROI, $r \geq 0.5$) and

Table 1: Real-artifact cohorts. Artifact masks: published (Wegley et al.), audited detector (95% precision on 40 images), or expert-trained probe (pixel accuracy 0.92).

Cohort	Modality	Images	Artifact (mask source)	ROI
ISIC 2019	dermoscopy	25,331	hair/ruler (published masks)	lesion
TN3K/TNCD	thyroid ultrasound	3,493	calipers (rule-based detector)	nodule
MMOTU	ovarian ultrasound	1,202	calipers/overlays (same detector)	tumour
SEE-AI	capsule endoscopy	5,481	debris/bubbles (patch-token probe)	lesion boxes
NIH ChestX-ray14	radiography	29,687	chest drains (image-level)	lungs

Trap B (out-of-ROI, $r < 0.1$); both share one artifact-free group, cells are matched on source and label, and splits are grouped by lesion/patient identifiers where published, otherwise by perceptual hashes and frame blocks (five seeds $\times$ five folds; a stricter embedding-based grouping is a sensitivity analysis). The primary metric is reversed-test AUROC; we always report correlated and clean AUROC, and summarise robustness by $\min(\text{reversed, correlated})$, which penalises methods that merely flip the shortcut. Uncertainty is a hierarchical paired bootstrap (10,000 replicates, clusters = seeds or folds).

Cohorts (Table 1; details in Appendix A.) Heads are linear on frozen DINOv2 ViT-B/14, MedSigLIP, DermLIP, RAD-DINO or ConvNeXt features; regularisation and threshold are chosen on clean validation data only. Controlled sweeps paste a synthetic artifact (ruler, caliper, debris, tube) at overlap $r \in \{0, 0.25, 0.5, 0.75, 1\}$.

4 Mask-then-erase with a generic insertion library

Removing an in-ROI artifact requires its pixels; inserting one does not. A procedural generator $\mathcal{G}$ (strokes, bands, patches, blobs, dashed lines, glyphs; random colour, opacity and position, half centred in the ROI) turns every masked training image $m(x_i)$ into a counterfactual $m(\mathcal{G}(x_i))$; the feature differences $\Delta_i = f(m(\mathcal{G}(x_i))) - f(m(x_i))$ are independent of Y . U-MtE erases their span,

$$U_k = \text{top-}k \text{ eigvecs} \left(\frac{1}{n} \sum_i \Delta_i \Delta_i^\top \right), \quad P(z) = z - U_k U_k^\top (z - \mu), \quad (1)$$

with k chosen without labels (90% held-out difference energy), and fits the head on $P(f(m(x)))$. The generator contains no ruler, hair, caliper or debris template. *Disease protection* erases $U'_k = \text{orth}((I - WW^\top)U_k)$, where W spans logistic label directions fitted on artifact-free training images, guaranteeing $w^\top P(z) = w^\top z$ for $w \in \text{span}(W)$. With image-level artifact labels the head is group-balanced (U-MtE_{bal}).

5 Why location matters

Let the disease block have class-separation SNR S and the artifact block SNR A per unit of presence. For the LDA-optimal linear head trained with (p_1, p_0) , the artifact’s effective SNR is $\tilde{A} = (p_1 - p_0)^2 A / (1 + p_1(1 - p_1)A)$ and

$$\text{AUROC}_{\text{rev}} = \Phi \left(\frac{S - \tilde{A}}{\sqrt{2(S + \tilde{A})}} \right), \quad (2)$$

increasing in S and decreasing in $\tilde{A}$ (simulation agreement ≤ 0.009 AUROC; Appendix B). Masking maps $(S, \tilde{A}) \mapsto (S_m, \tilde{A}_m)$ with $\tilde{A}_m = 0$ out of the ROI and $\tilde{A}_m \geq \tilde{A}$ inside. Hence in-ROI masking harms exactly when it removes disease context ($S_m < S$); the crossover is positive for every S_m ; erasure sets $\tilde{A} \rightarrow 0$ but scales S by $1 - \rho^2$ (ρ : overlap of the erased subspace with the disease direction), so it fails for pathology-like artifacts unless protected; and group reweighting yields $\Phi(\sqrt{S/2})$ regardless of location (simulation in Appendix Fig. 5).

Test on real encoders The same head gives the AUROC of any test composition (Appendix B). In a pre-registered test we fitted (S, A) for every trained model to its clean and correlated AUROC only,




Figure 2: Pre-registered test of the closed form on real encoders: (S, A) fitted to each model’s clean and correlated AUROC, reversed-test results held out. (a) Reversed-test AUROC (ERM, masking). (b) Location crossover per cohort and backbone (red: end-to-end fine-tuned, outside the model’s assumptions). (c) Mask gain along controlled sweeps.




Figure 3: Controlled overlap sweeps: Δ reversed-test AUROC versus ERM as a synthetic artifact moves into the ROI (95% CIs). Masking (blue) declines with overlap except in chest radiography, where the encoder ignores extra-pulmonary tubes; U-MtE with group balancing (green) is flat.

with no shared or tuned parameter, and compared the implied reversed-test AUROC with the held-out observation (Fig. 2): mean absolute error 0.039 over 148 ERM and masking models ($r = 0.92$; 0.108 for a linear-symmetry heuristic). The implied location crossover agrees with the observed one across 14 cohort–backbone pairs ($r = 0.997$, error 0.014, 13/14 signs; the miss is a chest-radiograph cell at -0.002) and also for fine-tuned networks (4/4 signs). One sign miss means that, by our pre-registered criterion, the model remains a qualitative account; it fails only for chest drains, whose shortcut is carried by correlates that the clean and correlated environments do not expose.

6 Results

6.1 A location law for ROI masking

In all four real-artifact cohorts the crossover is positive and survives Holm correction (Table 2); it replicates with MedSigLIP and ConvNeXt, with end-to-end fine-tuning (ResNet-50 capsule; ViT-S thyroid), and in the zero-shot scores of MedSigLIP and DermLIP, i.e. before any task training. Masking becomes *harmful* for in-ROI artifacts when it also removes disease context: real dermoscopic hair (-0.037 [$-0.057, -0.019$]), controlled sweeps at full overlap (ovary -0.197 [$-0.292, -0.100$];

Table 2: Masking gain (mask – ERM, reversed AUROC) is always larger for out-of-ROI than in-ROI artifacts. Crossover = gain(out) – gain(in), 95% CI. Zero-shot: no task training. Natural: shortcut-conflicting pairs of untouched test data.

Setting	Cohort (backbone)	Crossover [95% CI]
Real traps, frozen	ISIC hair (DINOv2)	+0.145 [+0.105, +0.181]
	Thyroid (DINOv2 / MedSigLIP / ConvNeXt)	+0.230 / +0.163 / +0.284
	Ovary (DINOv2)	+0.170 [+0.116, +0.224]
	Capsule (DINOv2 / MedSigLIP / ConvNeXt)	+0.368 / +0.332 / +0.292
Fine-tuned	Capsule (ResNet-50)	+0.584 [+0.469, +0.691]
	Thyroid (ViT-S)	+0.168 [+0.074, +0.265]
Zero-shot VLM	Thyroid / Capsule (MedSigLIP)	+0.103 / +0.049
	Ovary (MedSigLIP) / Hair (DermLIP)	+0.181 / +0.051
Natural, mask – ERM	Thyroid (official split)	–0.102 [–0.139, –0.066]
	ISIC held-out BCN / MSK	–0.016 / –0.056

Table 3: Unaltered data (no resampling): change in AUROC on shortcut-conflicting pairs. Thyroid: official TN3K split; ISIC: each hospital held out in turn; capsule: held-out frame blocks. 95% CIs.

Test set	mask – ERM	U-MtE – mask	U-MtE _{bal} – mask
Thyroid (official split)	–0.102 [–0.139, –0.066]	+0.051 [+0.033, +0.072]	+0.088 [+0.071, +0.108]
ISIC, BCN held out	–0.016 [–0.027, –0.005]	+0.003 [+0.000, +0.006]	+0.034 [+0.029, +0.039]
ISIC, MSK held out	–0.056 [–0.084, –0.029]	–0.008 [–0.015, –0.002]	+0.041 [+0.029, +0.053]
ISIC, HAM held out	+0.053 [+0.037, +0.068]	+0.023 [+0.017, +0.032]	+0.065 [+0.056, +0.075]
Capsule (held-out frames)	+0.021 [+0.000, +0.044]	+0.000 [–0.010, +0.009]	+0.014 [+0.003, +0.025]

capsule –0.102 [–0.162, –0.040]; Fig. 3), fine-tuned capsule (–0.284 [–0.347, –0.226]) and on unaltered data, where masking lowers AUROC on shortcut-conflicting cases for thyroid and two of three held-out ISIC hospitals (Table 3); at the validation-fixed operating point it halves thyroid sensitivity (0.734 → 0.347; AUPRC, Brier score and calibration in Appendix Table 6). The claims are robust to the artifact-label definition (large-marker and strict-location variants) to a visual audit of the masks and to a stricter, embedding-based leakage grouping (11 of 12 primary estimates unchanged; Appendix E). The law also holds across model scale: with DINOv2 ViT-S, ViT-B and ViT-L the crossover is positive in all nine cohort-size pairs (+0.17 to +0.41), and larger models exploit in-ROI calipers *more* (ERM correlated–reversed gap 0.547 → 0.621 on thyroid, 0.319 → 0.570 on ovary from ViT-S to ViT-L). Chest radiography is a boundary case: the encoder ignores extra-pulmonary tubes, so masking has nothing to remove outside the lungs.

Mechanism. Three explanations were tested separately in the controlled dermoscopy sweep (Appendix C). *Occlusion* is not the cause: with the artifact present but uncorrelated with the label, masking helps at every overlap. *Lost context* alone is not the cause: dilating the mask at 50% overlap increases harm as it restores artifact pixels, and at full overlap harm is flat although the artifact’s share of the visible pixels falls threefold. The governing quantity is artifact *retention*: masking keeps an in-ROI artifact, and the masked model’s counterfactual sensitivity to it roughly doubles.

6.2 Removing in-ROI shortcuts

Without any artifact example, artifact mask or label, U-MtE improves on masking where the artifact is a distinct overlay: thyroid +0.222 [+0.197, +0.251] (DINOv2), +0.204 [+0.174, +0.237] (MedSigLIP), +0.128 [+0.112, +0.145] (ConvNeXt); ovary +0.215 [+0.182, +0.246] (MedSigLIP); controlled sweeps +0.08 to +0.28; end-to-end (ResNet-50, thyroid) +0.143 [+0.099, +0.188]. The same overlays used as training augmentation gain at most +0.03; erasing their subspace gains +0.207 [+0.182, +0.235] more on thyroid. On capsule debris, which resembles erosion fibrin, unprotected erasure removes disease signal; disease protection removes the failure. With image-level artifact labels, U-MtE_{bal} is the most robust model on thyroid with all three backbones (Table 4; full values in Table 7). The gain is stable across model sizes (protected U-MtE on thyroid: +0.181, +0.217, +0.222 for ViT-S/B/L) and erasure ranks ≥ 16 (+0.18 to +0.22), but applied post hoc to




Figure 4: The 16 pre-registered primary tests with Holm correction (blue: supported, 12/16; red: not supported).

Table 4: In-ROI artifacts (Trap A): robustness = $\min(\text{reversed}, \text{correlated})$ AUROC, which penalises both shortcut use and shortcut flipping. Bold: best per column. Full reversed/correlated/clean values for all cohorts and backbones: Table 7.

Method	Needs	Hair DINOv2	Thyroid DINOv2	Thyroid MedSigLIP	Ovary DINOv2	Capsule DINOv2	Drain RAD-DINO
ROI masking	ROI	0.45	0.40	0.35	0.55	0.68	0.24
Mask + overlay augmentation	ROI	0.46	0.42	0.38	0.53	0.70	0.24
U-MtE (ours)	ROI	0.51	0.62	0.55	0.56	0.57	0.27
U-MtE, protected (ours)	ROI	0.49	0.62	0.61	0.59	0.67	0.27
Group-balanced	A	0.72	0.68	0.70	0.71	0.78	0.54
DFR	A	0.74	0.66	0.72	0.69	0.80	0.68
Mask + DFR	ROI + A	0.69	0.74	0.76	0.75	0.88	0.65
U-MtE + balanced (ours)	ROI + A	0.70	0.79	0.79	0.74	0.86	0.46

144 a fine-tuned ResNet-50 it does not help (+0.004): U-MtE is a method for frozen representations.
145 Where the shortcut is carried by correlates of the artifact (hair co-varies with body site, age and
146 sex; a drain comes with a treated lung) no pixel- or insertion-based method helps much and DFR
147 [Kirichenko et al., 2023] is best.

148 **Baselines** SPLINCE [Holstege et al., 2025], fitted with artifact labels, removes much disease sig-
149 nal when artifact and label are strongly associated (clean AUROC 0.56–0.78) and flips the short-
150 cut on capsule; U-MtE_{bal} is more robust on thyroid, capsule and hair (0.788 vs 0.596; 0.857 vs
151 0.683; 0.704 vs 0.564). Text-prompted erasure (biased prompts, Chuang et al., 2023) is a near no-op
152 ($\leq +0.008$ over masking): the direction along which a *described* artifact moves the text embed-
153 ding is not the one along which the real artifact moves the image embedding. Oracle-mask LaMa
154 inpainting [Suvorov et al., 2022] is strong when artifact masks exist; protected U-MtE matches it
155 on thyroid without masks. Label-free, U-MtE beats JTT [Liu et al., 2021] on thyroid (+0.198)
156 and ovary (+0.137). A validation-driven selector beats masking in all five cohorts but not DFR
157 everywhere (Appendix F).

7 Discussion and limitations

What carries the shortcut decides the method Across six real and controlled settings three regimes emerge. (1) *Distinct overlays that carry the shortcut in their pixels* (sonographer calipers, synthetic rulers, debris and tubes): masking removes an out-of-ROI copy but leaves an in-ROI copy intact, and U-MtE removes most of what masking leaves (thyroid +0.22, +0.20, +0.13 on three backbones; controlled sweeps +0.08 to +0.28). (2) *Artifacts that resemble the pathology* (capsule debris versus erosion fibrin): any subspace estimated from artifact appearance removes disease evidence; disease protection makes erasure safe, and with image-level artifact labels masking plus last-layer retraining is the most robust option. (3) *Artifacts whose shortcut is carried by correlates* (dermoscopic hair co-varies with body site, age and sex; a chest drain comes with a treated, re-expanded lung): removing even every artifact patch with an oracle mask gains little (+0.015 for hair), no pixel- or insertion-based method can close the gap, and label-based group methods (DFR, balancing, JTT on hair) are preferable. The oracle-removal ceiling is a cheap diagnostic of the regime.

Why erasure works where masking and patch removal do not In thyroid ultrasound, dropping every patch token that contains a caliper (oracle masks) before mean-pooling improves the reversed AUROC by only +0.047 over pooling all ROI tokens, whereas erasing the generic-overlay subspace from the class-token representation improves on masking by +0.22 (the two representations differ, so the comparison is indicative). Self-attention distributes the information about a salient overlay across the whole representation; an in-ROI artifact is therefore not contained in “its” pixels or patches, and a representation-level operation is needed to remove it.

Erasure, not augmentation The same overlays used as training augmentation gain at most +0.03 over masking in every cohort; used to estimate and erase a subspace they gain up to +0.21 more. Augmentation teaches invariance to the *generated* overlays; erasure removes the directions along which *any* overlay moves the representation, which transfers to real artifacts the generator never produced.

Reading reversed-environment gains A reversed-test gain can be manufactured by removing diagnosis-correlated signal together with the artifact: erasure fitted on unpaired image-level labels and post-hoc prevalence recalibration invert the correlated/reversed ordering. We report reversed, correlated and clean AUROC for every method and summarise robustness by the minimum of correlated and reversed AUROC.

Decision rule for practitioners (1) Measure where artifacts sit relative to the ROI; out-of-ROI artifacts are removed by masking. (2) For in-ROI artifacts without any artifact annotation, use protected U-MtE; it matches oracle-mask LaMa inpainting where the artifact is a distinct overlay. (3) With image-level artifact labels, combine masking or U-MtE with group balancing or last-layer retraining, choosing among them on a group-aware validation split; prefer label-based reweighting alone when the shortcut is carried by correlates of the artifact, which the oracle-removal ceiling (a handful of annotated masks) reveals.

Limitations (i) U-MtE is a representation-level intervention for frozen (foundation-model) features with linear heads, the setting in which it is evaluated throughout; end-to-end fine-tuning is a secondary check (the location law replicates, the U-MtE analogue works on thyroid with ResNet-50 but not with ViT-S, and post-hoc erasure of a fine-tuned network’s features does not help), not a full sweep. (ii) Artifact masks for thyroid calipers and capsule debris come from an audited rule-based detector (precision/recall 0.86–0.89) and a patch-token probe (IoU 0.62); errors act as label noise on A and on the overlap strata. (iii) Hair and ruler share one annotation channel in the ISIC 2019 archive. (iv) No patient or video identifiers are published for TN3K, MMOTU or SEE-AI; leakage groups use perceptual hashes and frame blocks, and a stricter embedding-based grouping leaves 11 of 12 primary estimates unchanged (Sec. E.6). (v) Chest-drain labels for pneumothorax-negative radiographs come from a detector. (vi) The controlled chest-X-ray sweep did not show the sign reversal: the encoder ignored extra-pulmonary tubes. (vii) Reversed environments are constructed by subsampling; real deployment shifts are less extreme, and the clean results bound the cost without shift. (viii) The disease-protection step needs diagnosis labels of artifact-free images, and the balanced variants need image-level artifact labels. (ix) The linear-Gaussian account is fitted per model;

211 it cannot anticipate a shortcut carried by correlates that the clean and correlated environments do not
212 expose (chest drains).

213 **Broader impact** Shortcut audits that report only overall accuracy can certify masked models that
214 still exploit in-ROI artifacts, and a mitigation that flips a shortcut can look like a success on a
215 reversed test. We recommend reporting artifact location relative to the ROI and performance on
216 shortcut-conflicting cases together with clean and correlated performance.

217 8 Conclusion

218 ROI masking is not neutral: it removes shortcuts outside the ROI and leaves—or amplifies—those
219 inside it. This holds across diseases, modalities, training regimes and foundation models, and a two-
220 parameter model fitted to clean and correlated performance matches the held-out results. Erasing
221 a subspace learned from generic synthetic overlays removes in-ROI shortcuts carried by distinct
222 artifacts without any annotation of the artifact.

223 References

- 224 N. Belrose, D. Schneider-Joseph, S. Ravfogel, R. Cotterell, E. Raff, and S. Biderman. LEACE:
225 Perfect linear concept erasure in closed form. In *NeurIPS*, 2023.
- 226 A. Bissoto, E. Valle, and S. Avila. Debiasing skin lesion datasets and models? not so fast. In *CVPR*
227 *Workshops*, pages 740–741, 2020.
- 228 A. Bissoto, C. Barata, E. Valle, and S. Avila. Test-time selection for robust skin lesion analysis. In
229 *MICCAI Workshops (ISIC)*, 2023.
- 230 C.-Y. Chuang, V. Jampani, Y. Li, A. Torralba, and S. Jegelka. Debiasing vision-language models via
231 biased prompts. *arXiv preprint arXiv:2302.00070*, 2023.
- 232 J. P. Cohen, J. D. Viviano, P. Bertin, et al. TorchXRyVision: A library of chest x-ray datasets and
233 models. *MIDL*, 2022.
- 234 A. J. DeGrave, J. D. Janizek, and S.-I. Lee. AI for radiographic COVID-19 detection selects short-
235 cuts over signal. *Nature Machine Intelligence*, 3:610–619, 2021.
- 236 R. Fong and A. Vedaldi. Occlusions for effective data augmentation in image classification. In *ICCV*
237 *Workshops*, 2019.
- 238 R. Geirhos, J.-H. Jacobsen, C. Michaelis, R. Zemel, W. Brendel, M. Bethge, and F. A. Wichmann.
239 Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2:665–673, 2020.
- 240 E. Germani, F. Zeineddine, C. Mourad, and S. Albarqouni. When multi-institutional dataset aggre-
241 gation masks shortcut-like behavior in foundation models for medical imaging. *Medical Image*
242 *Analysis*, 114:104258, 2026. doi: 10.1016/j.media.2026.104258.
- 243 F. Holstege, S. Ravfogel, and B. Wouters. Preserving task-relevant information under linear
244 concept removal. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
245 arXiv:2506.10703.
- 246 A. Jiménez-Sánchez, D. Juodelyte, B. Chamberlain, and V. Cheplygina. Detecting shortcuts in
247 medical images – a case study in chest x-rays. In *IEEE ISBI*, 2023.
- 248 Q. Jin, W. Gerych, and M. Ghassemi. MaskMedPaint: Masked medical image inpainting with
249 diffusion models for mitigation of spurious correlations, 2024.
- 250 M. A. Kina and E. Petersen. Prevalence calibration as shortcut mitigation, 2026.
- 251 P. Kirichenko, P. Izmailov, and A. G. Wilson. Last layer re-training is sufficient for robustness to
252 spurious correlations. In *ICLR*, 2023.

- 253 M. Lin et al. Shortcut learning in medical image segmentation. In *Medical Image Computing*
254 *and Computer Assisted Intervention (MICCAI)*, volume 15008 of *LNCS*, 2024. doi: 10.1007/
255 978-3-031-72111-3_59.
- 256 E. Z. Liu, B. Haghighi, A. S. Chen, A. Raghunathan, P. W. Koh, S. Sagawa, P. Liang, and C. Finn.
257 Just train twice: Improving group robustness without training group information. In *ICML*, 2021.
- 258 L. Oakden-Rayner, J. Dunnmon, G. Carneiro, and C. Ré. Hidden stratification causes clinically
259 meaningful failures in machine learning for medical imaging. In *ACM CHIL*, pages 151–159,
260 2020.
- 261 S. W. Pewton, B. Cassidy, C. Kendrick, and M. H. Yap. Dermoscopic dark corner artifacts removal:
262 Friend or foe? *Computer Methods and Programs in Biomedicine*, 244:107986, 2024. doi: 10.
263 1016/j.cmpb.2023.107986.
- 264 S. Ravfogel, Y. Elazar, H. Gonen, M. Twiton, and Y. Goldberg. Null it out: Guarding protected
265 attributes by iterative nullspace projection. In *ACL*, 2020.
- 266 S. Sagawa, P. W. Koh, T. B. Hashimoto, and P. Liang. Distributionally robust neural networks for
267 group shifts. In *ICLR*, 2020.
- 268 R. Suvorov, E. Logacheva, A. Mashikhin, A. Remizova, A. Ashukha, A. Silvestrov, N. Kong,
269 H. Goka, K. Park, and V. Lempitsky. Resolution-robust large mask inpainting with fourier convo-
270 lutions. In *WACV*, 2022.
- 271 P. Tschandl, C. Rosendahl, and H. Kittler. The HAM10000 dataset, a large collection of multi-source
272 dermatoscopic images of common pigmented skin lesions. *Scientific Data*, 5:180161, 2018.
- 273 X. Wang, Y. Peng, L. Lu, Z. Lu, M. Bagheri, and R. M. Summers. ChestX-ray8: Hospital-scale
274 chest x-ray database and benchmarks. In *CVPR*, 2017.
- 275 Z. Wang, Y. Xue, H. Xi, X. Tan, K. Lin, C. Wang, and J. Zhang. Influence of hair presence on
276 dermoscopic image analysis by AI in skin lesion diagnosis. *Computers in Biology and Medicine*,
277 183:109335, 2024. doi: 10.1016/j.combiomed.2024.109335.
- 278 V. Wegley et al. Dermatology skin lesion image analysis: artifact annotation masks for ISIC 2019.
279 Scholars’ Mine, Missouri S&T, doi:10.71674/man1-qa33, 2026.
- 280 J. K. Winkler, C. Fink, F. Toberer, et al. Association between surgical skin markings in dermoscopic
281 images and diagnostic performance of a deep learning convolutional neural network for melanoma
282 recognition. *JAMA Dermatology*, 155(10):1135–1141, 2019.
- 283 J. K. Winkler, K. Sies, C. Fink, et al. Association between different scale bars in dermoscopic images
284 and diagnostic performance of a market-approved deep learning convolutional neural network for
285 melanoma recognition. *European Journal of Cancer*, 145:146–154, 2021.
- 286 A. Yokote et al. The SEE-AI project dataset: capsule endoscopy images with lesion bounding boxes.
287 Kaggle dataset capsuleyolo/kyucapsule, CC BY 4.0, 2023.
- 288 Q. Zhao, S. Lyu, W. Bai, L. Cai, B. Liu, M. Wu, X. Sang, M. Yang, and L. Chen. A multi-modality
289 ovarian tumor ultrasound image dataset for unsupervised cross-domain semantic segmentation.
290 *arXiv preprint arXiv:2207.06799*, 2022.
- 291 Z. Zhong, L. Zheng, G. Kang, S. Li, and Y. Yang. Random erasing data augmentation. In *AAAI*,
292 2020.

293 A Data

294 **Controlled dermoscopy cohort** ISIC 2018 Task 1/2 images with manual lesion masks and
295 melanoma labels from the union of the ISIC 2016 and 2017 tables (three images with conflicting
296 labels dropped); perceptual-hash (Hamming ≤ 8) and lesion grouping; 2,437 images for which a
297 65×19 px ruler (0.46% of a 518×518 image) can be placed at 0/25/50/75/100% overlap with the
298 lesion within 10 percentage points (train/val/test 1,849/371/371 before common support).

299 **Real hair cohort (ISIC 2019)** All 25,331 ISIC 2019 training images; hair/ruler masks from the
 300 archive of Wegley et al. [2026]; lesion masks manual for HAM10000 [Tschandl et al., 2018] and pre-
 301 dicted elsewhere by a U-Net trained on ISIC 2018 Task 1. Two traps share one hair-free group (≤ 30
 302 hair pixels at native resolution): Trap A contains hair with overlap $r = |\text{hair} \cap \text{lesion}|/|\text{hair}| \geq 0.5$,
 303 Trap B $r < 0.1$. Cells are subsampled to the hair-free source mix (HAM/BCN/MSK) within
 304 each label, giving Trap A 704/157/2,738/989 and Trap B 704/157/4,327/930 images (hair-free be-
 305 nign/melanoma, hair benign/melanoma); five seeds of five-fold group-safe cross-validation (leakage
 306 groups: lesion identifiers and pHash); 172 reversed-test melanomas per trap and seed.

307 **Thyroid ultrasound** TN3K images with TNCD benign/malignant labels (3,493 images; 1,210
 308 malignant) and manual nodule masks. Sonographer calipers were segmented by a frozen rule-based
 309 detector (dotted measurement lines and ‘+’ ends; audit in Sec. 7). Trap A: calipers with $r \geq 0.5$
 310 (1,150 images); Trap B: $r < 0.1$ (326); caliper-free: 1,801. No patient identifiers are published;
 311 leakage groups are pHash chains (stricter embedding-based groups in Sec. E.6).

312 **Ovarian tumour ultrasound** MMOTU OTU_2d [Zhao et al., 2022] (CC BY 4.0): 1,202 images
 313 with tumour masks, labelled benign cystic (chocolate cyst, serous cystadenoma, simple cyst) versus
 314 lesions with solid components (teratoma, theca-cell tumour, mucinous cystadenoma, high-grade
 315 serous carcinoma). The thyroid detector is applied unchanged: Trap A 341/276, Trap B 91/103,
 316 caliper-free 170/178 (benign-cystic/solid).

317 **Capsule endoscopy** SEE-AI frames [Yokote et al., 2023] (CC BY 4.0) containing only erosions
 318 (3,550) or only polyp-like lesions (1,931); the ROI is the union of the expert lesion boxes. Contam-
 319 ination (debris, bubbles, bile) is segmented by a linear probe on DINOv2 patch tokens trained on
 320 2,160 expert-masked frames (held-out pixel accuracy 0.92, IoU 0.62; expert masks are used where
 321 available). Debris-free: $< 3\%$ of the frame; Trap A: debris $\geq 10\%$ with $r \geq 0.5$; Trap B: $r < 0.1$.
 322 Leakage groups: blocks of 100 consecutive frame numbers plus pHash duplicates.

323 **Chest radiography** NIH ChestX-ray14 [Wang et al., 2017]; lung ROIs from the TorchXRyVision
 324 PSPNet [Cohen et al., 2022]. *Controlled cohort*: 1,906 pneumothorax images without a chest drain
 325 (NEATX annotations) and 5,718 *No Finding* images from other patients; a 100×12 px radio-opaque
 326 tube is placed at controlled overlap with the lungs (7,610 images in common support). *Real drains*
 327 (in-ROI only): 29,687 images; drains are labelled by NEATX for pneumothorax-positive images
 328 and by a drain detector (precision 0.95) for negatives; environments are built within view position
 329 (AP/PA); five patient-grouped folds.

330 B Why location matters: a linear-Gaussian account

331 Let a frozen encoder yield a disease block with class-separation signal-to-noise ratio S and an artifact
 332 block with SNR A per unit of artifact presence; training uses $P(A=1 | Y=1) = p_1$, $P(A=1 |$
 333 $Y=0) = p_0$ and the reversed test swaps them. For the LDA-optimal linear head the artifact’s
 334 effective training SNR is $\tilde{A} = (p_1 - p_0)^2 A / (1 + p_1(1 - p_1)A)$ and

$$\text{AUROC}_{\text{rev}}(S, \tilde{A}) = \Phi \left(\frac{S - \tilde{A}}{\sqrt{2(S + \tilde{A})}} \right), \quad (3)$$

335 increasing in S and decreasing in $\tilde{A}$ (verified by simulation to within 0.009 AUROC; Fig. 5). Mask-
 336 ing maps $(S, \tilde{A}) \mapsto (S_m, \tilde{A}_m)$. Out of the ROI, $\tilde{A}_m = 0$; inside, $\tilde{A}_m \geq \tilde{A}$. Hence (i) in-ROI
 337 masking harms exactly when the mask removes disease-informative context ($S_m < S$) and helps
 338 only by denoising ($S_m > S$); (ii) for any S_m the benefit of masking is larger out of the ROI than
 339 inside it—the crossover is always positive; (iii) erasing a subspace that contains the artifact direc-
 340 tion sets $\tilde{A} \rightarrow 0$ but scales S by $1 - \rho^2$, where ρ is its overlap with the disease direction, so erasure
 341 beats in-ROI masking unless the artifact resembles the pathology, which disease protection prevents;
 342 (iv) group balancing removes the artifact–label association and everything correlated with it, giving
 343 $\Phi(\sqrt{S/2})$ irrespective of location; (v) a shortcut carried by correlates that neither the mask nor pixel-
 344 level erasure touch can only be removed by label-based reweighting. Each consequence is borne out
 345 in Sections C–F.




Figure 5: Reversed-test AUROC predicted by the linear-Gaussian model (lines) and simulated with logistic heads (points) as masking removes a growing share of disease signal ($S = 1$, $A = 2$, 90/10 training). Generated by `scripts/analysis/theory_sim.py`.

346 **Does the account match real encoders?** For any test composition $q_y = P(A=1 \mid Y=y)$ the
347 same head gives $\text{AUROC} = \Phi(\Delta / \sqrt{\sigma_1^2 + \sigma_0^2})$ with $\Delta = S + (p_1 - p_0)(q_1 - q_0)A / (1 + vA)$,
348 $\sigma_y^2 = S + w_a^2(1 + Aq_y(1 - q_y))$, $w_a = \sqrt{A}(p_1 - p_0) / (1 + vA)$ and $v = [p_1(1 - p_1) + p_0(1 - p_0)] / 2$.
349 In a pre-registered test we fitted (S, A) for every trained model to its clean and correlated AUROC
350 *only* and compared the implied reversed-test AUROC with the held-out observation; no parameter
351 is shared across models or tuned (Fig. 2). Over 148 ERM and masking models (real traps, chest
352 drains, controlled sweeps) the mean absolute error is 0.039 AUROC ($r = 0.92$), against 0.234 for
353 “reversed = clean” and 0.108 for a linear-symmetry heuristic (2-clean – correlated). The implied
354 location crossover agrees with the observed one across 14 cohort–backbone pairs ($r = 0.997$, mean
355 absolute error 0.014, 13/14 signs; the miss is a chest-radiograph cell with an observed crossover of
356 -0.002), the implied mask gains agree with the controlled sweeps ($r = 0.965$, 42/45 signs), and
357 the agreement extends to end-to-end fine-tuned networks (4/4 crossover signs), which violate the
358 model’s assumptions. Because our pre-registered criterion required every crossover sign to agree,
359 we still call the model a qualitative account. Its one clear failure is the chest-drain cohort (implied
360 0.70–0.76, observed 0.19–0.24): the clean and correlated environments do not expose a shortcut
361 carried by correlates of the drain, regime (v).

362 C Results: dermoscopy

363 C.1 Masking reverses with location (controlled)

364 With a synthetic ruler whose overlap with the lesion is swept from 0 to 100% (2,437 ISIC 2018
365 images, DINOv2 ViT-B/14 at 518 px, three seeds), ROI masking improves reversed-environment
366 AUROC over ERM by $+0.184$ [$+0.146, +0.224$] when the ruler lies outside the lesion and *low-*
367 *ers* it by -0.097 [$-0.138, -0.056$], -0.112 [$-0.155, -0.068$] and -0.129 [$-0.170, -0.087$] at 50,
368 75 and 100% overlap (location interaction $+0.313$ [$+0.272, +0.353$]; Table 5). The sign flip is
369 reproduced at 224 px ($+0.361$ at 0%; -0.095 at 100%) and under randomised ruler colour, opac-
370 ity, ticks and scale. At 0% overlap the masked model’s correlated and reversed AUROC coincide
371 (0.849/0.849): the shortcut is removed. At 100% the masked model’s correlated AUROC exceeds
372 ERM’s (0.961 vs 0.937): it uses the ruler *more*.

373 Insert-then-erase in this table is the artifact-specific precursor of U-MtE: its insertion pairs use tem-
374 plates of the artifact type (rulers) and it erases from unmasked features. U-MtE (Sec. 4) replaces the
375 templates by one generic overlay library for every modality and erases after masking; its dermoscopy
376 results are in Sec. F.

Table 5: Synthetic ruler, ISIC 2018, DINOv2 ViT-B/14 @518: AUROC and reversed-test vs ERM [95% CI]

Method	Needs	clean@0	rev@0	clean@50	rev@50	clean@100	rev@100	rev@0	rev@50	rev@100
ERM		0.797	0.665	0.793	0.573	0.792	0.515			
ROI masking	ROI mask (train+test)	0.849	0.849	0.829	0.476	0.825	0.386	+0.184 [+0.146, +0.224]	-0.097 [-0.138, -0.056]	-0.129 [-0.170, -0.087]
Inpainting (oracle)	artifact pixels (train+test)	0.810	0.806	0.809	0.783	0.807	0.759	+0.141 [+0.120, +0.162]	+0.211 [+0.187, +0.234]	+0.243 [+0.218, +0.269]
Group-balanced	image-level A (train)	0.814	0.802	0.815	0.792	0.806	0.796	+0.137 [+0.109, +0.168]	+0.219 [+0.188, +0.252]	+0.281 [+0.248, +0.315]
DFR	image-level A (val)	0.786	0.789	0.787	0.794	0.789	0.794	+0.124 [+0.091, +0.160]	+0.221 [+0.184, +0.261]	+0.279 [+0.240, +0.320]
LEACE (paired, oracle renders)	paired renders (train)	0.829	0.852	0.810	0.810	0.809	0.809	+0.167 [+0.139, +0.197]	+0.257 [+0.212, +0.263]	+0.294 [+0.265, +0.325]
Repair consistency	artifact pixels (train)	0.828	0.707	0.822	0.616	0.814	0.558	+0.042 [+0.015, +0.069]	+0.043 [+0.014, +0.072]	+0.042 [+0.012, +0.073]
Insert-then-erase (artifact templates)	artifact templates	0.848	0.844	0.850	0.833	0.847	0.818	+0.179 [+0.152, +0.206]	+0.260 [+0.229, +0.291]	+0.303 [+0.271, +0.336]
Insert-then-erase + balanced (artifact templates)	artifact templates + A (train)	0.841	0.842	0.843	0.840	0.844	0.842	+0.177 [+0.151, +0.205]	+0.267 [+0.238, +0.298]	+0.327 [+0.293, +0.362]
Insert-then-erase rank-1 (ablation)	artifact templates	0.811	0.799	0.809	0.770	0.806	0.750	+0.134 [+0.115, +0.154]	+0.197 [+0.174, +0.220]	+0.215 [+0.190, +0.240]
Insertion augmentation	artifact templates	0.814	0.726	0.812	0.654	0.810	0.603	+0.061 [+0.049, +0.074]	+0.081 [+0.068, +0.094]	+0.088 [+0.077, +0.099]




Figure 6: Mechanism: same-head counterfactual sensitivity $|p(x+\text{artifact}) - p(x)|$ versus the artifact’s overlap with the ROI in the five controlled sweeps (mean over seeds; band: range over seeds). Masking removes the out-of-ROI artifact ($r = 0$) but increasingly relies on an in-ROI one; erasure and balancing remove that reliance. Generated by `scripts/analysis/mechanism_figure.py`.

377 C.2 Mechanism: retention, not occlusion or context

378 Three explanations were tested separately. *Occlusion*: with the ruler present but uncorrelated with
379 the label (50/50), masking helps at every overlap (+0.036, +0.040, +0.039 at 0/50/100%); covering
380 lesion pixels costs nothing measurable, so the harm requires the correlation. *Lost context*: dilating
381 the mask at 50% overlap *increases* harm ($-0.097 \rightarrow -0.138 \rightarrow -0.154$ for margins 0/10/25 px) as
382 the margin restores ruler pixels (retention 0.50 $\rightarrow$ 0.73 $\rightarrow$ 0.95); at 100% overlap, where retention is
383 already 1, harm is flat (-0.129 to -0.144) although the ruler’s share of visible pixels falls threefold.
384 *Amplified reliance*: masking doubles the same-head counterfactual sensitivity to the ruler ($|\Delta p|$
385 $0.235 \rightarrow 0.491$ at 100%). Harm is largest for small lesions (-0.364 vs $+0.064$ for large lesions
386 at 100%), where the ruler is a larger share of what survives the mask. Figure 6 shows the same
387 retention mechanism in all five controlled sweeps: with the head held fixed, inserting the artifact
388 changes the masked model’s prediction more the deeper the artifact lies in the ROI. At full overlap
389 this sensitivity exceeds that of ERM for dermoscopic rulers (0.491 vs 0.235), ovarian calipers (0.218
390 vs 0.032) and capsule debris (0.302 vs 0.058) and matches it for thyroid calipers (0.393 vs 0.369).
391 In chest radiography ERM ignores a tube outside the lungs ($|\Delta p| = 0.005$ at $r = 0$), which is why
392 masking has nothing to remove there. U-MtE lowers the masked model’s sensitivity in every sweep
393 (full overlap: 0.338, 0.134, 0.120, 0.164, 0.059), and U-MtE with group balancing lowers it most
394 (≤ 0.083).

395 C.3 Real hair (ISIC 2019), pre-registered

396 Two traps share one hair-free group (at most 30 hair pixels): Trap A contains hair lying mostly
397 inside the lesion ($r \geq 0.5$), Trap B hair mostly outside ($r < 0.1$); cells are matched to the
398 hair-free source mix, and five seeds of five-fold group-safe cross-validation give 172 reversed-test
399 melanomas per trap. Masking helps in Trap B (+0.107 [+0.078, +0.135]), harms in Trap A (-0.037
400 [$-0.057, -0.019$], negative in all five seeds) and the crossover is +0.145 [+0.105, +0.181]. The
401 harm holds in BCN images ($-0.051 [-0.075, -0.027]$); in HAM it is not distinguishable from
402 zero ($-0.009 [-0.036, +0.017]$). Oracle hair inpainting helps in both traps (+0.103, +0.096).
403 Group-balanced training and DFR, which need only an image-level hair label, give the largest gains
404 (+0.225/+0.251 in Trap A, +0.209/+0.215 in Trap B). Paired linear erasure recovers less than
405 half of that (+0.086/+0.081). The magnitude of the in-ROI harm depends on the definition of the
406 comparison group: with a looser hair-free group ($< 0.1\%$ of pixels) and only heavily haired images,
407 masking gives no benefit but no measurable harm inside the lesion ($+0.012 [-0.042, +0.062]$) while
408 still helping outside (+0.178); the benefit shrinks to zero as more of the hair lies inside the lesion.

409 With the dermatology foundation model DermLIP, the controlled ruler sweep reverses as with DI-
410 NOv2 (+0.161 at no overlap, -0.263 to -0.309 at 25–100%), but on real hair masking harms
411 at both locations (Trap A -0.085 [-0.107, -0.063], Trap B -0.088): with this encoder masking
412 removes disease context regardless of where the hair lies, so no location gap appears.

413 C.4 What the real-hair shortcut is made of

414 Removing hair pixels at test time only closes 44–47% of the ERM clean–reversed gap: most of
415 what a frozen encoder exploits in a real hair trap is not the hair pixels but what accompanies hair
416 (hair-bearing images differ in sex, age and body site). This bounds every pixel-level method, in-
417 cluding oracle inpainting and U-MtE, and explains why image-level-label methods, which remove
418 everything correlated with the label, recover more.

419 C.5 Reversed-environment gains can be gamed

420 Erasing the artifact concept from unpaired image-level labels yields the largest reversed gain in the
421 study (+0.326 in Trap A) while *inverting* the correlated/reversed ordering (0.618 vs 0.817). Post-
422 hoc prevalence-equalised recalibration [Kina and Petersen, 2026] behaves the same way under the
423 trap’s 90/10 shift (+0.427; correlated 0.433 vs reversed 0.918). Both remove diagnosis-correlated
424 signal or over-correct the group prior rather than removing reliance on the artifact; we therefore read
425 every reversed gain together with clean AUROC and the correlated/reversed ordering.

426 D Results: ultrasound, capsule endoscopy and chest radiography

427 D.1 Thyroid ultrasound: sonographer calipers

428 TN3K images with TNCD labels (3,493 images, 1,210 malignant) carry measurement calipers that
429 a sonographer places on or around the nodule. Calipers were segmented by an audited rule-based
430 detector (image-level precision and recall 0.86–0.89). With the E13 protocol unchanged, masking
431 the nodule helps greatly when the calipers lie outside it (Trap B: +0.341 [+0.320, +0.363]) and
432 much less when they lie inside it (Trap A: +0.111 [+0.097, +0.124]); the crossover is +0.230
433 [+0.204, +0.255] with DINOv2 and +0.163 [+0.132, +0.194] with MedSigLIP. After masking,
434 the Trap A model still relies on the calipers (correlated 0.902 versus reversed 0.400). With syn-
435 thetic calipers swept across the nodule boundary (1,663 caliper-free images), masking gains +0.402
436 [+0.347, +0.458] at no overlap and -0.030 [-0.079, +0.016] at full overlap (interaction +0.432
437 [+0.377, +0.492]): the entire benefit of masking disappears, without significant harm.

438 D.2 Capsule endoscopy: intestinal debris

439 SEE-AI frames (erosion 3,550 versus polyp-like 1,931; expert lesion boxes as ROI) with contamina-
440 tion masks from a patch-token probe trained on 2,160 expert-masked frames (held-out pixel accuracy
441 0.92, IoU 0.62). Leakage groups are contiguous blocks of 100 frame numbers plus perceptual-hash
442 duplicates (no video identifiers are published). Masking helps against out-of-ROI debris (+0.465
443 [+0.440, +0.490]) and much less against debris on the lesion (+0.097), a crossover of +0.368
444 [+0.340, +0.397]. With synthetic debris on the 796 cleanest frames the reversal appears: +0.103
445 [+0.060, +0.149] at no overlap, -0.102 [-0.162, -0.040] at full overlap (interaction +0.205
446 [+0.150, +0.259]).

447 D.3 Chest radiography: a boundary case

448 With a synthetic tube swept into the lung fields of NIH pneumothorax radiographs (RAD-DINO),
449 the encoder does not use an extra-pulmonary tube at all (reversed 0.913 versus clean 0.915 at no
450 overlap), so masking has nothing to remove; inside the lungs the tube is used (reversed 0.614) and
451 lung masking only partly removes it (0.725). On real devices (RANZCR-CLiP linked to NIH; central
452 lines and endotracheal tubes) masking helped for both locations. The location-dependent *harm* of
453 masking is therefore not universal; what is universal in our data is that masking leaves an in-ROI
454 shortcut largely intact.

Table 6: Clinical secondary metrics on the unaltered test sets (Sec. E): mean (SD) over training seeds. Sensitivity and specificity at the operating point fixed on clean validation data; ECE with 10 equal-width bins; AUPRC baseline = prevalence.

Test set	Arm	AUROC	AUPRC	Brier	ECE	Sens	Spec
Thyroid (official split) ($\pi = 0.38$)	ERM	0.736 (0.009)	0.606 (0.012)	0.207 (0.005)	0.082 (0.008)	0.734 (0.047)	0.619 (0.061)
	mask	0.731 (0.006)	0.632 (0.007)	0.214 (0.006)	0.120 (0.026)	0.347 (0.095)	0.888 (0.041)
	U-MtE	0.742 (0.003)	0.638 (0.007)	0.203 (0.001)	0.085 (0.007)	0.425 (0.088)	0.854 (0.050)
	balanced	0.716 (0.007)	0.584 (0.012)	0.219 (0.004)	0.110 (0.007)	0.714 (0.052)	0.599 (0.061)
	U-MtE _{bal}	0.731 (0.004)	0.630 (0.007)	0.207 (0.001)	0.084 (0.006)	0.399 (0.069)	0.866 (0.034)
ISIC, BCN held out ($\pi = 0.23$)	ERM	0.786 (0.003)	0.586 (0.003)	0.239 (0.006)	0.279 (0.011)	0.842 (0.038)	0.519 (0.075)
	mask	0.769 (0.004)	0.571 (0.009)	0.286 (0.007)	0.332 (0.009)	0.860 (0.013)	0.439 (0.023)
	U-MtE	0.773 (0.004)	0.574 (0.008)	0.280 (0.009)	0.328 (0.012)	0.858 (0.039)	0.448 (0.089)
	balanced	0.781 (0.005)	0.585 (0.009)	0.228 (0.006)	0.257 (0.011)	0.811 (0.046)	0.554 (0.087)
	U-MtE _{bal}	0.774 (0.004)	0.568 (0.008)	0.279 (0.011)	0.325 (0.016)	0.860 (0.027)	0.447 (0.051)
ISIC, MSK held out ($\pi = 0.19$)	ERM	0.770 (0.008)	0.477 (0.012)	0.172 (0.007)	0.161 (0.019)	0.608 (0.049)	0.786 (0.039)
	mask	0.790 (0.004)	0.512 (0.009)	0.146 (0.003)	0.115 (0.009)	0.593 (0.052)	0.824 (0.037)
	U-MtE	0.783 (0.004)	0.500 (0.008)	0.150 (0.003)	0.126 (0.009)	0.566 (0.073)	0.831 (0.049)
	balanced	0.766 (0.008)	0.462 (0.014)	0.178 (0.008)	0.170 (0.020)	0.626 (0.047)	0.755 (0.046)
	U-MtE _{bal}	0.774 (0.006)	0.482 (0.009)	0.154 (0.004)	0.131 (0.010)	0.594 (0.061)	0.798 (0.051)
ISIC, HAM held out ($\pi = 0.11$)	ERM	0.750 (0.005)	0.353 (0.005)	0.162 (0.006)	0.234 (0.011)	0.535 (0.092)	0.813 (0.068)
	mask	0.789 (0.002)	0.404 (0.006)	0.143 (0.005)	0.217 (0.011)	0.559 (0.112)	0.842 (0.063)
	U-MtE	0.799 (0.008)	0.406 (0.009)	0.145 (0.007)	0.226 (0.016)	0.573 (0.067)	0.846 (0.041)
	balanced	0.747 (0.017)	0.346 (0.019)	0.156 (0.014)	0.228 (0.015)	0.539 (0.078)	0.803 (0.076)
	U-MtE _{bal}	0.802 (0.005)	0.401 (0.009)	0.139 (0.006)	0.215 (0.021)	0.606 (0.054)	0.825 (0.037)
Capsule (held-out frames) ($\pi = 0.56$)	ERM	0.889 (0.019)	0.897 (0.046)	0.130 (0.016)	0.042 (0.011)	0.770 (0.062)	0.844 (0.055)
	mask	0.967 (0.011)	0.971 (0.014)	0.069 (0.013)	0.029 (0.011)	0.859 (0.032)	0.933 (0.025)
	U-MtE	0.962 (0.009)	0.967 (0.013)	0.081 (0.017)	0.065 (0.045)	0.871 (0.042)	0.897 (0.043)
	balanced	0.881 (0.024)	0.889 (0.053)	0.136 (0.017)	0.051 (0.009)	0.758 (0.087)	0.838 (0.059)
	U-MtE _{bal}	0.959 (0.008)	0.964 (0.012)	0.079 (0.007)	0.044 (0.012)	0.866 (0.048)	0.897 (0.052)

E Results: natural distribution, fine-tuning and label robustness

E.1 Unaltered test distributions

Without any resampling, in-ROI artifacts are naturally associated with the label in every cohort (thyroid calipers: 37.6% of benign vs 24.1% of malignant nodules; in-lesion hair: 31% of melanomas vs 20% of benign lesions; capsule debris on the lesion: 23% of polyp-like lesions vs 10% of erosions). We trained on natural training data and tested on untouched data: the official TN3K test split, each ISIC source held out in turn (a real change of hospital and device), and held-out capsule frame blocks. On the shortcut-conflicting pairs, masking *lowers* AUROC for thyroid (-0.102 [$-0.139, -0.066$]) and for two of the three held-out hospitals (BCN -0.016 [$-0.027, -0.005$], MSK -0.056 [$-0.084, -0.029$]); it helps for HAM ($+0.053$) and does not differ significantly from ERM on capsule. U-MtE with group balancing improves on masking for these pairs in every cohort ($+0.014$ to $+0.088$) and gives the best or tied-best cross-group AUROC in three of five settings (HAM 0.786, BCN 0.734 vs 0.735, capsule 0.947); plain group balancing is best on thyroid and MSK. Without artifact labels, U-MtE improves on masking for thyroid ($+0.051$) and HAM ($+0.023$) only.

Clinical secondary metrics Table 6 reports AUPRC, Brier score, expected calibration error and sensitivity/specificity at the operating point fixed on validation data. On the official thyroid split masking lowers sensitivity from 0.734 to 0.347 while specificity rises from 0.619 to 0.888 (balanced accuracy $0.677 \rightarrow 0.618$); U-MtE recovers part of the sensitivity (0.425) and has the best AUPRC (0.638) and Brier score (0.203). For the held-out BCN hospital masking worsens both Brier score ($0.239 \rightarrow 0.286$) and calibration (ECE $0.279 \rightarrow 0.332$); for HAM, MSK and capsule it improves every metric, in line with the AUROC results.

E.2 End-to-end fine-tuning

Fine-tuning ResNet-50 end-to-end (five folds) reproduces the frozen-backbone findings. Capsule: masking harms in-ROI (-0.284 [$-0.347, -0.226$]) and helps out-of-ROI ($+0.301$), crossover $+0.584$ [$+0.469, +0.691$]. Thyroid: the crossover points the same way ($+0.057$ [$-0.021, +0.130$], five folds), and an end-to-end analogue of U-MtE—an invariance penalty between masked images

with and without a generic overlay—improves on masking by +0.143 [+0.099, +0.188] without shortcut flipping (reversed 0.679, correlated 0.703), matching the best label-using arm. It fails on capsule (−0.021), where the artifact resembles the pathology. Ovary (1,202 images) is underpowered: U-MtE end-to-end has the best reversed, correlated and clean AUROC, with CIs overlapping masking. A fine-tuned ViT-S reproduces the crossover on thyroid (+0.168 [+0.074, +0.265]), but the end-to-end U-MtE analogue does not improve on masking with this architecture (+0.009 [−0.021, +0.040]). Applying U-MtE post hoc to the penultimate features of the masked, fine-tuned ResNet-50 does not help either (pre-registered; +0.004 [+0.000, +0.012] over the same features without erasure): after task training the shortcut is encoded along task-specific directions that generic overlays do not move. U-MtE is therefore a method for frozen representations.

492 E.3 Zero-shot vision–language models

493 Without any task training, the foundation models’ own zero-shot scores (similarity to “benign” ver-
 494 sus “malignant” text prompts) are already swayed by artifacts: on capsule frames the correlated and
 495 reversed AUROC differ by 0.40. Masking changes these zero-shot scores exactly as the location law
 496 predicts: for MedSigLIP it leaves reversed AUROC unchanged for in-ROI artifacts but improves it
 497 for out-of-ROI artifacts (thyroid +0.013 vs +0.116, crossover +0.103 [+0.079, +0.127]; capsule
 498 crossover +0.049 [+0.024, +0.076]), and it harms more for in-ROI artifacts on ovarian ultrasound
 499 (crossover +0.181 [+0.135, +0.228]) and for DermLIP on hair (+0.051 [+0.031, +0.071]). The
 500 location effect is therefore a property of the representation, not of the classifier trained on top of it.

501 E.4 Robustness to artifact-label definitions

502 Artifact masks for calipers come from a detector (visual audit: 95% of thyroid detections lie on
 503 real calipers; 92–97% of caliper-free images contain no visible caliper). Re-running the thyroid and
 504 ovary analyses with large-marker-only (≥ 50 px) and strict-location ($r \geq 0.7 / r < 0.05$) definitions
 505 leaves the crossover (thyroid +0.215 to +0.262; ovary +0.170 to +0.195) and the U-MtE_{protect} gain
 506 over masking (thyroid +0.217 to +0.282; ovary +0.043 to +0.067) significant in every variant; both
 507 grow under stricter definitions, as expected if detector errors dilute rather than create the effects.

508 E.5 Backbone scale

509 Replacing DINOv2 ViT-B/14 (86 M parameters) by ViT-S/14 (21 M) or ViT-L/14 (304 M), with ev-
 510 erything else fixed (pre-registered), the location crossover is positive in all nine cohort–size combina-
 511 tions (ViT-S / ViT-B / ViT-L: thyroid +0.244 / +0.230 / +0.271; ovary +0.298 / +0.170 / +0.202;
 512 capsule +0.370 / +0.368 / +0.408; every CI excludes zero), and protected U-MtE beats masking
 513 on thyroid and ovary at every size (ViT-S +0.181 [+0.146, +0.218] and +0.103 [+0.078, +0.128];
 514 ViT-L +0.222 [+0.195, +0.256] and +0.060 [+0.043, +0.077]). Larger models exploit the in-ROI
 515 shortcut *more*: the ERM correlated–reversed gap grows from 0.547 to 0.621 (thyroid) and from
 516 0.319 to 0.570 (ovary) between ViT-S and ViT-L, so scale does not remove it. On capsule, disease
 517 protection is less complete with ViT-L (−0.040 [−0.046, −0.035] versus masking).

518 E.6 Stricter leakage grouping

519 TN3K/TNCD, MMOTU and SEE-AI publish no patient or video identifiers, so the main analyses
 520 group near-duplicates by perceptual hash (and, for capsule, blocks of 100 consecutive frames). As
 521 a pre-registered sensitivity analysis we united these groups with connected components of mean-
 522 centred DINOv2 embeddings at the most aggressive cosine threshold that keeps every group below
 523 5% of the cohort (thyroid: 3,455 $\rightarrow$ 3,081 groups, $\tau = 0.80$; ovary: 1,059 $\rightarrow$ 969, $\tau = 0.80$; capsule:
 524 137 $\rightarrow$ 135, $\tau = 0.93$) and re-ran all primary contrasts. Eleven of twelve keep their sign and sig-
 525 nificance: the location crossover is +0.235 [+0.211, +0.259] (thyroid), +0.157 [+0.080, +0.235]
 526 (ovary) and +0.394 [+0.359, +0.430] (capsule), and protected U-MtE still beats masking on thy-
 527 roid (+0.251 [+0.219, +0.279]) and ovary (+0.047 [+0.020, +0.072]). The exception is the small
 528 ovary advantage of erasure over augmentation, which loses significance (+0.020 [−0.011, +0.053]).
 529 Embedding grouping removes visually near-identical images of one patient or session; it cannot ex-
 530 clude differently looking images of the same patient.

Table 7: In-ROI artifacts (Trap A): reversed-test AUROC, with min(reversed, correlated) in parentheses (penalises shortcut flipping). Best robust value per column in bold. Generated by scripts/make_main_table.py.

Method	Needs	Hair/DINOv2	Hair/DeinLIP	Thyroid/DINOv2	Thyroid/MedSigLIP	Thyroid/ConcNEXt	Ovary/DINOv2	Ovary/MedSigLIP	Capsule/DINOv2	Capsule/MedSigLIP	Capsule/ConcNEXt	DeinRAD-DINO	Dein/DINOv2
ERM		0.49 (0.49)	0.67 (0.67)	0.29 (0.29)	0.28 (0.28)	0.35 (0.35)	0.41 (0.41)	0.42 (0.42)	0.59 (0.59)	0.59 (0.59)	0.55 (0.55)	0.19 (0.19)	0.28 (0.28)
ROI masking	ROI	0.45 (0.45)	0.58 (0.58)	0.40 (0.40)	0.35 (0.35)	0.36 (0.36)	0.55 (0.55)	0.43 (0.43)	0.68 (0.68)	0.70 (0.70)	0.64 (0.64)	0.24 (0.24)	0.29 (0.29)
Mask + overlay augmentation	ROI	0.46 (0.46)		0.42 (0.42)	0.38 (0.38)	0.37 (0.37)	0.53 (0.53)		0.70 (0.70)	0.72 (0.72)	0.65 (0.65)	0.24 (0.24)	0.29 (0.29)
U-MtE (ours)	ROI	0.51 (0.51)	0.51 (0.51)	0.62 (0.62)	0.55 (0.55)	0.49 (0.49)	0.56 (0.56)	0.64 (0.64)	0.57 (0.57)	0.47 (0.47)	0.56 (0.56)	0.27 (0.27)	0.28 (0.28)
U-MtE, protected (ours)	ROI + Y of A=0	0.49 (0.49)	0.59 (0.59)	0.62 (0.62)	0.61 (0.61)	0.55 (0.55)	0.59 (0.59)	0.57 (0.57)	0.67 (0.67)	0.67 (0.67)	0.64 (0.64)	0.27 (0.27)	0.28 (0.28)
Group-balanced	A	0.72 (0.72)	0.81 (0.81)	0.68 (0.68)	0.70 (0.70)	0.66 (0.66)	0.71 (0.71)	0.73 (0.73)	0.78 (0.78)	0.81 (0.81)	0.72 (0.72)	0.54 (0.54)	0.40 (0.40)
DFR	A	0.74 (0.74)	0.78 (0.78)	0.74 (0.74)	0.76 (0.72)	0.71 (0.67)	0.72 (0.69)		0.80 (0.80)	0.80 (0.80)	0.77 (0.77)	0.68 (0.68)	0.55 (0.55)
Mask + DFR	ROI + A	0.69 (0.69)		0.83 (0.74)	0.81 (0.76)	0.77 (0.72)	0.75 (0.75)		0.88 (0.88)	0.90 (0.90)	0.86 (0.86)	0.65 (0.65)	0.57 (0.57)
U-MtE + balanced (ours)	ROI + A	0.70 (0.70)	0.65 (0.65)	0.79 (0.79)	0.79 (0.79)	0.74 (0.74)	0.74 (0.74)	0.79 (0.79)	0.86 (0.86)	0.85 (0.85)	0.80 (0.80)	0.46 (0.46)	0.44 (0.44)

F Results: artifact-agnostic mask-then-erase

F.1 U-MtE repairs what masking leaves behind

Without any artifact example, artifact mask or artifact label—it uses only the ROI mask that masking already requires—U-MtE improves reversed AUROC over masking in the in-ROI regime of every modality: thyroid calipers +0.222 [+0.197, +0.251] (DINOv2) and +0.204 [+0.174, +0.237] (MedSigLIP); dermoscopic hair +0.052 [+0.044, +0.062]; synthetic capsule debris +0.081 [+0.050, +0.112]; synthetic calipers +0.278 [+0.232, +0.325]; synthetic chest tubes +0.134 [+0.120, +0.151] (DINOv2 / RAD-DINO). With MedSigLIP the unbalanced method *loses* to masking in both controlled sweeps (debris −0.114 [−0.147, −0.082], calipers −0.076 [−0.110, −0.043]) although it wins on real thyroid calipers (+0.204); the balanced variant beats masking in all four controlled sweeps (+0.193 to +0.471) and has no location interaction in any of them. Disease protection turns the MedSigLIP debris failure into a gain (+0.061 [+0.032, +0.092] over masking) but not the caliper one (−0.041 [−0.089, −0.005]).

Sensitivity to the erasure rank In a pre-registered ablation (Trap A, DINOv2), the thyroid gain over masking is +0.18 to +0.22 for every setting that erases at least 16 directions (held-out energy 0.5–0.99, or fixed $k \in \{16, 64\}$), but only +0.06 to +0.07 with one to four directions: the shift caused by overlays is multi-dimensional. The pre-specified 90% rule reached the rank cap ($k = 64$) in all three cohorts. The protected variant never harms capsule by more than 0.010 and helps ovary at every rank (+0.007 to +0.043), whereas unprotected erasure harms capsule more the more directions it removes, as the overlap term ρ of Sec. B predicts.

F.2 Erasure, not augmentation

Using the same generic overlays as training augmentation (overlaid copies of a random half of the training images, in the masked view) gains at most +0.03 over masking; erasing their subspace instead gains +0.207 [+0.182, +0.235] (thyroid, DINOv2), +0.177 [+0.148, +0.211] (thyroid, MedSigLIP) and +0.045 [+0.037, +0.054] (hair) more than augmentation. Under the stricter near-duplicate grouping (Sec. E.6) the thyroid advantage is unchanged (+0.192) but the small ovary advantage (+0.035) loses significance, so we regard erase > augment as established for thyroid and hair only.

F.3 When the artifact resembles the disease

On real capsule debris U-MtE is worse than masking (−0.115): erosions carry yellow-white fibrin that resembles debris, and a subspace estimated from real-debris templates is worse still (−0.395). Disease protection removes the failure (−0.010 [−0.020, −0.001] generic, −0.012 [−0.028, +0.003] template) at a cost of at most 0.012 where erasure already works (thyroid −0.005 [−0.036, +0.024], hair −0.012 [−0.020, −0.003]).

F.4 With image-level artifact labels

Adding group balancing, U-MtE_{bal} is the most robust model in thyroid ultrasound by the minimum of correlated and reversed AUROC (0.788 DINOv2, 0.792 MedSigLIP; DFR 0.659 and 0.719); in

real capsule endoscopy masking plus last-layer retraining is higher (0.885 versus 0.857); on dermoscopic hair DFR alone is best (0.741) because most of the hair shortcut is not in the hair pixels (removing every hair patch with the oracle mask gains only +0.015). In both controlled sweeps the benefit of U-MtE_{bal} no longer depends on location (interaction $-0.000 [-0.023, +0.023]$ capsule; $-0.054 [-0.098, -0.004]$ thyroid), whereas that of masking does (+0.205, +0.432). The capsule findings replicate with MedSigLIP (crossover +0.332 [+0.296, +0.365]; masking plus DFR 0.901). All thyroid and capsule findings replicate with a CNN (ConvNeXt-Base): crossovers +0.284 and +0.292; U-MtE – mask +0.128 [+0.112, +0.145] and erase – augment +0.120 [+0.102, +0.139] on thyroid.

577 F.5 Label-free and erasure baselines

Among methods that need no artifact labels, U-MtE beats JTT [Liu et al., 2021] by +0.198 [+0.167, +0.231] (thyroid), +0.137 [+0.098, +0.179] (ovary) and +0.094 [+0.077, +0.109] (chest drains, RAD-DINO; not with DINOv2, -0.010); JTT on the unmasked view is best on hair. Task-preserving concept removal (SPLINCE; Holstege et al., 2025), fitted with artifact labels in the whitened, minimal-distortion form of the original method, removes the artifact covariance at a large cost in disease signal (clean AUROC 0.56–0.78 versus 0.69–0.87 for masking) and still flips the shortcut on capsule endoscopy (reversed 0.868, correlated 0.683): when artifact and label are strongly associated in training, removing one removes much of the other. By min(reversed, correlated) U-MtE_{bal} exceeds SPLINCE on thyroid, capsule and hair (0.788 vs 0.596; 0.857 vs 0.683; 0.704 vs 0.564); on chest drains SPLINCE is higher (0.609 vs 0.460) but below DFR (0.683).

588 F.6 Text-prompted erasure

With vision–language backbones an artifact direction can be obtained from text alone (“a thyroid ultrasound image with measurement calipers” minus “a thyroid ultrasound image”; biased-prompt construction, Chuang et al., 2023). Erasing that subspace after masking is a near no-op in every cohort (+0.001 to +0.008 over masking; MedSigLIP thyroid, ovary, capsule and DermLIP hair), whereas U-MtE with the same MedSigLIP features gains +0.204 [+0.174, +0.237] (thyroid) and +0.215 [+0.182, +0.246] (ovary). The direction along which a described artifact moves the *text* embedding is not the direction along which the real artifact moves the *image* embedding; synthetic insertions measure the latter directly.

597 F.7 A single procedure: validation-driven selection

Selecting, per fold, the candidate with the best group-aware validation score (image-level artifact labels needed on validation images only) beats masking in all five cohorts (+0.17 to +0.40), is within 0.02 of the best candidate in each, and beats DFR in three (thyroid +0.057, capsule +0.047, ovary +0.010); it chooses U-MtE_{bal} where the artifact is an overlay and plain balancing on hair and drains, where DFR remains better. A split-validation variant that admits DFR as a candidate is reported in the supplement.

604 F.8 Modern inpainting

Inpainting the detected caliper pixels with LaMa [Suvorov et al., 2022] and then masking is a strong baseline when artifact masks are available (thyroid +0.226 [+0.213, +0.238], ovary +0.123 [+0.089, +0.155] over masking). U-MtE_{protect}, which needs no artifact mask, matches it on thyroid ($-0.009 [-0.045, +0.023]$) and is below it on ovary (-0.080) and on dermoscopic hair, where published hair masks are available (-0.086); with image-level artifact labels, U-MtE_{bal} exceeds it in all three (+0.162, +0.067, +0.123).

611 G Statistical plan and pre-registration

Every experiment was specified in a dated, publicly versioned document before any model was fitted on its data. The 16 primary tests (location crossover, protected U-MtE vs masking, erase vs augment, U-MtE_{bal} vs balanced; four cohorts each) are Holm-corrected; 12/16 are supported (Fig. 4). Everything else is secondary. Datasets considered and not used (CANDID-PTX: no access; breast

616 ultrasound: no artifact-free group; colonoscopy: no reflection-free images; ISIC ink, cardiomegaly
617 lines: too few images; pathology pen marks: unreliable masks) and the leakage analysis are docu-
618 mented in the repository.

619 H Compute

620 All experiments ran on one RTX 3090 (24 GB), 8 CPU cores and 31 GB RAM; the full pipeline
621 takes about two days, dominated by feature extraction and fine-tuning.

622 NeurIPS Paper Checklist

623 1. Claims

624 Question: Do the main claims made in the abstract and introduction accurately reflect the
625 paper’s contributions and scope?

626 Answer: [\[Yes\]](#)

627 Justification: The abstract and Sec. 1 state the location law, the theory, U-MtE and its limits;
628 each claim is backed in Sec. 6 and Appendices C–F, including the negative results.

629 Guidelines:

- 630 • The answer NA means that the abstract and introduction do not include the claims
631 made in the paper.
- 632 • The abstract and/or introduction should clearly state the claims made, including the
633 contributions made in the paper and important assumptions and limitations. A No or
634 NA answer to this question will not be perceived well by the reviewers.
- 635 • The claims made should match theoretical and experimental results, and reflect how
636 much the results can be expected to generalize to other settings.
- 637 • It is fine to include aspirational goals as motivation as long as it is clear that these
638 goals are not attained by the paper.

639 2. Limitations

640 Question: Does the paper discuss the limitations of the work performed by the authors?

641 Answer: [\[Yes\]](#)

642 Justification: Sec. 7 (Limitations) and the regime analysis in Sec. 6 state where masking
643 does not harm, where U-MtE fails and which datasets were infeasible.

644 Guidelines:

- 645 • The answer NA means that the paper has no limitation while the answer No means
646 that the paper has limitations, but those are not discussed in the paper.
- 647 • The authors are encouraged to create a separate "Limitations" section in their paper.
- 648 • The paper should point out any strong assumptions and how robust the results are to
649 violations of these assumptions (e.g., independence assumptions, noiseless settings,
650 model well-specification, asymptotic approximations only holding locally). The au-
651 thors should reflect on how these assumptions might be violated in practice and what
652 the implications would be.
- 653 • The authors should reflect on the scope of the claims made, e.g., if the approach was
654 only tested on a few datasets or with a few runs. In general, empirical results often
655 depend on implicit assumptions, which should be articulated.
- 656 • The authors should reflect on the factors that influence the performance of the ap-
657 proach. For example, a facial recognition algorithm may perform poorly when image
658 resolution is low or images are taken in low lighting. Or a speech-to-text system might
659 not be used reliably to provide closed captions for online lectures because it fails to
660 handle technical jargon.
- 661 • The authors should discuss the computational efficiency of the proposed algorithms
662 and how they scale with dataset size.
- 663 • If applicable, the authors should discuss possible limitations of their approach to ad-
664 dress problems of privacy and fairness.

- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: Sec. 5 states the assumptions (Gaussian class-conditional features, LDA-optimal linear head, additive artifact block); the derivation and a simulation check are in Appendix B.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [\[Yes\]](#)

Justification: Protocol, splits, seeds, thresholds and arms are specified in Sec. 3 and the appendices; every experiment has a dated pre-registration and a make target in the released code.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.

- (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All datasets are public; code, frozen splits, pre-registrations and a two-command tool are released (anonymised link at submission).

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so No is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: Sec. 3 and Appendix A: cohorts, backbones, regularisation grid, threshold selection on clean validation data, fine-tuning hyper-parameters.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All effects are reported with 95% hierarchical paired bootstrap CIs (10,000 replicates; clusters = seeds/folds); the primary family is Holm-corrected (Appendix G).

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix H: one RTX 3090, 8 CPU cores, 31 GB RAM; about two days end to end.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: Only public, de-identified datasets under their licences; no human subjects; results reported whichever way they went.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.

- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [\[Yes\]](#)

Justification: Sec. 7: masked models can be wrongly certified as shortcut-free; we recommend reporting artifact location and shortcut-conflicting performance. No foreseeable misuse beyond general model deployment risks.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [\[NA\]](#)

Justification: No high-risk models or data are released; the released artifacts are code and feature-level tools.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: Datasets are cited with licences in Appendix A and the repository (ISIC CC BY-NC, MMOTU and SEE-AI CC BY 4.0, NIH ChestX-ray14, TN3K/TNCD, NEATX, RANZCR-CLiP); backbones are used under their model licences.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The code base and the U-MtE/SLAS tools are documented in the repository README with usage examples and tests.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[NA\]](#)

Justification: No crowdsourcing or new human-subject data.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [NA]

Justification: Only publicly released, de-identified retrospective datasets are used; no new data were collected.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [NA]

Justification: LLMs are not a component of the proposed methods. (An AI coding assistant supported implementation and drafting under author supervision; authors should state this per venue policy.)

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.